

```

private JScrollPane createScroller(JTextArea textArea) {
    JScrollPane scroller = new JScrollPane(textArea);
    scroller.setVerticalScrollBarPolicy(ScrollPaneConstants.VERTICAL_SCROLLBAR_ALWAYS);
    scroller.setHorizontalScrollBarPolicy(ScrollPaneConstants.HORIZONTAL_SCROLLBAR_NEVER);
    return scroller;
}

private JTextArea createTextArea(Font font) {
    JTextArea textArea = new JTextArea(6, 20);
    textArea.setLineWrap(true);
    textArea.setWrapStyleWord(true);
    textArea.setFont(font);
    return textArea;
}

private void nextCard() {
    QuizCard card = new QuizCard(question.getText(), answer.getText());
    cardList.add(card);
    clearCard();
}

private void saveCard() {
    QuizCard card = new QuizCard(question.getText(), answer.getText());
    cardList.add(card);

    JFileChooser fileSave = new JFileChooser();
    fileSave.showSaveDialog(frame);
    saveFile(fileSave.getSelectedFile()); ←
}

```

Creating a scroll pane or a text area needs a lot of similar-looking code. We've put the code into a couple of helper methods that we can call when we need a text area or scroll pane.

Brings up a file dialog box and waits on this line until the user chooses 'Save' from the dialog box. All the file dialog navigation and selecting a file, etc., is done for you by the JFileChooser! It really is this easy.

private void clearAll() {

```

    cardList.clear();
    clearCard();
}

```

When we want a new set of cards, we need to clear out the card list AND the text areas.

private void clearCard() {

```

    question.setText("");
    answer.setText("");
    question.requestFocus();
}

```

The method that does the actual file writing (called by the SaveMenuItem's event handler). The argument is the 'File' object the user is saving. We'll look at the File class on the next page.

private void saveFile(File file) {

```

try {
    BufferedWriter writer = new BufferedWriter(new FileWriter(file));
    for (QuizCard card : cardList) {
        writer.write(card.getQuestion() + "/");
        writer.write(card.getAnswer() + "\n");
    }
    writer.close();
} catch (IOException e) {
    System.out.println("Couldn't write the cardList out: " + e.getMessage());
}
}

```

We chain a BufferedWriter on to a new FileWriter to make writing more efficient. (We'll talk about that in a few pages.)

Walk through the ArrayList of cards and write them out, one card per line, with the question and answer separated by a "/", and then add a newline character ("\n").

The `java.io.File` class

The `java.io.File` class is another example of an older class in the Java API. It's been "replaced" by two classes in the newer `java.nio.file` package, but you'll undoubtedly encounter code that uses the `File` class. **For new code, we recommend using the `java.nio.file` package instead of the `java.io.File` class.** In a few pages, we'll take a look at a few of the most important capabilities in the `java.nio.file` package. With that said...

The `java.io.File` class *represents* a file on disk but doesn't actually represent the *contents* of the file. What? Think of a File object as something more like a *path name* of a file (or even a *directory*) rather than The Actual File Itself. The File class does not, for example, have methods for reading and writing. One VERY useful thing about a File object is that it offers a much safer way to represent a file than just using a String filename. For example, most classes that take a String filename in their constructor (like `FileWriter` or `FileInputStream`) can take a File object instead. You can construct a File object, verify that you've got a valid path, etc., and then give that File object to the `FileWriter` or `FileInputStream`.

Some things you can do with a File object:

- ① Make a File object representing an existing file

```
File f = new File("MyCode.txt");
```

- ② Make a new directory

```
File dir = new File("Chapter7");
dir.mkdir();
```

- ③ List the contents of a directory

```
if (dir.isDirectory()) {
    String[] dirContents = dir.list();
    for (String dirContent : dirContents) {
        System.out.println(dirContent);
    }
}
```

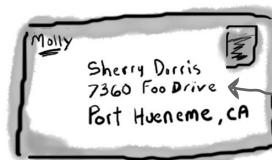
- ④ Delete a file or directory (returns true if successful)

```
boolean isDeleted = f.delete();
```

A File object represents the name and path of a file or directory on disk, for example:

`/Users/Kathy/Data/Game.txt`

But it does NOT represent, or give you access to, the data in the file!



An address is NOT the same as the actual house! A File object is like a street address... it represents the name and location of a particular file, but it isn't the file itself.

A File object represents the filename "GameFile.txt"

GameFile.txt

```
50,Elf,bow,sword,dust
200,Troll,bare hands,big ax
120,Magician,spells,invisibility
```

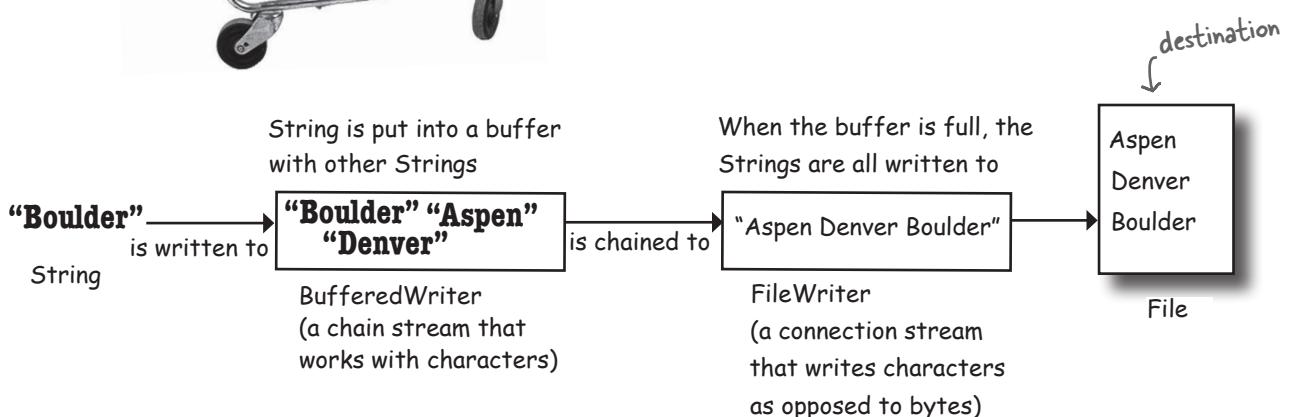
A File object does NOT represent (or give you direct access to) the data inside the file!

The beauty of buffers

If there were no buffers, it would be like shopping without a cart. You'd have to carry each thing out to your car, one soup can or toilet paper roll at a time.



Buffers give you a temporary holding place to group things until the holder (like the cart) is full. You get to make far fewer trips when you use a buffer.



```
BufferedWriter writer = new BufferedWriter(new FileWriter(aFile));
```

Using buffers is *much* more efficient than working without them. You can write to a file using FileWriter alone, by calling `write(someString)`, but FileWriter writes each and every thing you pass to the file each and every time. That's overhead you don't want or need, since every trip to the disk is a Big Deal compared to manipulating data in memory. By chaining a BufferedWriter onto a FileWriter, the BufferedWriter will hold all the stuff you write to it until it's full. *Only when the buffer is full will the FileWriter actually be told to write to the file on disk.*

If you do want to send data *before* the buffer is full, you do have control.

Just Flush It. Calls to `writer.flush()` say, “send whatever’s in the buffer, **now!**”

Notice that we don't even need to keep a reference to the FileWriter object. The only thing we care about is the BufferedWriter, because that's the object we'll call methods on, and when we close the BufferedWriter, it will take care of the rest of the chain.

Reading from a text file

Reading text from a file is simple, but this time we'll use a File object to represent the file, a FileReader to do the actual reading, and a BufferedReader to make the reading more efficient.

The read happens by reading lines in a *while* loop, ending the loop when the result of a readLine() is null. That's the most common style for reading data (pretty much anything that's not a Serialized object): read stuff in a while loop (actually a while loop *test*), terminating when there's nothing left to read (which we know because the result of whatever read method we're using is null).

```

import java.io.*;           Don't forget the import.

class ReadAFile {
    public static void main(String[] args) {
        try {
            File myFile = new File("MyText.txt");
            FileReader fileReader = new FileReader(myFile);

            BufferedReader reader = new BufferedReader(fileReader);
            Make a String variable to hold
            each line as the line is read
            String line;
            while ((line = reader.readLine()) != null) {
                System.out.println(line);
            }
            reader.close();
        } catch (IOException e) {
            e.printStackTrace();
        }
    }
}
  
```

A FileReader is a connection stream for characters that connects to a text file.

Chain the FileReader to a BufferedReader for more efficient reading. It'll go back to the file to read only when the buffer is empty (because the program has read everything in it).

This says, "Read a line of text, and assign it to the String variable "line." While that variable is not null (because there WAS something to read), print out the line that was just read."

Or another way of saying it, "While there are still lines to read, read them and print them."

A file with two lines of text.

What's $2 + 27/4$

What's $20+22/42$

MyText.txt

Java 8 Streams and I/O

If you're using Java 8 and you feel comfortable using the Streams API, you can replace all the code inside the try block with the following:

```

Files.lines(Path.of("MyText.txt"))
    .forEach(line -> System.out.println(line));
  
```

We'll see the Files and Path classes later in this chapter.

Quiz Card Player (code outline)

```
public class QuizCardPlayer {  
  
    public void go() {  
        // build and display gui  
    }  
  
    private void nextCard() {  
        // if this is a question, show the answer, otherwise show  
        // next question set a flag for whether we're viewing a  
        // question or answer  
    }  
  
    private void open() {  
        // bring up a file dialog box  
        // let the user navigate to and choose a card set to open  
    }  
  
    private void loadFile(File file) {  
        // must build an ArrayList of cards, by reading them from  
        // a text file called from the OpenMenuItem event handler,  
        // reads the file one line at a time and tells the makeCard()  
        // method to make a new card out of the line (one line in the  
        // file holds both the question and answer, separated by a "/")  
    }  
  
    private void makeCard(String lineToParse) {  
        // called by the loadFile method, takes a line from the text file  
        // and parses into two pieces—question and answer—and creates a  
        // new QuizCard and adds it to the ArrayList called CardList  
    }  
}
```

Quiz Card Player code

```
import javax.swing.*;
import java.awt.*;
import java.io.*;
import java.util.ArrayList;

public class QuizCardPlayer {
    private ArrayList<QuizCard> cardList;
    private int currentCardIndex;
    private QuizCard currentCard;
    private JTextArea display;
    private JFrame frame;
    private JButton nextButton;
    private boolean isShowAnswer;

    public static void main(String[] args) {
        QuizCardPlayer reader = new QuizCardPlayer();
        reader.go();
    }

    public void go() {
        frame = new JFrame("Quiz Card Player");
        JPanel mainPanel = new JPanel();
        Font bigFont = new Font("sanserif", Font.BOLD, 24);

        display = new JTextArea(10, 20);
        display.setFont(bigFont);
        display.setLineWrap(true);
        display.setEditable(false);

        JScrollPane scroller = new JScrollPane(display);
        scroller.setVerticalScrollBarPolicy(ScrollPaneConstants.VERTICAL_SCROLLBAR_ALWAYS);
        scroller.setHorizontalScrollBarPolicy(ScrollPaneConstants.HORIZONTAL_SCROLLBAR_NEVER);
        mainPanel.add(scroller);

        nextButton = new JButton("Show Question");
        nextButton.addActionListener(e -> nextCard());
        mainPanel.add(nextButton);

        JMenuBar menuBar = new JMenuBar();
        JMenu fileMenu = new JMenu("File");
        JMenuItem loadMenuItem = new JMenuItem("Load card set");
        loadMenuItem.addActionListener(e -> open());
        fileMenu.add(loadMenuItem);
        menuBar.add(fileMenu);
        frame.setJMenuBar(menuBar);

        frame.getContentPane().add(BorderLayout.CENTER, mainPanel);
        frame.setSize(500, 400);
        frame.setVisible(true);
    }
}
```

*Just GUI code on this page;
nothing special.*

```

private void nextCard() {
    if (isShowAnswer) {
        // show the answer because they've seen the question
        display.setText(currentCard.getAnswer());
        nextButton.setText("Next Card");
        isShowAnswer = false;
    } else { // show the next question
        if (currentCardIndex < cardList.size()) {
            showNextCard();
        } else {
            // there are no more cards!
            display.setText("That was last card");
            nextButton.setEnabled(false);
        }
    }
}

```

Check the isShowAnswer boolean flag to see if they're currently viewing a question or an answer, and do the appropriate thing depending on the answer.

```

private void open() {
    JFileChooser fileOpen = new JFileChooser();
    fileOpen.showOpenDialog(frame);
    loadFile(fileOpen.getSelectedFile());
}

```

Bring up the file dialog box and let them navigate to and choose the file to open.

```

private void loadFile(File file) {
    cardList = new ArrayList<>();
    currentCardIndex = 0;
    try {
        BufferedReader reader = new BufferedReader(new FileReader(file));
        String line;
        while ((line = reader.readLine()) != null) {
            makeCard(line);
        }
        reader.close();
    } catch (IOException e) {
        System.out.println("Couldn't write the cardList out: " + e.getMessage());
    }
    showNextCard(); ← Now time to start,
}

```

Make a BufferedReader chained to a new FileReader, giving the FileReader the File object the user chose from the open file dialog.

} Read a line at a time, passing the line to the makeCard() method that parses it and turns it into a real QuizCard and adds it to the ArrayList.

```

private void makeCard(String lineToParse) {
    String[] result = lineToParse.split("/");
    QuizCard card = new QuizCard(result[0], result[1]);
    cardList.add(card);
    System.out.println("made a card");
}

```

```

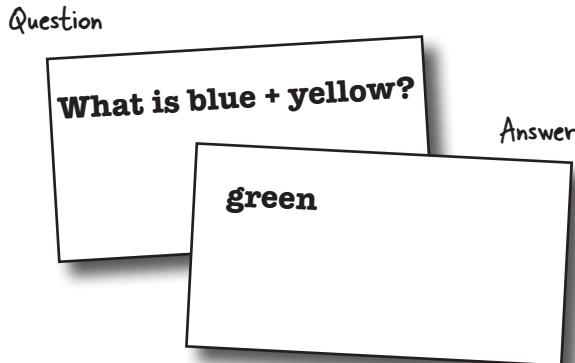
private void showNextCard() {
    currentCard = cardList.get(currentCardIndex);
    currentCardIndex++;
    display.setText(currentCard.getQuestion());
    nextButton.setText("Show Answer");
    isShowAnswer = true;
}

```

Each line of text corresponds to a single flashcard, but we have to parse out the question and answer as separate pieces. We use the String split() method to break the line into two tokens (one for the question and one for the answer). We'll look at the split() method on the next page.

Parsing with String split()

Imagine you have a flashcard like this:



Saved in a question file like this:

What is blue + yellow?/green
What is red + blue?/purple

How do you separate the question and answer?

When you read the file, the question and answer are smooshed together in one line, separated by a forward slash "/" (because that's how we wrote the file in the QuizCardBuilder code).

String split() lets you break a String into pieces.

The split() method says, "give me a separator, and I'll break out all the pieces of this String for you and put them in a String array."

What is blue + yellow?

token 1



token 2

green

separator

In the QuizCardPlayer app, this is what a single line looks like when it's read in from the file.

```
String toTest = "What is blue + yellow?/green";
String[] result = toTest.split("/");
for (String token : result) {
    System.out.println(token);
}
```

The split() method takes the "/" and uses it to break apart the String into (in this case) two pieces, token 1 and token 2. (Note: split() is FAR more powerful than what we're using it for here. It can do extremely complex parsing with filters, wildcards, etc.)

Loop through the array and print each token (piece). In this example, there are only two tokens: "What is blue + yellow?" and "green."

there are no Dumb Questions

Q: OK, I look in the API and there are about five million classes in the `java.io` package. How the heck do you know which ones to use?

A: The I/O API uses the modular “chaining” concept so that you can hook together connection streams and chain streams (also called “filter” streams) in a wide range of combinations to get just about anything you could want.

The chains don’t have to stop at two levels; you can hook multiple chain streams to one another to get just the right amount of processing you need.

Most of the time, though, you’ll use the same small handful of classes. If you’re writing text files, `BufferedReader` and `BufferedWriter` (chained to `FileReader` and `FileWriter`) are probably all you need. If you’re writing serialized objects, you can use `ObjectOutputStream` and `ObjectInputStream` (chained to `FileInputStream` and `FileOutputStream`).

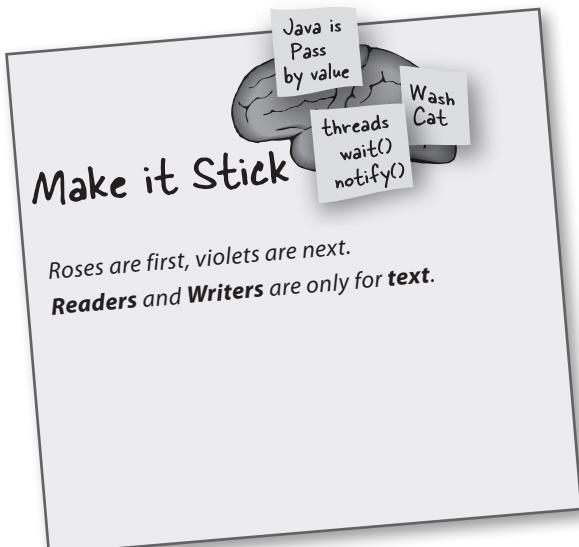
In other words, 90% of what you might typically do with Java I/O can use what we’ve already covered.

Q: You just said we’ve already learned 90% of what we’ll probably use, but we haven’t seen the fabled NIO.2 stuff yet. What gives?

A: NIO.2 is coming up on the very next page! But for reading and writing text files, `BufferedReader`s and `BufferedWriter`s are still usually the way to go. So we’ll be looking at how NIO.2 makes using them easier.

Q: My brain is a little tired, and I’ve heard NIO.2 is pretty complicated.

A: We’re going to focus on a few key concepts in the `java.nio.file` package.



BULLET POINTS

- To write a text file, start with a `FileWriter` connection stream.
- Chain the `FileWriter` to a `BufferedWriter` for efficiency.
- A `File` object represents a file at a particular path, but does not represent the actual contents of the file.
- With a `File` object you can create, traverse, and delete directories.
- Most streams that can use a `String` filename can use a `File` object as well, and a `File` object can be safer to use.
- To read a text file, start with a `FileReader` connection stream.
- Chain the `FileReader` to a `BufferedReader` for efficiency.
- To parse a text file, you need to be sure the file is written with some way to recognize the different elements. A common approach is to use some kind of character to separate the individual pieces.
- Use the `String split()` method to split a `String` up into individual tokens. A `String` with one separator will have two tokens, one on each side of the separator. *The separator doesn’t count as a token.*

NIO.2 and the `java.nio.file` package

Java NIO.2 is usually taken to mean two packages added in Java 7:

`java.nio.file`

`java.nio.file.attribute`

The `java.nio.file.attribute` package lets you manipulate the *metadata* associated with a computer's files and directories. For example, you would use the classes in this package if you wanted to read or change a file's permissions settings. We WON'T be discussing this package further. (phew)

The `java.nio.file` package is all you need to do common text file reading and writing, and it also provides you with the ability to manipulate a computer's directories and directory structure. Most of the time you'll use three types in `java.nio.file`:

- The Path interface: You'll always need a Path object to locate the directories or files you want to work with.
- The Paths class: You'll use the `Paths.get()` method to make the Path object you'll need when you use methods in the Files class.
- The Files class: This is the class whose (static) methods do all the work you'll want to do: making new Readers and Writers, and creating, modifying, and searching through directories and files on file systems.

A mini-tutorial, creating a `BufferedWriter` with NIO.2

- ① Import Path, Paths, and Files:

```
import java.nio.file.*;
```

- ② Make a Path object using the Paths class:

```
Path myPath = Paths.get("MyFile.txt");
```

Or, if the file is in a subdirectory like:

/myApp/files/MyFile.txt :

```
Path myPath = Paths.get("/myApp", "files", "MyFile.txt");
```

A Path object is used to locate a file on a computer (i.e., in the file system). A path can be used to locate files in the current directory or in other directories.

- ③ Make a new BufferedWriter using a Path and the Files class:

```
BufferedWriter writer = Files.newBufferedWriter(myPath);
```

The "/" in "/myApp" is called the name-separator. Depending on which OS you're using, your name-separator might be different; for example, it might be "\".

Somewhere—under the covers—some method is saying:

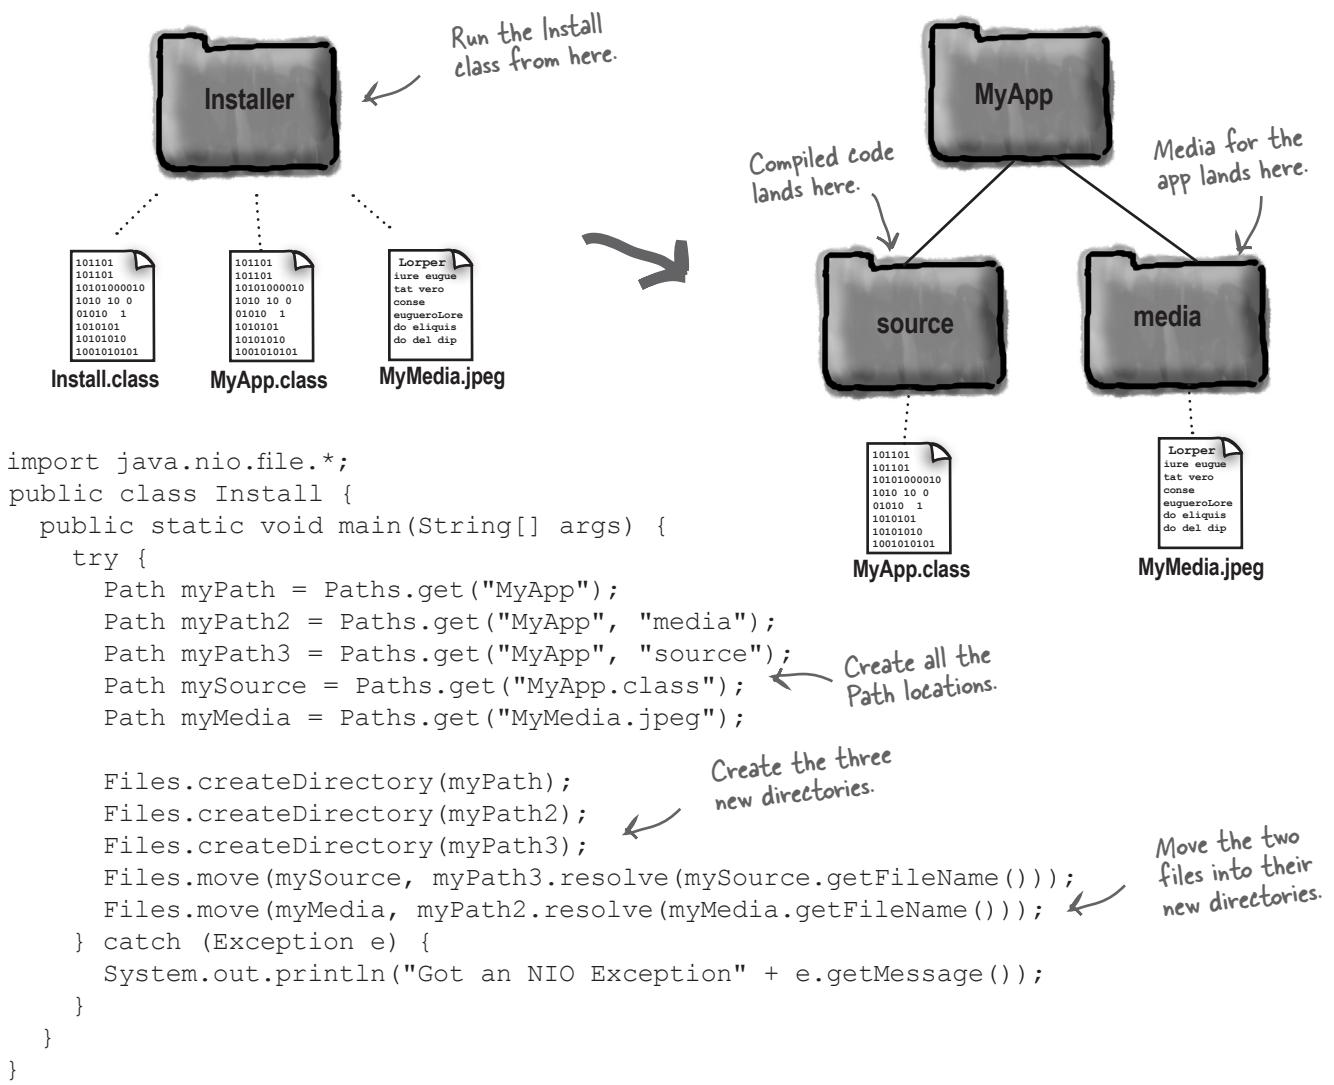
```
BufferedWriter writer =
new BufferedWriter(..)
```

Path, Paths, and Files (messing with directories)

In Appendix B, we'll be discussing how to split your Java app into packages. This includes creating the proper directory structure for all of your app's files. In most cases you'll make and move directories and files by hand, using the command line or utilities like the Finder or Windows Explorer. But you can also do it from within your Java code.

Warning! Goofing around with directories in a Java program is a real “can of worms” topic. To do it correctly you need to learn about paths, absolute paths, relative paths, OS permissions, file attributes, and on and on. Below is a greatly simplified example of messing around with directories, just to give you a feel for what's possible.

Suppose you wanted to make an installer program to install your killer app. You start with the directory and files on the left, and want to end up with the directory structure and files on the right.



Finally, a closer look at finally

Several chapters ago we looked at how try-catch-finally worked. Kind of. All we said about finally was that it was a good place to put your “cleanup code.” That’s true, but let’s get more specific. Most of the time, when we talk about “cleanup code,” we mean closing resources we borrowed from the operating system. When we open a file or a socket, the OS is giving us some of its resources. When we’re done with them, we need to give them back. Below is a snippet of code from the QuizCardBuilder class. We highlighted a call to a constructor and three separate method calls...

That's FOUR places an exception can be thrown!

```
private void saveFile(File file) {  
    try {  
        BufferedWriter writer = new BufferedWriter(new FileWriter(file));  
        for (QuizCard card : cardList) {  
            writer.write(card.getQuestion() + "/");  
            writer.write(card.getAnswer() + "\n");  
        }  
        writer.close();  
    } catch (IOException e) {  
        System.out.println("Couldn't write the cardList out: " + e.getMessage());  
    }  
}
```

The code is annotated with four arrows pointing to specific lines:

- A curved arrow points to the line `new FileWriter(file))`, indicating it's a constructor call.
- A horizontal arrow points to the line `writer.write(card.getAnswer() + "\n");`, indicating it's a method call.
- A horizontal arrow points to the line `writer.close();`, indicating it's another method call.
- A curved arrow points to the line `catch (IOException e)`, indicating it's a catch block.

A handwritten note to the right of the annotations reads: **All the places an exception could be thrown!**

If the call to make a new `FileWriter` fails, if ANY of the many `write()` invocations fail, or the `close()` itself fails, an exception will be thrown, the JVM will jump to the catch block, and the writer will never be closed. Yikes!

Remember, finally **ALWAYS** runs!!

Since we REALLY want to make sure we close the writer file, let’s put the `close()` invocation in a finally block.



Sharpen your pencil

→ Yours to solve.

Coding the new finally block

What changes will we have to make to the code above to move the `close()` to a finally block?

There might be more than you first imagine.

Finally, a closer look at finally, cont.

The amount of code required to put the close() in the finally block might surprise you; let's take a look.

```
private void saveFile(File file) {
    BufferedWriter writer = null; ←
    try {
        writer = new BufferedWriter(new FileWriter(file));
        for (QuizCard card : cardList) {
            writer.write(card.getQuestion() + "/");
            writer.write(card.getAnswer() + "\n");
        }
        writer.close();
    } catch (IOException e) {
        System.out.println("Couldn't write the cardList out: " + e.getMessage());
    } finally {
        try {
            writer.close(); ←
        } catch (Exception e) {
            System.out.println("Couldn't close writer: " + e.getMessage());
        }
    }
}
```

We had to declare the writer reference outside of the try block so that it's visible in the finally block.

Yup, we had to put the close() in yet another try-catch block!

Are you kidding me right now?
I have to write all of this code
every time I want to do a little
I/O? Verbose much?



There IS a better way!

In the early days of Java, this is how you had to make sure you were really closing a file. You are very likely to encounter finally blocks that look like this when you're looking at existing code. But for new code, there is a better way:

Try-With-Resources

We'll look at that next.

The try-with-resources (TWR) statement

If you're using Java 7 or later (and we sure hope you are!), you can use the try-with-resources version of try statements to make doing I/O easier. Let's compare the try code we've been looking at with try-with-resources code that does the same thing:

```
private void saveFile(File file) {
    BufferedWriter writer = null;
    try {
        writer = new BufferedWriter(new FileWriter(file));

        for (QuizCard card : cardList) {
            writer.write(card.getQuestion() + "/");
            writer.write(card.getAnswer() + "\n");
        }
    } catch (IOException e) {
        System.out.println("Couldn't write the cardList out: " + e.getMessage());
    } finally {
        try {
            writer.close();
        } catch (Exception e) {
            System.out.println("Couldn't close writer: " + e.getMessage());
        }
    }
}
```

*Old style,
try-catch-finally
code*

```
private void saveFile(File file) {
    try (BufferedWriter writer =
        new BufferedWriter(new FileWriter(file))) {

        for (QuizCard card : cardList) {
            writer.write(card.getQuestion() + "/");
            writer.write(card.getAnswer() + "\n");
        }
    } catch (IOException e) {
        System.out.println("Couldn't write the cardList out: " + e.getMessage());
    }
}
```

*Modern,
try-with-resources
code*

*there are no
Dumb Questions*

Q: Wait, what? You told us that a try statement needs a catch and/or a finally?

A: Nice catch! It turns out that when you use try-with-resources, the compiler makes a finally block for you. You can't see it, but it's there.

Autocloseable, the very small catch

On the last page we saw a different kind of try statement, the try-with-resources statement (TWR). Let's take a look at how to write and use TWR statements by first, deconstructing the following:

```
try (BufferedWriter writer =
      new BufferedWriter(new FileWriter(file))) {
```

Writing a try-with-resources statement

- ① Add a set of parentheses between "try" and "{":

```
try ( ... ) {
```

- ② Inside the parentheses, declare an object whose type implements Autocloseable:

```
try (BufferedWriter writer =
      new BufferedWriter(new FileWriter(file))) {
```

Like all of the I/O classes we've been using this chapter, BufferedWriter implements Autocloseable.

- ③ Use the object you declared inside the try block (just like you always did):

```
writer.write(card.getQuestion() + "/");
writer.write(card.getAnswer() + "\n");
```

Autocloseable, it's everywhere you do I/O

Autocloseable is an interface that was added to `java.lang` in Java 7. Almost all of the I/O you're ever going to do uses classes that implement Autocloseable. You mostly won't have to think about it.

There are a few more things worth knowing about TWR statements:

- You can declare and use more than one I/O resource in a single TWR block:

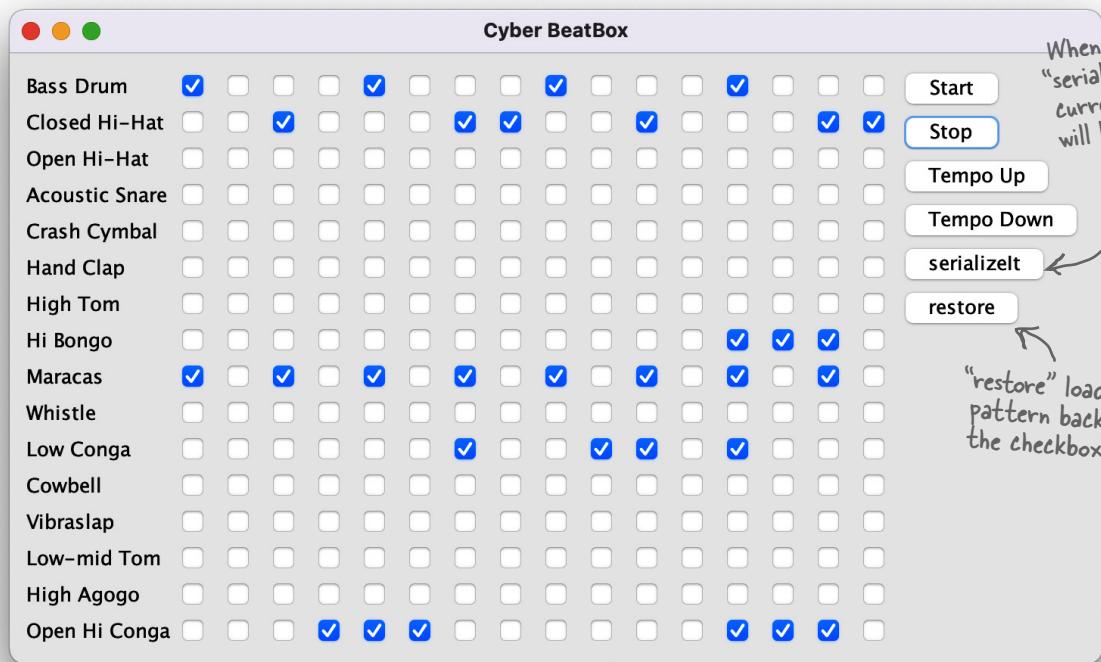
```
try (BufferedWriter writer =
      new BufferedWriter(new FileWriter(file));
      BufferedReader reader =
      new BufferedReader(new FileReader(file))) {
```

Separate the resources using semicolons, ;.

- If you declare more than one resource, they will be closed in the order OPPOSITE to which they were declared; i.e., first declared is last closed.
- If you add catch or finally blocks, the system will handle multiple close() invocations gracefully.

ONLY classes that implement Autocloseable can be used in TWR statements!

Code Kitchen



Let's make the BeatBox save and restore our favorite pattern.

Saving a BeatBox pattern

Remember, in the BeatBox, a drum pattern is nothing more than a bunch of checkboxes. When it's time to play the sequence, the code walks through the checkboxes to figure out which drums sounds are playing at each of the 16 beats. So to save a pattern, all we need to do is save the state of the checkboxes.

We can make a simple boolean array, holding the state of each of the 256 checkboxes. An array object is serializable as long as the things *in* the array are serializable, so we'll have no trouble saving an array of booleans.

To load a pattern back in, we read the single boolean array object (deserialize it) and restore the checkboxes. Most of the code you've already seen, in the Code Kitchen where we built the BeatBox GUI, so in this chapter, we look at only the save and restore code.

This CodeKitchen gets us ready for the next chapter, where instead of writing the pattern to a *file*, we send it over the *network* to the server. And instead of loading a pattern *in* from a file, we get patterns from the *server*, each time a participant sends one to the server.

Serializing a pattern

This is a method in the BeatBox code. We can call this from a lambda expression when we add an ActionListener to the serialize button, or create an ActionListener inner class that calls this.

```
private void writeFile() {
    boolean[] checkboxState = new boolean[256];
    for (int i = 0; i < 256; i++) {
        JCheckBox check = checkboxList.get(i);
        if (check.isSelected()) {
            checkboxState[i] = true;
        }
    }
    try (ObjectOutputStream os =
        new ObjectOutputStream(new FileOutputStream("Checkbox.ser"))) {
        os.writeObject(checkboxState);
    } catch (IOException e) {
        e.printStackTrace();
    }
}
```

← Make a boolean array to hold the state of each checkbox.

Walk through the checkboxList (ArrayList of checkboxes), get the state of each one, and add it to the boolean array.

Try-with-resources

This part's a piece of cake. Just write/serialize the one boolean array!

Restoring a BeatBox pattern

This is pretty much the save in reverse...read the boolean array and use it to restore the state of the GUI checkboxes. It all happens when the user hits the “restore” button.

Restoring a pattern

This is another method in the BeatBox class.

```
private void readFile() {
    boolean[] checkboxState = null;
    try (ObjectInputStream is =
        new ObjectInputStream(new FileInputStream("Checkbox.ser"))) {
        checkboxState = (boolean[]) is.readObject(); ← Try-with-resources
    } catch (Exception e) {
        e.printStackTrace();
    }

    for (int i = 0; i < 256; i++) {
        JCheckBox check = checkboxList.get(i);
        check.setSelected(checkboxState[i]); } Now restore the state of each of the checkboxes in the ArrayList of actual JCheckBox objects (checkboxList).
    }

    sequencer.stop();
    buildTrackAndStart(); } Now stop whatever is currently playing, and rebuild the sequence using the new state of the checkboxes in the ArrayList.
}
```

Read the single object in the file (the boolean array) and cast it back to a boolean array (remember, readObject() returns a reference of type Object).



→ Yours to solve.

This version has a huge limitation! When you hit the “serializelt” button, it serializes automatically, to a file named “Checkbox.ser” (which gets created if it doesn’t exist). But each time you save, you overwrite the previously saved file.

Improve the save and restore feature by incorporating a JFileChooser so that you can name and save as many different patterns as you like, and load/restore from *any* of your previously saved pattern files.



→ Yours to solve.

Can they be saved?

Which of these do you think are, or should be, serializable? If not, why not? Not meaningful?

Security risk? Only works for the current execution of the JVM? Make your best guess, without looking it up in the API.

Object type	Serializable?	If not, why not?
Object	Yes / No	_____
String	Yes / No	_____
File	Yes / No	_____
Date	Yes / No	_____
OutputStream	Yes / No	_____
JFrame	Yes / No	_____
Integer	Yes / No	_____
System	Yes / No	_____

What's Legal?

Circle the code fragments that would compile (assuming they're within a legal class).



→ Yours to solve.

```
FileReader fileReader = new FileReader();
BufferedReader reader = new BufferedReader(fileReader);
```

```
FileOutputStream f = new FileOutputStream("Foo.ser");
ObjectOutputStream os = new ObjectOutputStream(f);
```

```
BufferedReader reader = new BufferedReader(new FileReader(file));
String line;
while ((line = reader.readLine()) != null) {
    makeCard(line);
}
```

```
FileOutputStream f = new FileOutputStream("Game.ser");
ObjectInputStream is = new ObjectInputStream(f);
GameCharacter oneAgain = (GameCharacter) is.readObject();
```

exercise: True or False



This chapter explored the wonderful world of Java I/O. Your job is to decide whether each of the following I/O-related statements is true or false.

TRUE OR FALSE

1. Serialization is appropriate when saving data for non-Java programs to use.
2. Object state can be saved only by using serialization.
3. ObjectOutputStream is a class used to save serializable objects.
4. Chain streams can be used on their own or with connection streams.
5. A single call to writeObject() can cause many objects to be saved.
6. All classes are serializable by default.
7. The java.nio.file.Path class can be used to locate files.
8. If a superclass is not serializable, then the subclass can't be serializable.
9. Only classes that implement AutoCloseable can be used in try-with-resources statements.
10. When an object is deserialized, its constructor does not run.
11. Both serialization and saving to a text file can throw exceptions.
12. BufferedWriter can be chained to FileWriter.
13. File objects represent files, but not directories.
14. You can't force a buffer to send its data before it's full.
15. Both file readers and file writers can optionally be buffered.
16. The methods on the Files class let you operate on files and directories.
17. Try-with-resources statements cannot include explicit finally blocks.

—————> Answers on page 584.



Code Magnets

This one's tricky, so we promoted it from an Exercise to full Puzzle status. Reconstruct the code snippets to make a working Java program that produces the output listed below. (You might not need all of the magnets, and you may reuse a magnet more than once.)

```

class DungeonGame implements Serializable {
    try {
        FileOutputStream fos = new
            FileOutputStream("dg.ser");
        e.printStackTrace();
        short getZ() {
            return z;
        }
        oos.close();
        int getX() {
            return x;
        }
        System.out.println(d.getX() + d.getY() + d.getZ());
    }
    FileInputStream fis = new
        FileInputStream("dg.ser");
    public int x = 3;
    transient long y = 4;
    private short z = 5;
    long getY() {
        return y;
    }
    class DungeonTest {
        import java.io.*;
        ois.close();
        fos.writeObject(d);
        } catch (Exception e) {
        d = (DungeonGame) ois.readObject();
    }
}

```

```

File Edit Window Help Torture
% java DungeonTest
12
8

```

```

ObjectOutputStream oos = new
    ObjectOutputStream(fos);
    oos.writeObject(d);
    public static void main(String[] args) {
        DungeonGame d = new DungeonGame();
}

```

→ Answers on page 585.



Exercise Solutions

TRUE OR FALSE

(from page 582)

- | | |
|--|--------------|
| 1. Serialization is appropriate when saving data for non-Java programs to use. | False |
| 2. Object state can be saved only by using serialization. | False |
| 3. ObjectOutputStream is a class used to save serializable objects. | True |
| 4. Chain streams can be used on their own or with connection streams. | False |
| 5. A single call to writeObject() can cause many objects to be saved. | True |
| 6. All classes are serializable by default. | False |
| 7. The java.nio.file.Path class can be used to locate files. | False |
| 8. If a superclass is not serializable, then the subclass can't be serializable. | False |
| 9. Only classes that implement AutoCloseable can be used in try-with-resources statements. | True |
| 10. When an object is deserialized, its constructor does not run. | True |
| 11. Both serialization and saving to a text file can throw exceptions. | True |
| 12. BufferedWriter can be chained to FileWriters. | True |
| 13. File objects represent files, but not directories. | False |
| 14. You can't force a buffer to send its data before it's full. | False |
| 15. Both file readers and file writers can optionally be buffered. | True |
| 16. The methods on the Files class let you operate on files and directories. | True |
| 17. Try-with-resources statements cannot include explicit finally blocks. | False |



Good thing we're
finally at the answers.
I was gettin' kind of
tired of this chapter.

Code Magnets

(from page 583)

```

import java.io.*;

class DungeonGame implements Serializable {
    public int x = 3;
    transient long y = 4;
    private short z = 5;

    int getX() {
        return x;
    }
    long getY() {
        return y;
    }
    short getZ() {
        return z;
    }
}

class DungeonTest {
    public static void main(String[] args) {
        DungeonGame d = new DungeonGame();
        System.out.println(d.getX() + d.getY() + d.getZ());
        try {
            FileOutputStream fos = new FileOutputStream("dg.ser");
            ObjectOutputStream oos = new ObjectOutputStream(fos);
            oos.writeObject(d);
            oos.close();

            FileInputStream fis = new FileInputStream("dg.ser");
            ObjectInputStream ois = new ObjectInputStream(fis);
            d = (DungeonGame) ois.readObject();
            ois.close();
        } catch (Exception e) {
            e.printStackTrace();
        }
        System.out.println(d.getX() + d.getY() + d.getZ());
    }
}

```

```

File Edit Window Help Escape
% java DungeonTest
12
8
}

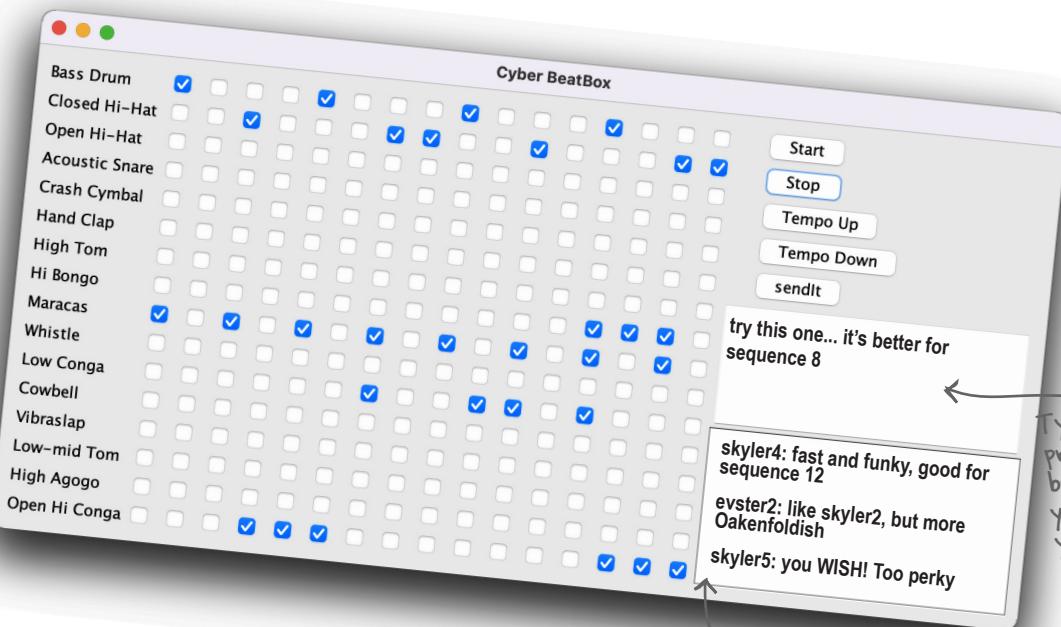
```


Make a Connection



Connect with the outside world. Your Java program can talk to a program on another machine. It's easy. All the low-level networking details are taken care of by the built-in Java libraries. One of Java's big benefits is that sending and receiving data over a network can be just I/O with a slightly different connection at the end of the I/O chain. In this chapter we'll connect to the outside world with *channels*. We'll make *client* channels. We'll make *server* channels. We'll make *clients* and *servers*, and we'll make them talk to each other. And we'll also have to learn how to do more than one thing at once. Before the chapter's done, you'll have a fully functional, multithreaded chat client. Did we just say *multithreaded*? Yes, now you *will* learn the secret of how to talk to Bob while simultaneously listening to Suzy.

Real-time BeatBox chat

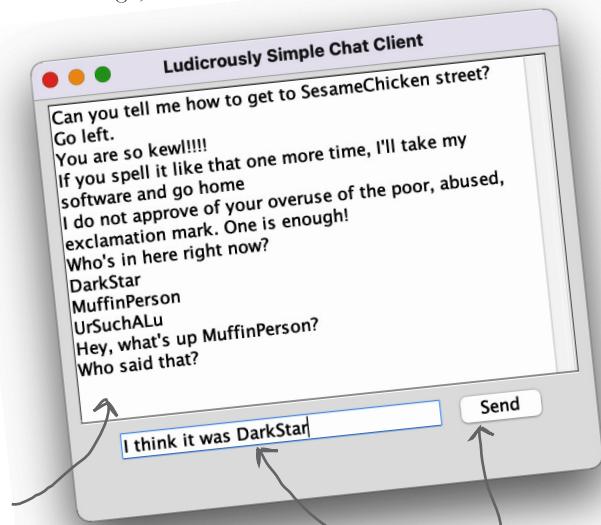


Type a message and press the sendit button to send your message AND your current beat pattern.

Clicking on a received message loads the pattern that went with it.

You're working on a computer game. You and your team are doing the sound design for each part of the game. Using a “chat” version of the BeatBox, your team can collaborate—you can send a beat pattern along with your chat message, and everybody in the BeatBox Chat gets it. So you don’t just get to *read* the other participants’ messages; you get to load and *play* a beat pattern simply by clicking the message in the incoming messages area.

In this chapter we’re going to learn what it takes to make a chat client like this. We’re even going to learn a little about making a chat *server*. We’ll save the full BeatBox Chat for the Code Kitchen, but in this chapter you *will* write a Ludicrously Simple Chat Client and Very Simple Chat Server that send and receive text messages.



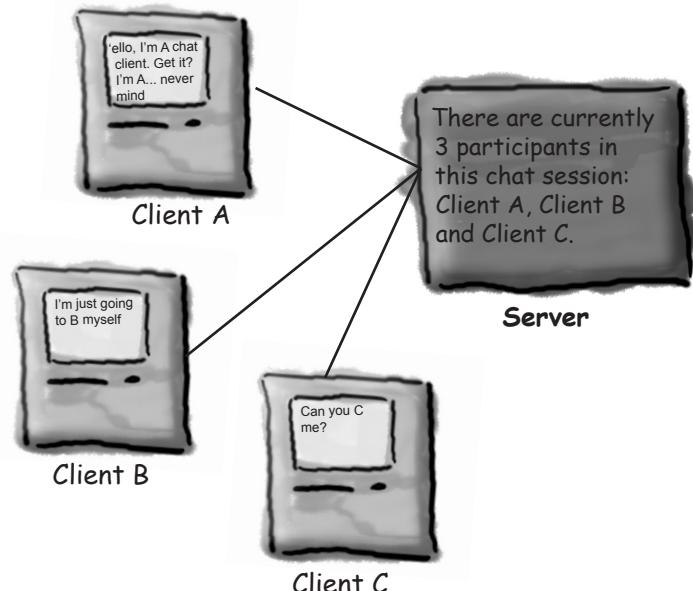
You can have completely authentic, intellectually stimulating chat conversations. Every message is sent to all participants.

Send your message to the server.

Chat program overview

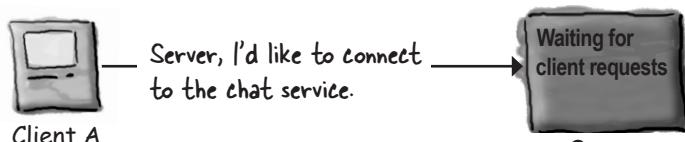
Each Client has to know about the Server.

The Server has to know about ALL the Clients.

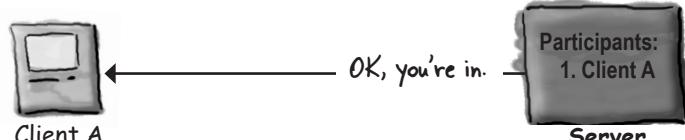


How it works:

- 1 Client connects to the server



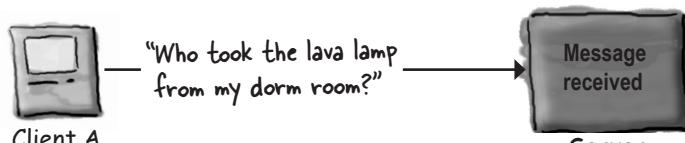
- 2 The server makes a connection and adds the client to the list of participants



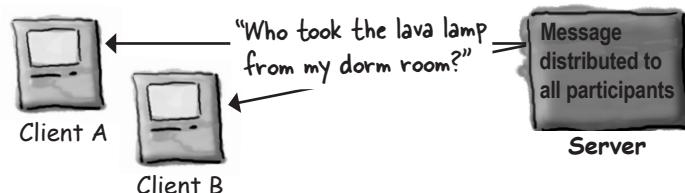
- 3 Another client connects



- 4 Client A sends a message to the chat service



- 5 The server distributes the message to ALL participants (including the original sender)



Connecting, sending, and receiving

The three things we have to learn to get the client working are:

1. How to establish the initial **connection** between the client and server
2. How to **receive** messages *from* the server
3. How to **send** messages *to* the server

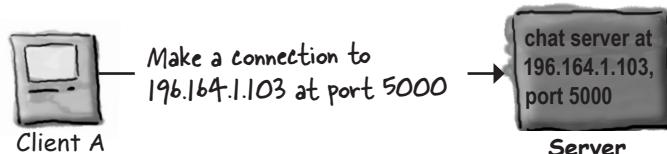
There's a lot of low-level stuff that has to happen for these things to work. But we're lucky, because the Java APIs make it a piece of cake for programmers. You'll see a lot more GUI code than networking and I/O code in this chapter.

And that's not all.

Lurking within the simple chat client is a problem we haven't faced so far in this book: doing two things at the same time. Establishing a connection is a one-time operation (that either works or fails). But after that, a chat participant wants to *send outgoing messages* and **simultaneously** receive incoming messages from the other participants (via the server). Hmm...that one's going to take a little thought, but we'll get there in just a few pages.

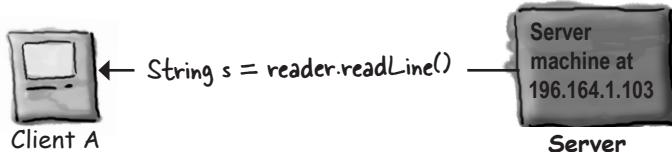
➊ Connect

Client **connects** to the server



➋ Receive

Client **reads** a message from the server



➌ Send

Client **writes** a message to the server



1. Connect

To talk to another machine, we need an object that represents a network connection between two machines. We can open a java.nio.channels.SocketChannel to give us this connection object.

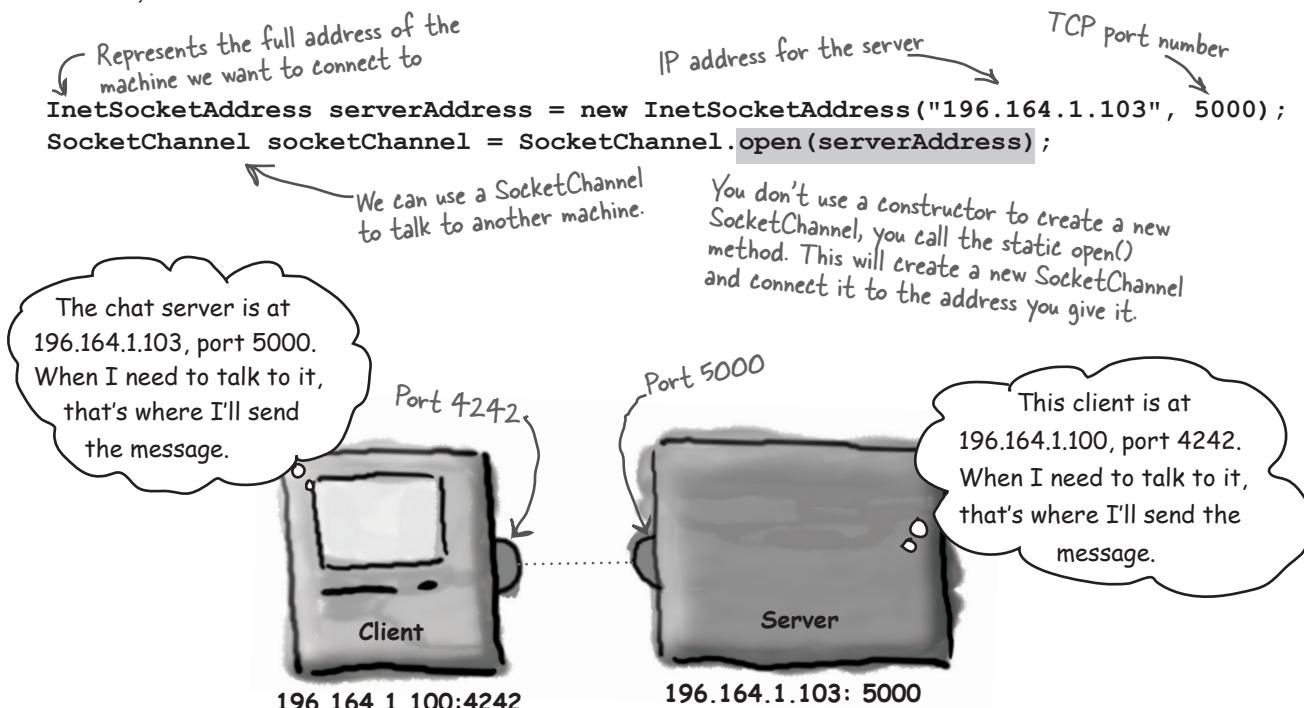
What's a connection? A *relationship* between two machines, where **two pieces of software know about each other**. Most importantly, those two pieces of software know how to *communicate* with each other. In other words, how to send *bits* to each other.

We don't care about the low-level details, thankfully, because they're handled at a much lower place in the “networking stack.” If you don't know what the “networking stack” is, don't worry about it. It's just a way of looking at the layers that information (bits) must travel through to get from a Java program running in a JVM on some OS, to physical hardware (Ethernet cables, for example), and back again on some other machine.

The part that you have to worry about is high-level. You just have to create an object for the server's address and then open a channel to that server. Ready?

To make a connection,
you need to know two
things about the server:
where it is and which
port it's running on.

In other words,
**IP address and TCP
port number**.



A connection means the two machines have information about each other, including network location (IP address) and TCP port.

A TCP port is just a number... a 16-bit number that identifies a specific program on the server

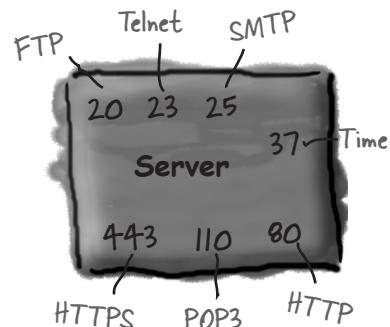
Your internet web (HTTP) server runs on port 80. That's a standard. If you've got a Telnet server, it's running on port 23. FTP? 20. POP3 mail server? 110. SMTP? 25. The Time server sits at 37. Think of port numbers as unique identifiers. They represent a logical connection to a particular piece of software running on the server. That's it. You can't spin your hardware box around and find a TCP port. For one thing, you have 65,536 of them on a server (0–65535). So they obviously don't represent a place to plug in physical devices. They're just a number representing an application.

Without port numbers, the server would have no way of knowing which application a client wanted to connect to. And since each application might have its own unique protocol, think of the trouble you'd have without these identifiers. What if your web browser, for example, landed at the POP3 mail server instead of the HTTP server? The mail server won't know how to parse an HTTP request! And even if it did, the POP3 server doesn't know anything about servicing the HTTP request.

When you write a server program, you'll include code that tells the program which port number you want it to run on (you'll see how to do this in Java a little later in this chapter). In the Chat program we're writing in this chapter, we picked 5000. Just because we wanted to. And because it met the criteria that it be a number between 1024 and 65535. Why 1024? Because 0 through 1023 are reserved for the well-known services like the ones we just talked about.

And if you're writing services (server programs) to run on a company network, you should check with the sysadmins to find out which ports are already taken. Your sysadmins might tell you, for example, that you can't use any port number below, say, 3000. In any case, if you value your limbs, you won't assign port numbers with abandon. Unless it's your *home* network. In which case you just have to check with your *kids*.

Well-known TCP port numbers
for common server applications:



A server can have up to 65,536 different server apps running, one per port.

The TCP port numbers from 0 to 1023 are reserved for well-known services. Don't use them for your own server programs!*

The chat server we're writing uses port 5000. We just picked a number between 1024 and 65535.

*Well, you *might* be able to use one of these, but the sysadmin where you work will write you a strongly worded message and CC your boss.

there are no Dumb Questions

Q: How do you know the port number of the server program you want to talk to?

A: That depends on whether the program is one of the well-known services. If you're trying to connect to a well-known service, like the ones on the opposite page (HTTP, SMTP, FTP, etc.), you can look these up on the internet (Google "Well-Known TCP Port"). Or ask your friendly neighborhood sysadmin.

But if the program isn't one of the well-known services, you need to find out from whoever is deploying the service. Ask them. Typically, if someone writes a network service and wants others to write clients for it, they'll publish the IP address, port number, and protocol for the service. For example, if you want to write a client for a GO game server, you can visit one of the GO server sites and find information about how to write a client for that particular server.

Q: Can there ever be more than one program running on a single port? In other words, can two applications on the same server have the same port number?

A: No! If you try to bind a program to a port that is already in use, you'll get a `BindException`. To *bind* a program to a port just means starting up a server application and telling it to run on a particular port. Again, you'll learn more about this when we get to the server part of this chapter.

IP address is the mall



Port number is the specific store in the mall



IP address is like specifying a particular shopping mall, say, "Flatirons Marketplace"

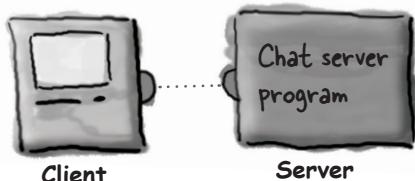
Port number is like naming a specific store, say, "Bob's CD Shop"



Brain Barbell

OK, you got a connection. The client and the server know the IP address and TCP port number for each other. Now what? How do you communicate over that connection? In other words, how do you move bits from one to the other? Imagine the kinds of messages your chat client needs to send and receive.

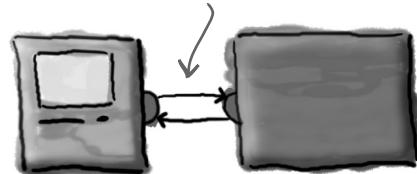
How do these two actually talk to each other?



2. Receive

To communicate over a remote connection, you can use regular old I/O streams, just like we used in the previous chapter. One of the coolest features in Java is that most of your I/O work won't care what your high-level chain stream is actually connected to. In other words, you can use a **BufferedReader** just like you did when you were reading from a file; the difference is that the underlying connection stream is connected to a *Channel* rather than a *File*!

Channels between the client and server



Reading from the network with BufferedReader

1 Make a connection to the server

```
SocketAddress serverAddr = new InetSocketAddress("127.0.0.1", 5000);
```

```
SocketChannel socketChannel = SocketChannel.open(serverAddr);
```

You need to open a SocketChannel
that connects to this address.

127.0.0.1 is the IP address for "localhost," in other words, the one this code is running on. You can use this when you're testing your client and server on a single, stand-alone machine. You could also use "localhost" here instead.

The port number, which you know because we TOLD you that 5000 is the port number for our chat server.

2 Create or get a Reader from the connection

```
Reader reader = Channels.newReader(socketChannel, StandardCharsets.UTF_8);
```

This Reader is a "bridge" between a low-level byte stream (like the one coming from the Channel) and a high-level character stream (like the BufferedReader we're after as our top of the chain stream).

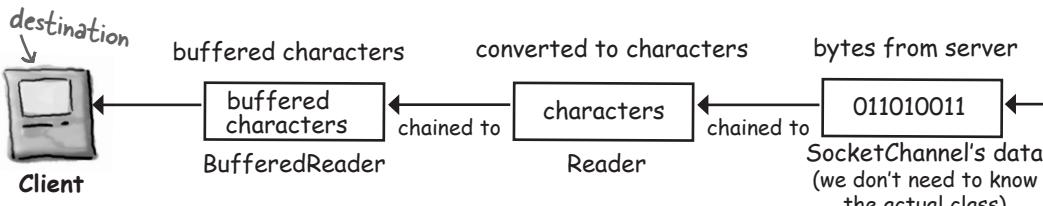
You can use the static helper methods on the Channels class to create a Reader from your SocketChannel.

You need to say which Charset to use for reading the values from the network. UTF 8 is a common one to use.

Chain the BufferedReader to the Reader (which is from our SocketChannel).

3 Make a BufferedReader and read!

```
BufferedReader bufferedReader = new BufferedReader(reader);
String message = bufferedReader.readLine();
```



3. Send

In the previous chapter, we used BufferedWriter. We have a choice here, but when you're writing one String at a time, **PrintWriter** is a standard choice. And you'll recognize the two key methods in PrintWriter, print() and println(). Just like good ol' System.out.

Writing to the network with PrintWriter

1 Make a connection to the server

```
SocketAddress serverAddr = new InetSocketAddress("127.0.0.1", 5000);
SocketChannel socketChannel = SocketChannel.open(serverAddr);
```

This part's the same as it was on the last page—to write to the server, we still have to connect to it.

2 Create or get a Writer from the connection

```
Writer writer = Channels.newWriter(socketChannel, StandardCharsets.UTF_8);
```

Writer acts as a bridge between character data and the bytes to be written to the Channel.

The Channels class contains utility methods to create a Writer.

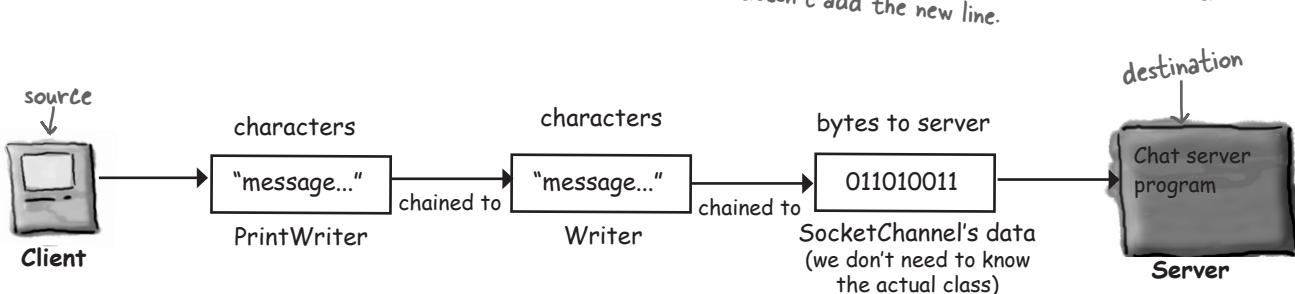
You need to say which Charset to use to write Strings. You should use the same one for reading as for writing!

3 Make a PrintWriter and write (print) something

```
PrintWriter printWriter = new PrintWriter(writer);
```

```
writer.println("message to send"); ← println() adds a new line at the end of what it sends.
writer.print("another message"); ← print() doesn't add the new line.
```

By chaining a PrintWriter to the Channel's Writer, we can write Strings to the Channel, which will be sent over the connection.



There's more than one way to make a connection

If you look at real life code that talks to a remote machine, you'll probably see a number of different ways to make connections and to read from and write to a remote computer.

Which approach you use depends on a number of things, including (but not limited to) the version of Java you're using and the needs of the application (for example, how many clients connect at once, the size of messages sent, frequency or message, etc). One of the simplest approaches is to use a **java.net.Socket** instead of a Channel.

Using a Socket

You can get an *InputStream* or *OutputStream* from a Socket, and read and write from it in a very similar way to what we've already seen.



```
Instead of using an InetSocketAddress and
opening a SocketChannel, you can create a
Socket with the host and port number.
```

```
Socket chatSocket = new Socket("127.0.0.1", 5000);
```

To read from the Socket, we need to get an *InputStream* from the Socket.

```
InputStreamReader in = new InputStreamReader(chatSocket.getInputStream());
```

Reader code is exactly the same as we've already seen.

```
BufferedReader reader = new BufferedReader(in);}
```

```
String message = reader.readLine();
```

```
PrintWriter writer = new PrintWriter(chatSocket.getOutputStream());
```

To write to the socket, we need to get an *OutputStream* from the Socket, which we can chain to the *PrintWriter*.

```
writer.println("message to send");}
writer.print("another message");}
```

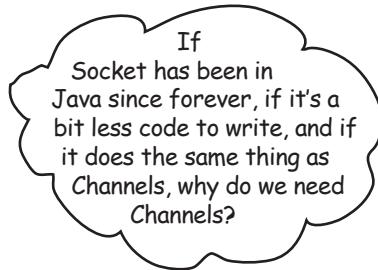
Writer code is exactly the same as we've already seen.

The **java.net.Socket** class is available in all versions of Java.

It supports simple network I/O via the I/O streams we've already used for file I/O.



o o



**As we've become an increasingly connected world,
Java has evolved to offer more ways to work with
remote machines.**

Remember that Channels are in the `java.nio.channels` package? The `java.nio` package (NIO) was introduced in Java 1.4, and there were more changes and additions made (sometimes called NIO.2) in Java 7.

There are ways to use Channels and NIO to get better performance when you're working with lots of network connections, or lots of data coming over those connections.

In this chapter, we're using Channels to provide the same very basic connection functionality we could get from

Sockets. However, if our application needed to work well with a very busy network connection (or lots of them!), we could configure our Channels differently and use them to their full potential, and our program would cope better with a high network I/O load.

We've chosen to teach you the **simplest way to get started** with network I/O using *Channels* so that if you need to "level up" to working with more advanced features, it shouldn't be such a big step.

If you do want to learn more about NIO, read *Java NIO* by Ron Hitchens and *Java I/O, NIO and NIO.2* by Jeff Friesen.



Channels support advanced networking features that you don't need for these exercises.

Channels can support nonblocking I/O, reading and writing via ByteBuffers, and asynchronous I/O. We're not going to show you any of this!

But at least now you have some keywords to put into your search engine when you want to know more.

The DailyAdviceClient

Before we start building the Chat app, let's start with something a little smaller. The Advice Guy is a server program that offers up practical, inspirational tips to get you through those long days of coding.

We're building a client for The Advice Guy program, which pulls a message from the server each time it connects.

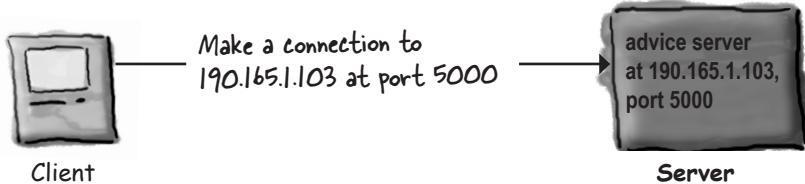
What are you waiting for? Who *knows* what opportunities you've missed without this app.



The Advice Guy

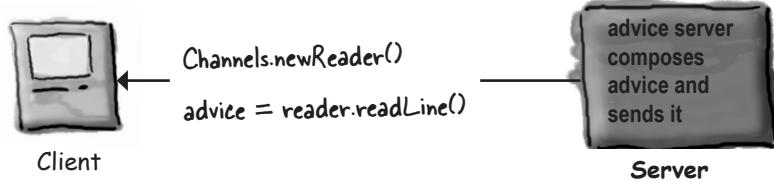
➊ Connect

Client connects to the server



➋ Read

Client gets a Reader for the Channel, and reads a message from the server



DailyAdviceClient code

This program makes a SocketChannel, makes a BufferedReader (with the help of the channel's Reader), and reads a single line from the server application (whatever is running at port 5000).

```

import java.io.*;
import java.net.InetSocketAddress;
import java.nio.channels.Channels;
import java.nio.channels.SocketChannel;
import java.nio.charset.StandardCharsets;

public class DailyAdviceClient {
    public void go() {
        InetSocketAddress serverAddress = new InetSocketAddress("127.0.0.1", 5000);
        try (SocketChannel socketChannel = SocketChannel.open(serverAddress)) {
            Reader channelReader = Channels.newReader(socketChannel, StandardCharsets.UTF_8);
            BufferedReader reader = new BufferedReader(channelReader);
            String advice = reader.readLine();
            System.out.println("Today you should: " + advice);
        } catch (IOException e) {
            e.printStackTrace();
        }
    }

    public static void main(String[] args) {
        new DailyAdviceClient().go();
    }
}

```

This uses try-with-resources to automatically close the SocketChannel when the code is complete.

Define the server address as being port 5000, on the same host this code is running on (the "localhost").

Create a SocketChannel by opening one for the server's address.

Create a Reader that reads from the SocketChannel.

Chain a BufferedReader to the Reader from the SocketChannel.

This readLine() is EXACTLY the same as if you were using a BufferedReader chained to a FILE.. In other words, by the time you call a BufferedReader method, the reader doesn't know or care where the characters came from.

exercise: sharpen your pencil



Sharpen your pencil

Test your memory of the classes for reading and writing from a SocketChannel. Try not to look at the opposite page!

→ Yours to solve.

To **read** text from a SocketChannel:



Client

Write/draw in the chain of classes the client uses to read from the server.



Server

To **send** text to a SocketChannel:



Client

Write/draw in the chain of classes the client uses to send something to the server.



Server



Sharpen your pencil

Fill in the blanks:

→ Yours to solve.

What two pieces of information does the client need in order to make a connection with a server?

Which TCP port numbers are reserved for “well-known services” like HTTP and FTP? _____

TRUE or FALSE: The range of valid TCP port numbers can be represented by a short primitive.

Writing a simple server application

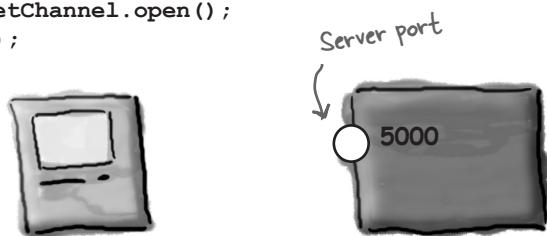
So what's it take to write a server application? Just a couple of Channels. Yes, a couple as in two. A ServerSocketChannel, which waits for client requests (when a client connects) and a SocketChannel to use for communication with the client. If there's more than one client, we'll need more than one channel, but we'll get to that later.

How it works:

- 1 Server application makes a ServerSocketChannel and binds it to a specific port

```
ServerSocketChannel serverChannel = ServerSocketChannel.open();
serverChannel.bind(new InetSocketAddress(5000));
```

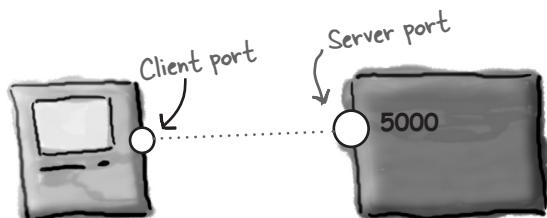
This starts the server application listening for client requests coming in for port 5000.



- 2 Client makes a SocketChannel connected to the server application

```
SocketChannel svr = SocketChannel.open(new InetSocketAddress("190.165.1.103", 5000));
```

Client knows the IP address and port number (published or given to them by whomever configures the server app to be on that port).

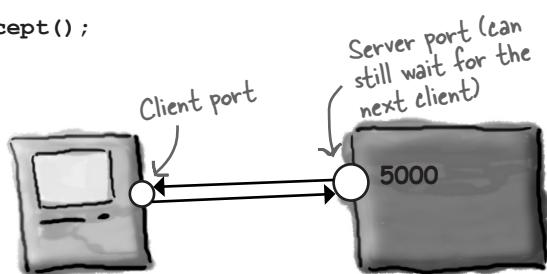


- 3 Server makes a new SocketChannel to communicate with this client

```
SocketChannel clientChannel = serverChannel.accept();
```

The accept() method blocks (just sits there) while it's waiting for a client connection. When a client finally connects, the method returns a SocketChannel that knows how to communicate with this client.

The ServerSocketChannel can go back to waiting for other clients. The server has just one ServerSocketChannel, and a SocketChannel per client.



DailyAdviceServer code

This program makes a ServerSocketChannel and waits for client requests. When it gets a client request (i.e., client created a new SocketChannel to this server), the server makes a new SocketChannel to that client. The server makes a PrintWriter (using a Writer created from the SocketChannel) and sends a message to the client.

```

import java.io.*;
import java.net.InetSocketAddress;
import java.nio.channels.*;
import java.util.Random;

public class DailyAdviceServer {
    final private String[] adviceList = {
        "Take smaller bites",           Remember the imports.
        "Go for the tight jeans. No they do NOT make you look fat.",
        "One word: inappropriate",
        "Just for today, be honest. Tell your boss what you *really* think",
        "You might want to rethink that haircut."};

    private final Random random = new Random();

    public void go() {
        try (ServerSocketChannel serverChannel = ServerSocketChannel.open()) {
            serverChannel.bind(new InetSocketAddress(5000));      Daily advice comes from this array.

            while (serverChannel.isOpen()) {                      ServerSocketChannel makes this server
                SocketChannel clientChannel = serverChannel.accept();  application "listen" for client requests on the
                PrintWriter writer = new PrintWriter(Channels.newOutputStream(clientChannel));   port it's bound to.

                String advice = getAdvice();
                writer.println(advice);                                You have to bind the ServerSocketChannel to
                writer.close();                                         the port you want to run the application on.

                System.out.println(advice);                           The accept method blocks (just sits there) until a
                }                                                       request comes in, and then the method returns a
            }                                                       SocketChannel for communicating with the client.

            catch (IOException ex) {                                Create an output stream for the client's
                ex.printStackTrace();                               channel, and wrap it in a PrintWriter. You can
            }                                                       use newOutputStream or newWriter here.

            Print in the server console, so
            we can see what's happening.                         Send the client a String advice message.

        }
    }

    private String getAdvice() {
        int nextAdvice = random.nextInt(adviceList.length);
        return adviceList[nextAdvice];
    }

    public static void main(String[] args) {
        new DailyAdviceServer().go();
    }
}

```

The server goes into a permanent loop, waiting for (and servicing) client requests.

The accept method blocks (just sits there) until a request comes in, and then the method returns a SocketChannel for communicating with the client.

Create an output stream for the client's channel, and wrap it in a PrintWriter. You can use newOutputStream or newWriter here.

Close the writer, which will also close the client SocketChannel.

How does the server know how to communicate with the client?

Think about how/when/where the server gets knowledge about the client.



Brain Barbell



Yes, that's right, **the server can't accept a request from a client until it has finished with the current client**. At which point, it starts the next iteration of the infinite loop, sitting, waiting, at the `accept()` call until a new request comes in, at which point it makes a `SocketChannel` to send data to the new client and starts the process over again.

To get this to work with multiple clients *at the same time*, we need to use separate threads.

We'd give each new client's `SocketChannel` to a new thread, and each thread can work independently.

We're just about to learn how to do that!

BULLET POINTS

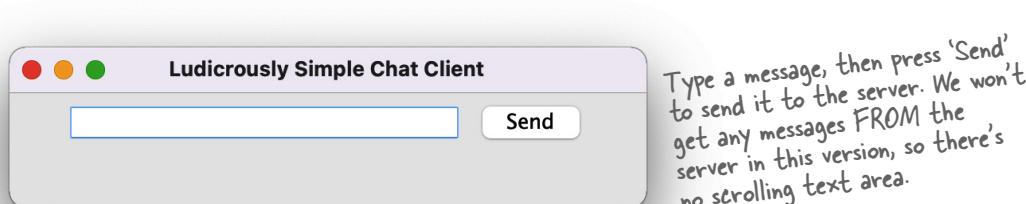
- Client and server applications communicate using Channels.
- A Channel represents a connection between two applications that may (or may not) be running on two different physical machines.
- A client must know the IP address (or host name) and TCP port number of the server application.
- A TCP port is a 16-bit unsigned number assigned to a specific server application. TCP port numbers allow different server applications to run on the same machine; clients connect to a specific application using its port number.
- The port numbers from 0 through 1023 are reserved for "well-known services" including HTTP, FTP, SMTP, etc.
- A client connects to a server by opening a `SocketChannel`:
`SocketChannel.open(
 new InetSocketAddress("127.0.0.1", 4200))`
- Once connected, a client can create readers (to read data from the server) and writers (to send data to the server) for the channel:
`Reader reader = Channels.newReader(sockCh,
 StandardCharsets.UTF_8);`
`Writer writer = Channels.newWriter(sockCh,
 StandardCharsets.UTF_8);`
- To read text data from the server, create a `BufferedReader`, chained to the Reader. The Reader is a "bridge" that takes in bytes and converts them to text (character) data. It's used primarily to act as the middle chain between the high-level `BufferedReader` and the low-level connection.
- To write text data to the server, create a `PrintWriter` chained to the Writer. Call the `print()` or `println()` methods to send Strings to the server.
- Servers use a `ServerSocketChannel` that waits for client requests on a particular port number.
- When a `ServerSocketChannel` gets a request, it "accepts" the request by making a `SocketChannel` for the client.

Writing a Chat Client

We'll write the Chat Client application in two stages. First we'll make a send-only version that sends messages to the server but doesn't get to read any of the messages from other participants (an exciting and mysterious twist to the whole chat room concept).

Then we'll go for the full chat monty and make one that both sends *and* receives chat messages.

Version One: send-only



Code outline

Here's an outline of the main functionality the chat client needs to provide. The full code is on the next page.

```
public class SimpleChatClientA {
    private JTextField outgoing;
    private PrintWriter writer;

    public void go() {
        // call the setUpNetworking() method
        // make gui and register a listener with the send button
    }

    private void setUpNetworking() {
        // open a SocketChannel to the server
        // make a PrintWriter and assign to writer instance variable
    }

    private void sendMessage() {
        // get the text from the text field and
        // send it to the server using the writer (a PrintWriter)
    }
}
```

```

import javax.swing.*;
import java.awt.*;
import java.io.*;
import java.net.InetSocketAddress;
import java.nio.channels.*;
import static java.nio.charset.StandardCharsets.UTF_8; This is a static import; we looked at static imports in Chapter 10.

public class SimpleChatClientA {
    private JTextField outgoing;
    private PrintWriter writer;

    public void go() { Call the method that will
        setUpNetworking(); connect to the server.

        outgoing = new JTextField(20);

        JButton sendButton = new JButton("Send");
        sendButton.addActionListener(e -> sendMessage());
    }

    JPanel mainPanel = new JPanel();
    mainPanel.add(outgoing);
    mainPanel.add(sendButton);
    JFrame frame = new JFrame("Ludicrously Simple Chat Client");
    frame.getContentPane().add(BorderLayout.CENTER, mainPanel);
    frame.setSize(400, 100);
    frame.setVisible(true);
    frame.setDefaultCloseOperation(WindowConstants.EXIT_ON_CLOSE);
}

private void setUpNetworking() { We're using localhost so you
    try { can test the client and
        InetSocketAddress serverAddress = new InetSocketAddress("127.0.0.1", 5000); server on one machine.

        SocketChannel socketChannel = SocketChannel.open(serverAddress); Open a SocketChannel
        writer = new PrintWriter(Channels.newWriter(socketChannel, UTF_8)); that connects to the
        System.out.println("Networking established."); server.

    } catch (IOException e) {
        e.printStackTrace();
    }
}

private void sendMessage() { Now we actually do the writing. Remember,
    writer.println(outgoing.getText()); } the writer is chained to the writer from the
    writer.flush(); SocketChannel, so whenever we do a println(),
    outgoing.setText(""); it goes over the network to the server!
    outgoing.requestFocus();

public static void main(String[] args) {
    new SimpleChatClientA().go();
}
}

```

Imports for writing (java.io), network connections (java.nio.channels) and the GUI stuff (awt and swing)

Build the GUI, nothing new here, and nothing related to networking or I/O.

This is where we make the PrintWriter from a writer that writes to the SocketChannel.

If you want to try this now, type in the Ready-bake chat server code listed on the next page.

First, start the server in one terminal. Next, use another terminal to start this client.



Ready-Bake Code

The really, really simple Chat Server

You can use this server code for all versions of the Chat Client. Every possible disclaimer ever disclaimed is in effect here. To keep the code stripped down to the bare essentials, we took out a lot of parts that you'd need to make this a real server. In other words, it works, but there are at least a hundred ways to break it. If you want to really sharpen your skills after you've finished this book, come back and make this server code more robust.

After you finish this chapter, you should be able to annotate this code yourself. You'll understand it much better if *you* work out what's happening than if we explained it to you. Then again, this is Ready-Bake Code, so you really don't have to understand it at all. It's here just to support the two versions of the Chat Client.

To run the Chat Client, you need two terminals. First, launch this server from one terminal, and then launch the client from another terminal.

```

import java.io.*;
import java.net.InetSocketAddress;
import java.nio.channels.*;
import java.util.*;
import java.util.concurrent.*;

import static java.nio.charset.StandardCharsets.UTF_8;

public class SimpleChatServer {
    private final List<PrintWriter> clientWriters = new ArrayList<>();

    public static void main(String[] args) {
        new SimpleChatServer().go();
    }

    public void go() {
        ExecutorService threadPool = Executors.newCachedThreadPool();
        try {
            ServerSocketChannel serverSocketChannel = ServerSocketChannel.open();
            serverSocketChannel.bind(new InetSocketAddress(5000));

            while (serverSocketChannel.isOpen()) {
                SocketChannel clientSocket = serverSocketChannel.accept();
                PrintWriter writer = new PrintWriter(Channels.newWriter(clientSocket, UTF_8));
                clientWriters.add(writer);
                threadPool.submit(new ClientHandler(clientSocket));
                System.out.println("got a connection");
            }
        } catch (IOException ex) {
            ex.printStackTrace();
        }
    }
}

```

```

private void tellEveryone(String message) {
    for (PrintWriter writer : clientWriters) {
        writer.println(message);
        writer.flush();
    }
}

public class ClientHandler implements Runnable {
    BufferedReader reader;
    SocketChannel socket;

    public ClientHandler(SocketChannel clientSocket) {
        socket = clientSocket;
        reader = new BufferedReader(Channels.newReader(socket, UTF_8));
    }

    public void run() {
        String message;
        try {
            while ((message = reader.readLine()) != null) {
                System.out.println("read " + message);
                tellEveryone(message);
            }
        } catch (IOException ex) {
            ex.printStackTrace();
        }
    }
}
}

```

```

File Edit Window Help TakesTwoToTango
%java SimpleChatServer
got a connection
read Nice to meet you

```

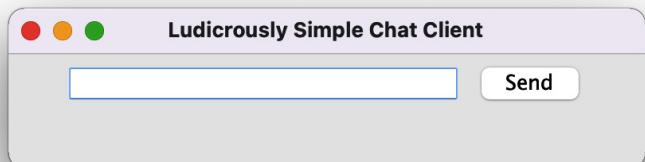
Runs in the background

```

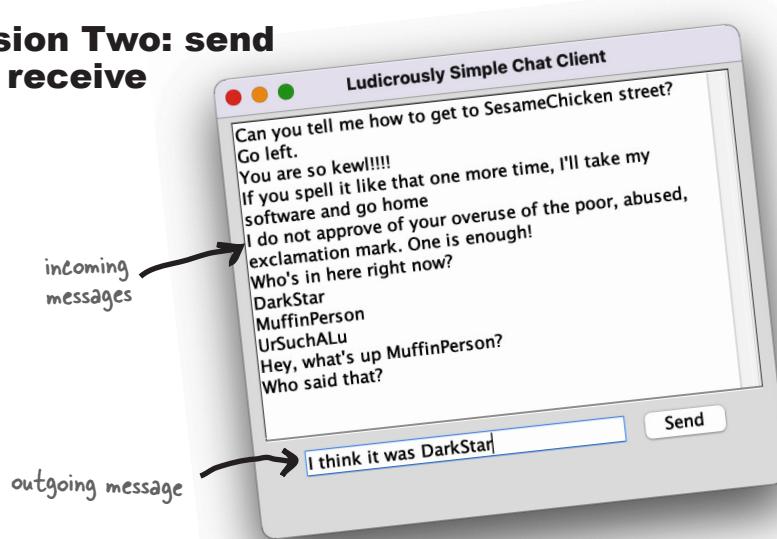
File Edit Window Help MayIHaveThisDance?
%java SimpleChatClientA
Networking established. Client
running at: /127.0.0.1:57531

```

Connects to the server and launches GUI



Version Two: send and receive



The Server sends a message to all client participants, as soon as the message is received by the server. When a client sends a message, it doesn't appear in the incoming message display area until the server sends it to everyone.

Big Question: HOW do you get messages from the server?

Should be easy; when you set up the networking, make a Reader as well. Then read messages using `readLine`.

Bigger Question: WHEN do you get messages from the server?

Think about that. What are the options?

① Option One: Read something in from the server each time the user sends a message.

Pros: Doable, very easy.

Cons: Stupid. Why choose such an arbitrary time to check for messages? What if a user is a lurker and doesn't send anything?

② Option Two: Poll the server every 20 seconds.

Pros: It's doable, and it fixes the lurker problem.

Cons: How does the server know what you've seen and what you haven't? The server would have to store the messages, rather than just doing a distribute-and-forget each time it gets one. And why 20 seconds? A delay like this affects usability, but as you reduce the delay, you risk hitting your server needlessly. Inefficient.

③ Option Three: Read messages as soon as they're sent from the server.

Pros: Most efficient, best usability.

Cons: How do you do two things at the same time? Where would you put this code? You'd need a loop somewhere that was always waiting to read from the server. But where would that go? Once you launch the GUI, nothing happens until an event is fired by a GUI component.



In Java you really CAN walk and chew gum at the same time.

You know by now that we're going with option three

We want something to run continuously, checking for messages from the server, but *without interrupting the user's ability to interact with the GUI!* So while the user is happily typing new messages or scrolling through the incoming messages, we want something *behind the scenes* to keep reading in new input from the server.

That means we finally need a new thread. A new, separate stack.

We want everything we did in the Send-Only version (version one) to work the same way, while a new *process* runs alongside that reads information from the server and displays it in the incoming text area.

Well, not quite. Each new Java thread is not actually a separate process running on the OS. But it almost *feels* as though it is.

We're going to take a break from the chat application for a bit while we explore how this works. Then we'll come back and add it to our chat client at the end of the chapter.

Multithreading in Java

Java has support for multiple threads built right into the fabric of the language. And it's a snap to make a new thread of execution:

```
Thread t = new Thread();
t.start();
```

That's it. By creating a new *Thread object*, you've launched a separate *thread of execution*, with its very own call stack.

Except for one problem.

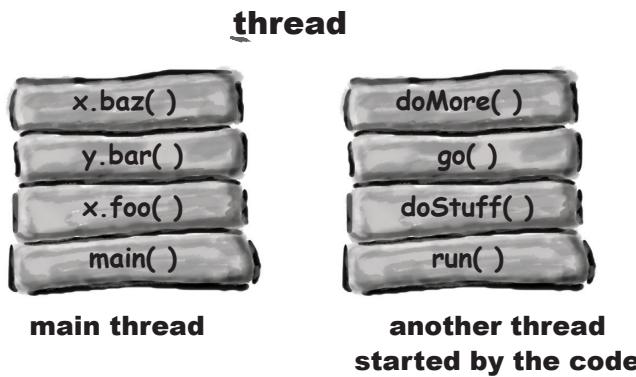
That thread doesn't actually *do* anything, so the thread "dies" virtually the instant it's born. When a thread dies, its new stack disappears again. End of story.

So we're missing one key component—the thread's *job*. In other words, we need the code that you want to have run by a separate thread.

Multiple threading in Java means we have to look at both the *thread* and the *job* that's *run* by the thread. In fact, **there's more than one way to run multiple jobs in Java**, not just with the Thread *class* in the java.lang package. (Remember, java.lang is the package you get imported for free, implicitly, and it's where the classes most fundamental to the language live, including String and System.)

Java has multiple threads but only one Thread class

We can talk about *thread* with a lowercase “t” and **Thread** with a capital “T.” When you see *thread*, we’re talking about a separate thread of execution. In other words, a separate call stack. When you see **Thread**, think of the Java naming convention. What, in Java, starts with a capital letter? Classes and interfaces. In this case, **Thread** is a class in the `java.lang` package. A **Thread** object represents a *thread of execution*. In older versions of Java, you always had to create an instance of class **Thread** each time you wanted to start up a new *thread* of execution. Java has evolved over time, and now using the **Thread** class directly is not the only way. We’ll see this in more detail as we go through the rest of the chapter.



A *thread* (lowercase “t”) is a separate thread of execution. That means a separate call stack. Every Java application starts up a main thread—the thread that puts the `main()` method on the bottom of the stack. The JVM is responsible for starting the main thread (and other threads, as it chooses, including the garbage collection thread). As a programmer, you can write code to start other threads of your own.

A **thread** is a separate “thread of execution,” a separate call stack.

A **Thread** is a Java class that represents a thread.

Using the **Thread** class is not the only way to do multithreading in Java.

Thread

Thread
<code>void join()</code>
<code>void start()</code>

`static void sleep()`

`java.lang.Thread` class

Thread (capital “T”) is a class that represents a thread of execution. It has methods for starting a thread, joining one thread with another, putting a thread to sleep, and more.

What does it mean to have more than one call stack?

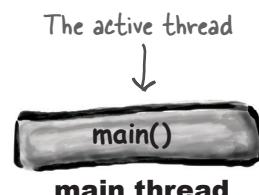
With more than one call stack, you can have multiple things happen at the same time. If you're running on a multiprocessor system (like most modern computers and phones), you can actually do more than one thing at a time. With Java threads, even if you're not running on a multiprocessor system or if you're running more processes than available cores, it can *appear* that you're doing all these things simultaneously. In other words, execution can move back and forth between stacks so rapidly that you feel as though all stacks are executing at the same time. Remember, Java is just a process running on your underlying OS. So first, Java *itself* has to be "the currently executing process" on the OS. But once Java gets its turn to execute, exactly *what* does the JVM *run*? Which bytecodes execute? Whatever is on the top of the currently running stack! And in 100 milliseconds, the currently executing code might switch to a *different* method on a *different* stack.

One of the things a thread must do is keep track of which statement (of which method) is currently executing on the thread's stack.

It might look something like this:

- 1** The JVM calls the main() method.

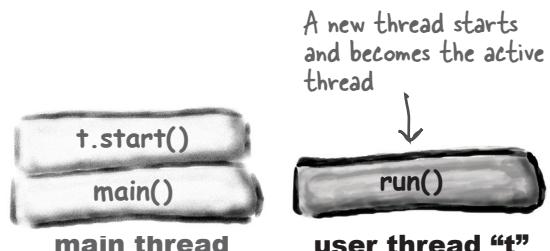
```
public static void main(String[] args) {  
    ...  
}
```



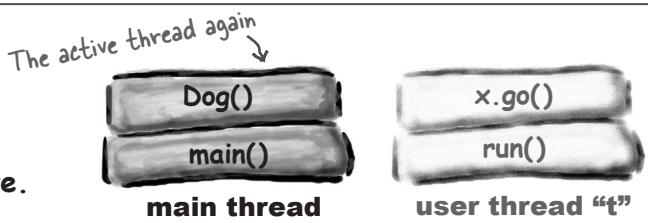
- 2** main() starts a new thread. The main thread may be temporarily frozen while the new thread starts running.

```
Runnable r = new MyThreadJob();  
Thread t = new Thread(r);  
t.start();  
Dog d = new Dog();
```

You'll learn what this means in just a moment...



- 3** The JVM switches between the new thread (user thread A) and the original main thread, until both threads complete.



To create a new call stack you need a job to run



Runnable is to a thread what a job is to a worker. A Runnable is the job a thread is supposed to run.

A Runnable holds the method that goes on the bottom of the new call stack: `run()`.

To start a new call stack the thread needs a job—a job the thread will run when it's started. That job is actually the first method that goes on the new thread's stack, and it must always be a method that looks like this:

```
public void run() {  
    // code that will be run by the new thread  
}
```

How does the thread know which method to put at the bottom of the stack? Because Runnable defines a contract. Because Runnable is an interface. A thread's job can be defined in any class that implements the Runnable interface, or a lambda expression that is the right shape for the `run` method.

Once you have a Runnable class or lambda expression, you can tell the JVM to run this code in a separate thread; you're giving the thread its job.

The Runnable interface defines only one method, `public void run()`. Since it has only a single method, it's a SAM type, a Functional Interface, and you can use a lambda instead of creating a whole class that implements Runnable if you want.

To make a job for your thread, implement the Runnable interface

Runnable is in the `java.lang` package,
so you don't need to import it.

```
public class MyRunnable implements Runnable {
```

```
    public void run() {
        go();
    }
```

```
    public void go() {
        doMore();
    }
```

```
    public void doMore() {
        System.out.println(Thread.currentThread().getName() +
            ": top o' the stack");
        Thread.dumpStack();
    }
```

We'll see the stack for this thread on the next page. Yes, we've gone straight to "2," you should see why over the page...

Runnable has only one method to implement: `public void run()` (with no arguments). This is where you put the JOB the thread is supposed to run. This is the method that goes at the bottom of the new stack.

`dumpStack` will output the current call stack, just like an Exceptions stack trace. Using it here only use this for debugging (it might slow real code down).

This Runnable doesn't really need three tiny methods, which all call each other like this; we're using it to demonstrate what the call stack running this code looks like.

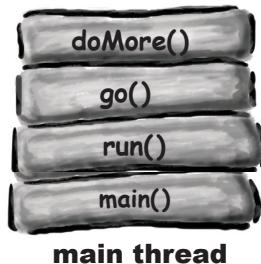
How NOT to run the Runnable

It may be tempting to create a new instance of the Runnable and call the run method, but that's **not enough to create a new call stack**.

```
class RunTester {
    public static void main(String[] args) {
        MyRunnable runnable = new MyRunnable();
        runnable.run();
        System.out.println(Thread.currentThread().getName() +
            ": back in main");
        Thread.dumpStack();
    }
}
```

This will NOT do what we want!

The `run()` method was called directly from inside the `main()` method, so it's part of the call stack of the main thread.



How we used to launch a new thread

The simplest way to launch a new thread is with the Thread class that we mentioned earlier. This method has been around in Java since the very beginning, but **it is no longer the recommended approach to use**. We're showing it here because a) it's simple, and b) you'll see it in the Real World. We will talk later about why it might not be the best approach.

The diagram illustrates the creation of a new thread and its call stack.

Code Snippet:

```

class ThreadTester {
    public static void main(String[] args) {
        Runnable threadJob = new MyRunnable();
        Thread myThread = new Thread(threadJob);
        myThread.start();
        System.out.println(
            Thread.currentThread().getName() +
            ": back in main");
        Thread.dumpStack();
    }
}

```

Annotations:

- ①** A brace groups the line `myThread.start();` and the block containing `System.out.println` and `Thread.dumpStack()`. A callout notes: "This thread's stack is below."
- ②** An annotation points to the line `new Thread(threadJob);` with the text: "Pass the new Runnable instance into the new Thread constructor. This tells the thread what job to run. In other words, the Runnable's run() method will be the first method that the new thread will run."
- A callout for the `myThread.start();` line states: "You won't get a new thread of execution until you call start() on the Thread instance. A thread that's just a Thread instance, like any other object, but it won't have any real 'threadness.'"
- main thread:** A call stack diagram showing the sequence: `main()`.
- new thread:** A call stack diagram showing the sequence: `doMore()`, `go()`, `run()`.
- Call stacks:**
 - Call stack for the "main" thread of the application (started by the "public static void main" method).
 - Call stack for the new thread we started with the `MyRunnable` job.
- Terminal Output:**

```

File Edit Window Help Duo
% java ThreadTester

main: back in main
Thread-0: top o' the stack
java.lang.Exception: Stack trace
    at java.base/java.lang.Thread.dumpStack(Thread.java:1383)
    at ThreadTester.main(MyRunnable.java:38)

java.lang.Exception: Stack trace
    at java.base/java.lang.Thread.dumpStack(Thread.java:1383)
    at MyRunnable.doMore(MyRunnable.java:15)
    at MyRunnable.go(MyRunnable.java:10)
    at MyRunnable.run(MyRunnable.java:6)
    at java.base/java.lang.Thread.run(Thread.java:829)

```
- Annotations on terminal output:**
 - An arrow from the text "dumpStack() called from doMore() in MyRunnable" points to the line `at MyRunnable.doMore(MyRunnable.java:15)`.
 - An arrow from the text "dumpStack() called from main() method" points to the line `at ThreadTester.main(MyRunnable.java:38)`.
 - An arrow from the text "Note the main method is NOT the bottom of the call stack of the Runnable." points to the line `at java.base/java.lang.Thread.run(Thread.java:829)`.

A better alternative: don't manage the Threads at all

Creating and starting a new Thread gives you a lot of control over that Thread, but the downside is you *have* to control it. You have to keep track of all the Threads and make sure they're shut down at the end. Wouldn't it be better to have something else that starts, stops, and even reuses the Threads so you don't have to?

Allow us to introduce an interface in `java.util.concurrent.ExecutorService`. Implementations of this interface will *execute* jobs (Runnables). Behind the scenes the ExecutorService will create, reuse, and kill threads in order to run these jobs.

The `java.util.concurrent.Executors` class has *factory methods* to create the ExecutorService instances we'll need.

Executors have been around since Java 5 and so should be available to you even if you're working with quite an old version of Java. There's no real need to use Thread directly at all these days.

Running one job

For the simple cases we're going to get started with, we'll want to run only one job in addition to our main class. There's a *single thread executor* that we can use to do this.

```
class ExecutorTester {
    public static void main(String[] args) {
        Runnable job = new MyRunnable();
        ExecutorService executor = Executors.newSingleThreadExecutor();
        executor.execute(job);
        System.out.println(Thread.currentThread().getName() +
            ": back in main");
        Thread.dumpStack();
        executor.shutdown();
    }
}
```

Tell the ExecutorService to run the job. It will take care of starting a new thread for the job if it needs to.

Remember to shut down the ExecutorService when you've finished with it. If you don't shut it down, the program will hang around waiting for more jobs.

Instead of creating a Thread instance, use a method on the Executors class to create an ExecutorService.

In our case, we only want to start a single job, so it's logical to create a single thread executor.

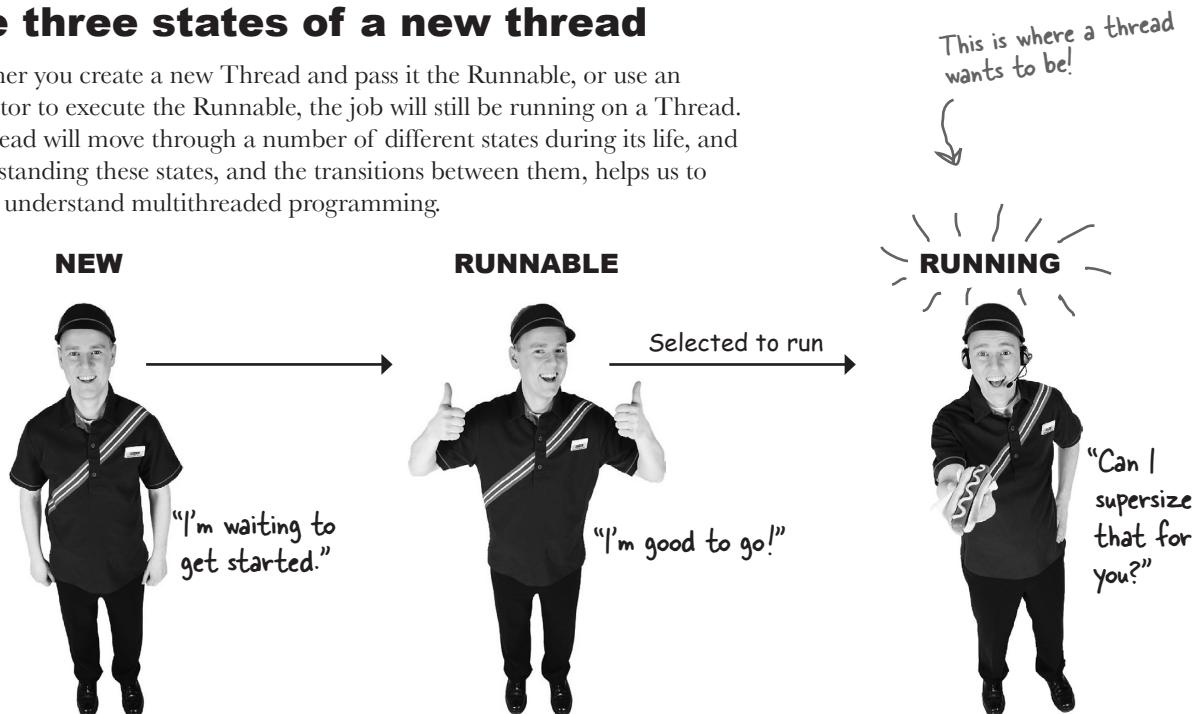
We'll come back to the Executors factory methods later, and we'll see why it might be better to use ExecutorServices rather than managing the Thread itself.

Static factory methods can be used instead of constructors.

Factory methods return exactly the implementation of an interface that we need. We don't need to know the concrete classes or how to create them.

The three states of a new thread

Whether you create a new Thread and pass it the Runnable, or use an Executor to execute the Runnable, the job will still be running on a Thread. A Thread will move through a number of different states during its life, and understanding these states, and the transitions between them, helps us to better understand multithreaded programming.



A Thread instance has been created but not started. In other words, there is a Thread *object*, but no *thread of execution*.

When you start the thread, it moves into the runnable state. This means the thread is ready to run and just waiting for its Big Chance to be selected for execution. At this point, there is a new call stack for this thread.

This is the state all threads lust after! To be Up And Running. Only the JVM thread scheduler can make that decision. You can sometimes *influence* that decision, but you cannot force a thread to move from runnable to running. In the running state, a thread (and ONLY this thread) has an active call stack, and the method on the top of the stack is executing.

But there's more. Once the thread becomes runnable, it can move back and forth between runnable, running, and an additional state: temporarily not runnable.

Typical runnable/running loop

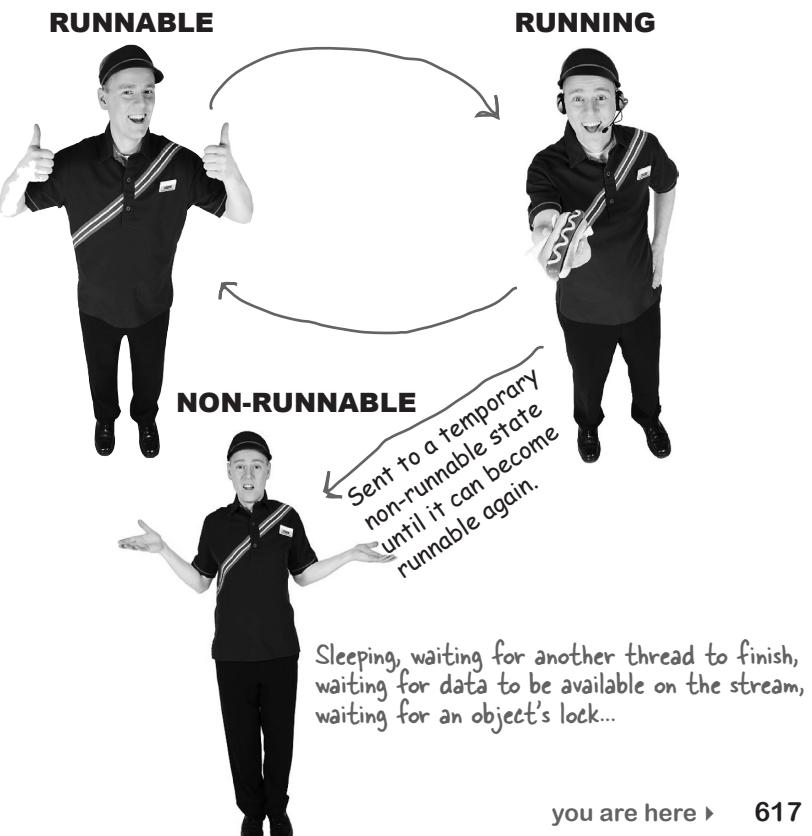
Typically, a thread moves back and forth between runnable and running, as the JVM thread scheduler selects a thread to run and then kicks it back out so another thread gets a chance.



A thread can be made temporarily not-runnable

The thread scheduler can move a running thread into a blocked state, for a variety of reasons. For example, the thread might be executing code to read from an input stream, but there isn't any data to read. The scheduler will move the thread out of the running state until something becomes available. Or the executing code might have told the thread to put itself to sleep (`sleep()`). Or the thread might be waiting because it tried to call a method on an object, and that object was "locked." In that case, the thread can't continue until the object's lock is freed by the thread that has it.

All of those conditions (and more) cause a thread to become temporarily not-runnable.



The thread scheduler

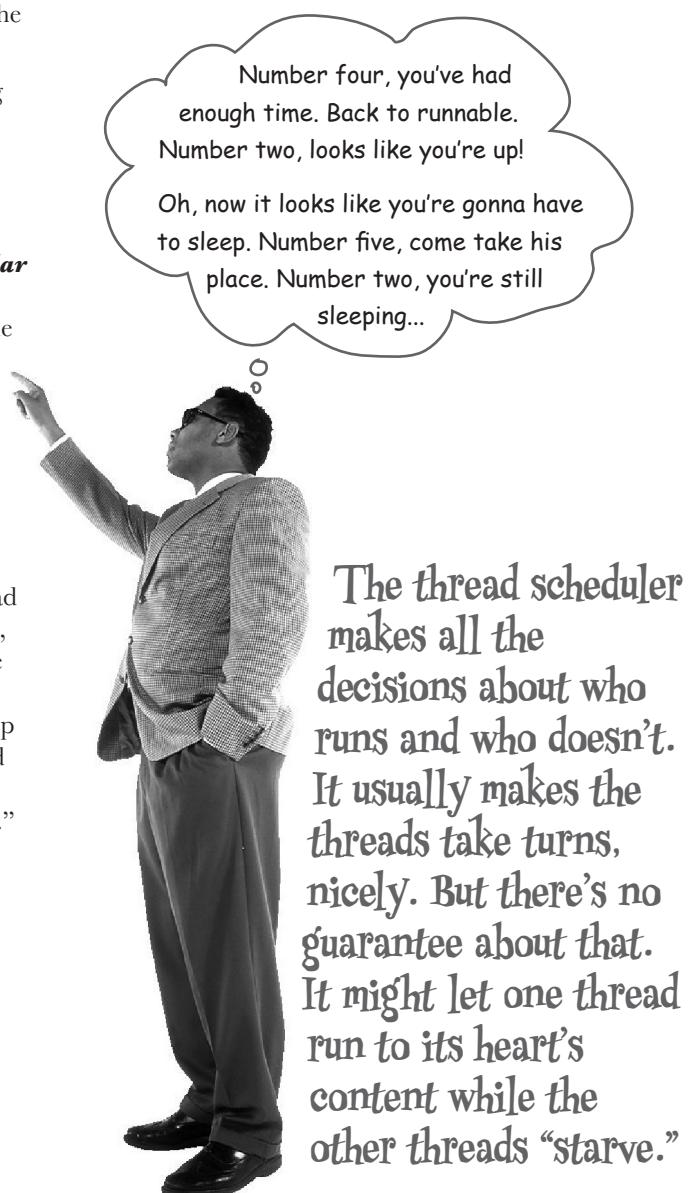
The thread scheduler makes all the decisions about who moves from runnable to running, and about when (and under what circumstances) a thread leaves the running state. The scheduler decides who runs, for how long, and where the threads go when it decides to kick them out of the currently running state.

You can't control the scheduler. There is no API for calling methods on the scheduler. Most importantly, there are no guarantees about scheduling! (There are a few *almost*-guarantees, but even those are a little fuzzy.)

The bottom line is this: ***do not base your program's correctness on the scheduler working in a particular way!*** The scheduler implementations are different for different JVMs, and even running the same program on the same machine can give you different results. One of the worst mistakes new Java programmers make is to test their multithreaded program on a single machine, and assume the thread scheduler will always work that way, regardless of where the program runs.

So what does this mean for write-once-run-anywhere? It means that to write platform-independent Java code, your multithreaded program must work no matter *how* the thread scheduler behaves. That means you can't be dependent on, for example, the scheduler making sure all the threads take nice, perfectly fair, and equal turns at the running state.

Although highly unlikely today, your program might end up running on a JVM with a scheduler that says, “OK, thread five, you're up, and as far as I'm concerned, you can stay here until you're done, when your `run()` method completes.”



An example of how unpredictable the scheduler can be...

Running this code on one machine:

```
class ExecutorTestDrive {
    public static void main (String[] args) {
        ExecutorService executor =
            Executors.newSingleThreadExecutor();

        executor.execute(() ->
            System.out.println("top o' the stack"));

        System.out.println("back in main");
        executor.shutdown();
    }
}
```

Runnable is a Functional Interface and can be represented as a lambda expression. Our job is just a single line of code, so a lambda expression makes sense here.

It doesn't matter if you run this using an ExecutorService, like the code above, or with Threads directly, like the code below; both show the same symptoms.

```
class ThreadTestDrive {
    public static void main (String[] args) {
        Thread myThread = new Thread(() ->
            System.out.println("top o' the stack"));
        myThread.start();
        System.out.println("back in main");
    }
}
```

Notice how the order changes randomly. Sometimes the new thread finishes first, and sometimes the main thread finishes first.

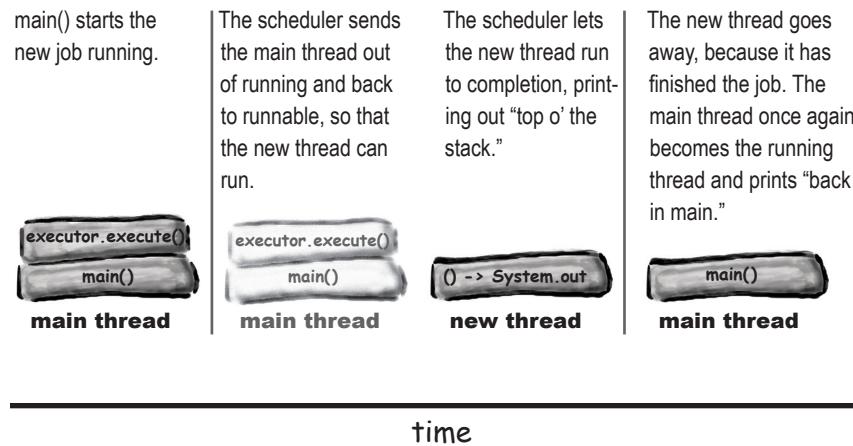
Produced this output:

```
File Edit Window Help PickMe
% java ExecutorTestDrive
back in main
top o' the stack
% java ExecutorTestDrive
top o' the stack
back in main
% java ExecutorTestDrive
top o' the stack
back in main
% java ExecutorTestDrive
top o' the stack
back in main
% java ExecutorTestDrive
top o' the stack
back in main
% java ExecutorTestDrive
top o' the stack
back in main
% java ExecutorTestDrive
top o' the stack
back in main
% java ExecutorTestDrive
top o' the stack
back in main
% java ExecutorTestDrive
top o' the stack
back in main
% java ExecutorTestDrive
top o' the stack
```

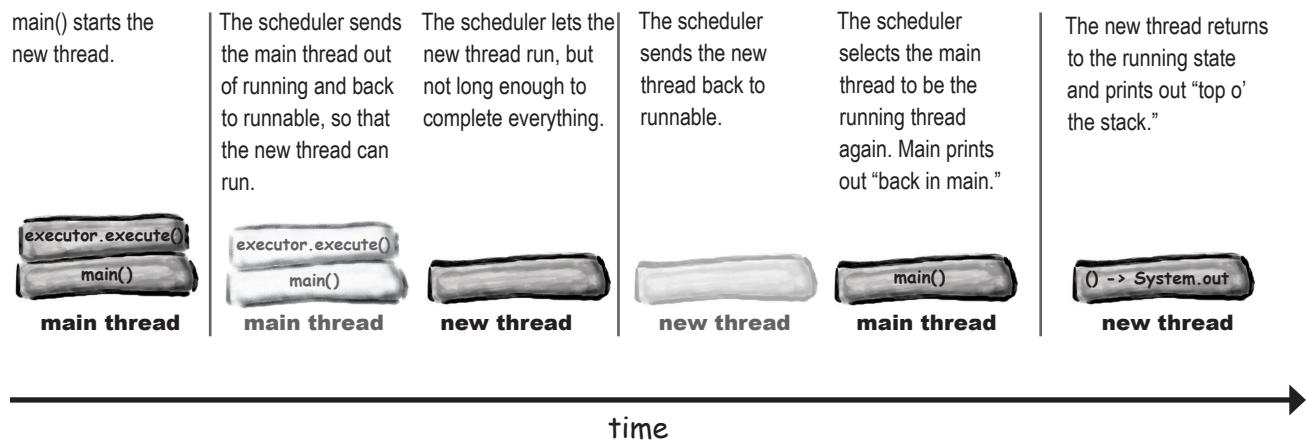
How did we end up with different results?

Multithreaded programs are *not deterministic*; they don't run the same way every time. The thread scheduler can schedule each thread differently each time the program runs.

Sometimes it runs like this:



And sometimes it runs like this:



Even if the new thread is tiny, if it has only one line of code to run like our lambda expression, it can still be interrupted by the thread scheduler.

there are no Dumb Questions

Q: Should I use a lambda expression for my Runnable or create a new class that implements Runnable?

A: It depends upon how complicated your job is, and also on whether you think it's easier to understand as a lambda expression or a class. Lambda expressions are great for when the job is really tiny, like our single-line "print" example. Lambda expressions (or method references) may also work if you have a few lines of code in another method that you want to turn into a job:

```
executor.execute(() -> printMsg());
```

You'll most likely want to use a full Runnable class if your job needs to store things in fields and/or if your job is made up of a number of methods. This is more likely when your jobs are more complex.

Q: What's the advantage of using an ExecutorService? So far, it works the same as creating a Thread and starting it.

A: It's true that for these simple examples, where we're starting just one thread, letting it run, and then stopping our application, the two approaches seem similar. ExecutorServices become really helpful when we're starting lots of independent jobs. We don't necessarily want to create a new Thread for each of these jobs, and we don't want to keep track of all these Threads. There are different ExecutorService implementations depending upon how many threads we'll want to start (or especially if we don't know how many Threads we'll need), including ExecutorServices that create Thread pools. Thread pools let us reuse Thread instances, so we don't have to pay the cost of starting up new Threads for every job. We'll explore this in more detail later.

BULLET POINTS

- A thread with a lowercase "t" is a separate thread of execution in Java.
- Every thread in Java has its own call stack.
- A Thread with a capital "T" is the java.lang.Thread class. A Thread object represents a thread of execution.
- A thread needs a job to do. The job can be an instance of something that implements the Runnable interface.
- The Runnable interface has just a single method, run(). This is the method that goes on the bottom of the new call stack. In other words, it is the first method to run in the new thread.
- Because the Runnable interface has just a single method, you can use lambda expressions where a Runnable is expected.
- Using the Thread class to run separate jobs is no longer the preferred way to create multithreaded applications in Java. Instead, use an Executor or an ExecutorService.
- The Executors class has helper methods that can create standard ExecutorServices to use to start new jobs.
- A thread is in the NEW state when it has not yet started.
- When a thread has been started, a new stack is created, with the Runnable's run() method on the bottom of the stack. The thread is now in the RUNNABLE state, waiting to be chosen to run.
- A thread is said to be RUNNING when the JVM's thread scheduler has selected it to be the currently running thread. On a single-processor machine, there can be only one currently running thread.
- Sometimes a thread can be moved from the RUNNING state to a temporarily NON-RUNNABLE state. A thread might be blocked because it's waiting for data from a stream, because it has gone to sleep, or because it is waiting for an object's lock. We'll see locks in the next chapter.
- Thread scheduling is not guaranteed to work in any particular way, so you cannot be certain that threads will take turns nicely.



Putting a thread to sleep

One way to help your threads take turns is to put them to sleep periodically. All you need to do is call the static `sleep()` method, passing it the amount of time you want the thread to sleep for, in milliseconds.

For example:

```
Thread.sleep(2000);
```

will knock a thread out of the running state and keep it out of the runnable state for two seconds. The thread *can't* become the running thread again until after at least two seconds have passed.

A bit unfortunately, the sleep method throws an `InterruptedException`, a checked exception, so all calls to sleep must be wrapped in a try/catch (or declared). So a sleep call really looks like this:

```
try {
        Thread.sleep(2000);
} catch(InterruptedException ex) {
        ex.printStackTrace();
}
```

Now you know that your thread won't wake up *before* the specified duration, but is it possible that it will wake up some time *after* the “timer” has expired? Effectively, yes. The thread won't automatically wake up at the designated time and become the currently running thread. When a thread wakes up, the thread is once again at the mercy of the thread scheduler; therefore, there are no guarantees about how long the thread will be out of action.

Putting a thread to sleep gives the other threads a chance to run.

When the thread wakes up, it always goes back to the runnable state and waits for the thread scheduler to choose it to run again.

It can be hard to understand how much time a number of milliseconds represents. There's a convenience method on `java.util.concurrent.TimeUnit` that we can use to make a more readable sleep time:

```
TimeUnit.MINUTES.sleep(2);
```

which may be easier to understand than:

```
Thread.sleep(120000);
```

(You still need to wrap both in a try-catch, though.)

Using sleep to make our program more predictable

Remember our earlier example that kept giving us different results each time we ran it? Look back and study the code and the sample output. Sometimes main had to wait until the new thread finished (and printed “top o’ the stack”), while other times the new thread would be sent back to runnable before it was finished, allowing the main thread to come back in and print out “back in main.” How can we fix that? Stop for a moment and answer this question: “Where can you put a sleep() call, to make sure that “back in main” always prints before “top o’ the stack”?

```
class PredictableSleep {
    public static void main (String[] args) {
        ExecutorService executor =
            Executors.newSingleThreadExecutor();
        executor.execute(() -> sleepThenPrint());
        System.out.println("back in main");
        executor.shutdown();
    }

    private static void sleepThenPrint() {
        try {
            TimeUnit.SECONDS.sleep(2);
        } catch (InterruptedException e) {
            e.printStackTrace();
        }
        System.out.println("top o' the stack");
    }
}
```

Thread.sleep() throws a checked Exception that we need to catch or declare. Because catching the Exception makes the job’s code a bit longer, we’ve put it into its own method.

There will be a pause (for about two seconds) before we get to this line, which prints out “top o’ the stack.”

This is what we want—a consistent order of print statements:

```
File Edit Window Help SnoozeButton
% java PredictableSleep
back in main
top o' the stack
% java PredictableSleep
back in main
top o' the stack
% java PredictableSleep
back in main
top o' the stack
% java PredictableSleep
back in main
top o' the stack
% java PredictableSleep
back in main
top o' the stack
% java PredictableSleep
back in main
top o' the stack
```

Instead of putting a lambda with an ugly try-catch inside, we’ve put the job code inside a method. We’re calling the method from this lambda

Calling sleep here will force the new thread to leave the currently running state. The main thread will get a chance to print out “back in main.”



Brain Barbell

Can you think of any problems with forcing your threads to sleep for a set amount of time? How long will it take to run this code 10 times?

There are downsides to forcing the thread to sleep

① The program has to wait for at least that amount of time.

If we put the thread to sleep for two seconds, the thread will be non-runnable for that time.

When it wakes up, it won't automatically become the currently running thread. When a thread wakes up, the thread is once again at the mercy of the thread scheduler. Our application is going to be hanging around for at least those two seconds, probably more. This might not sound like a big deal, but imagine a bigger program full of these pauses *intentionally* slowing down the application.

② How do you know the other job will finish in that time?

We put the new thread to sleep for two seconds, assuming that the main thread would be the running thread, and complete its work in that time. But what if the main thread took longer to finish than that? What if another thread, running a longer job, was scheduled instead? One of the ways people deal with this is to set sleep times that are much longer than they'd expect a job to take, but then our first problem becomes even more of a problem.



Is there a way for one thread to tell another that it has finished what it's working on? That way, the main thread could just wait for that signal and then carry on when it knows it's safe to go.

A better alternative: wait for the perfect time.

What we really wanted in our example was to wait until a specific thing had happened in our main thread before carrying on with our new thread. Java supports a number of different mechanisms to do this, like Future, CyclicBarrier, Semaphore, and CountDownLatch.

To coordinate events happening on multiple threads, one thread may need to wait for a specific signal from another thread before it can continue.

Counting down until ready

You can make threads *count down* when significant events have happened. A thread (or threads) can wait for all these events to complete before continuing. You might be counting down until a minimum number of clients have connected, or a number of services have been started.

This is what **`java.util.concurrent.CountDownLatch`** is for. You set a number to count down from. Then any thread can tell the latch to count down when a relevant event has happened.

In our example, we have only one thing we want to count—our new thread should wait until the main thread has printed “back in main” before it can continue.

```
import java.util.concurrent.*;
class PredictableLatch {
    public static void main (String[] args) {
        ExecutorService executor = Executors.newSingleThreadExecutor();
        CountDownLatch latch = new CountDownLatch(1); ←

        executor.execute(() -> waitForLatchThenPrint(latch));
        System.out.println("back in main");
        latch.countDown(); ← Tell the latch to count down when the
        executor.shutdown(); main method has printed its message.
    }
}
```

```
private static void waitForLatchThenPrint(CountDownLatch latch) {  
    try {  
        latch.await(); // Wait for the main thread to print out its message. This  
    } catch (InterruptedException e) { // thread will be in a non-runnable state while it's waiting.  
        e.printStackTrace();  
    }  
    System.out.println("top o' the stack");  
}
```

await() can throw an InterruptedException, which needs to be caught or declared.

The code is really similar to the code that performs a sleep; the main difference is the latch.countDown in the main method. The performance difference is significant, though. Instead of having to wait *at least* two seconds to make sure main has printed its message, the new thread waits only until the main method has printed its “back in main” message.

To get an idea of the performance difference this might make on a real system, when this latch code was run on a MacBook 100 times, it took around 50 milliseconds to finish *all* one hundred runs, and the output was in the correct order *every time*. If running the sleep() version just one time takes over 2 seconds (2000 milliseconds), imagine how long it took to run 100 times*....

* 200331 milliseconds. That's over 4000x slower.

you are here ►

Making and starting two threads (or more!)

What happens if we want to start more than one job in addition to our main thread? Clearly, we can't use Executors.newSingleThreadExecutor() if we want to run more than one thread. What else is available?

(Just a few of the factory methods)

java.util.concurrent.Executors

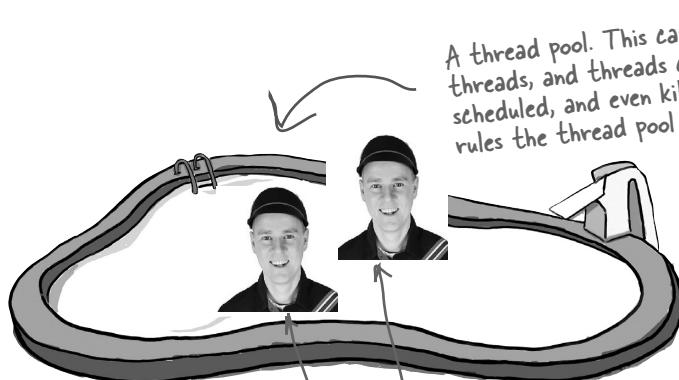
- ExecutorService newCachedThreadPool()**
Creates a thread pool that creates new threads as needed, but will reuse previously constructed threads when they are available.
- ExecutorService newFixedThreadPool(int nThreads)**
Creates a thread pool that reuses a fixed number of threads operating off a shared unbounded queue.
- ScheduledExecutorService newScheduledThreadPool(int corePoolSize)**
Creates a thread pool that can schedule commands to run after a given delay, or to execute periodically.
- ExecutorService newSingleThreadExecutor()**
Creates an Executor that uses a single worker thread operating off an unbounded queue.
- ScheduledExecutorService newSingleThreadScheduledExecutor()**
Creates a single-threaded executor that can schedule commands to run after a given delay, or to execute periodically.
- ExecutorService newWorkStealingPool()**
Creates a work-stealing thread pool using the number of available processors as its target parallelism level.

These ExecutorServices use some form of Thread Pool. This is a collection of Thread instances that can be used (and reused) to perform jobs.

How many threads are in the pool, and what to do if there are more jobs to run than threads available, depends on the ExecutorService implementation.

Pooling Threads

Using a pool of resources, especially ones that are expensive to create like Threads or database connections, is a common pattern in application code.



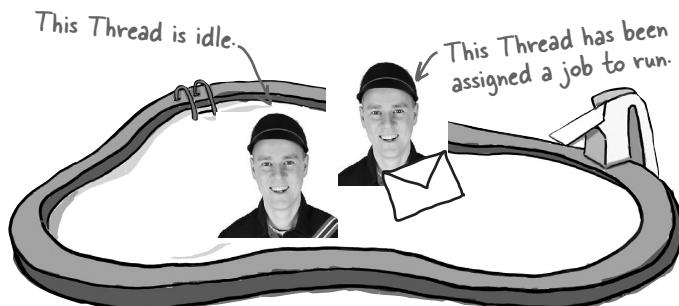
A thread pool. This can contain one or more threads, and threads can be added, used, reused, scheduled, and even killed according to whatever rules the thread pool was set up with.

The threads available for running jobs. How many threads are allowed and how they are used is determined by the pool.

When you create a new ExecutorService, its pool may be started with some threads to begin with, or the pool may be empty.

You can create an ExecutorService with a thread pool using one of the helper methods from the Executors class.

```
ExecutorService threadPool =
    Executors.newCachedThreadPool();
```



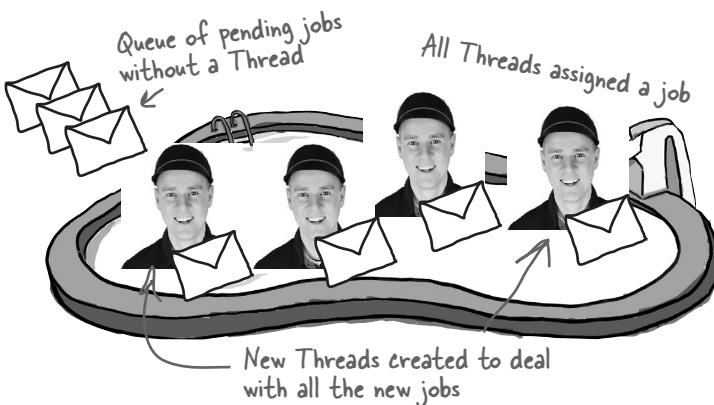
This Thread is idle.

This Thread has been assigned a job to run.

You can use the pool's threads to run your job by giving the job to the ExecutorService. The ExecutorService can then figure out if there's a free Thread to run the job.

```
threadPool.execute(() -> run("Job 1"));
```

This means an ExecutorService can **reuse** threads; it doesn't just create and destroy them.



Queue of pending jobs without a Thread

All Threads assigned a job

New Threads created to deal with all the new jobs

As you give the ExecutorService more jobs to run, it *may* create and start new Threads to handle the jobs. It *may* store the jobs in a queue if there are more jobs than Threads.

How an ExecutorService deals with additional jobs depends on how it is set up.

```
threadPool.execute(() -> run("Job 324"));
```

The ExecutorService may also **terminate** Threads that have been idle for some period of time. This can help to minimize the amount of hardware resources (CPU, memory) your application needs.

Running multiple threads

The following example runs two jobs, and uses a fixed-sized thread pool to create two threads to run the jobs. Each thread has the same job: run in a loop, printing the currently running thread's name with each iteration.

```

import java.util.concurrent.ExecutorService;
import java.util.concurrent.Executors;

public class RunThreads {
    public static void main(String[] args) {
        ExecutorService threadPool = Executors.newFixedThreadPool(2);
        threadPool.execute(() -> runJob("Job 1"));
        threadPool.execute(() -> runJob("Job 2"));
        threadPool.shutdown();
    }

    public static void runJob(String jobName) {
        for (int i = 0; i < 25; i++) {
            String threadName = Thread.currentThread().getName();
            System.out.println(jobName + " is running on " + threadName);
        }
    }
}

```

Create an ExecutorService with a fixed-sized thread pool (we know we're going to run only two jobs).

A lambda expression that represents our Runnable job. If you don't want to use lambdas, here you'd pass in a new instance of your Runnable class, like we did when we created MyRunnable earlier in the chapter.

The job is to run through this loop, printing the thread's name each time.

What will happen?

Will the threads take turns? Will you see the thread names alternating? How often will they switch? With each iteration? After five iterations?

You already know the answer: *we don't know!* It's up to the scheduler. And on your OS, with your particular JVM, on your CPU, you might get very different results.

Running this on a modern multicore system, the two jobs will likely run in parallel, but there's no guarantee that this means they will complete in the same amount of time or output values at the same rate.

```

File Edit Window Help Globetrotter
% java RunThreads
Job 1 is running on pool-1-thread-1

```

Closing time at the thread pool

You may have noticed that our examples have a `threadPool.shutdown()` at the end of the main methods. Although the thread pools will take care of our individual Threads, we do need to be responsible adults and close the pool when we're finished with it. That way, the pool can empty its job queue and shut down all of its threads to free up system resources.

ExecutorService has two shutdown methods. You can use either, but to be safe we'd use both:

① ExecutorService.shutdown()

Calling `shutdown()` asks the ExecutorService nicely if it wouldn't mind awfully wrapping things up so everyone can go home.

All of the Threads that are currently running jobs are allowed to finish those jobs, and any jobs waiting in the queue will also be finished off. The ExecutorService will reject any new jobs too.

If you need your code to wait until all of those things are finished, you can use `awaitTermination` to sit and wait until it's finished. You give `awaitTermination` a maximum amount of time to wait for everything to end, so `awaitTermination` will hang around until either the ExecutorService has finished everything or the timeout has been reached, whichever is earlier.

② ExecutorService.shutdownNow()

Everybody out! When this is called, the ExecutorService will try to stop any Threads that are running, will not run any waiting jobs, and definitely won't let anyone else into the pool.

Use this if you need to put a stop to everything. This is sometimes used after first calling `shutdown()` to give the jobs a chance to finish before pulling the plug entirely.

```

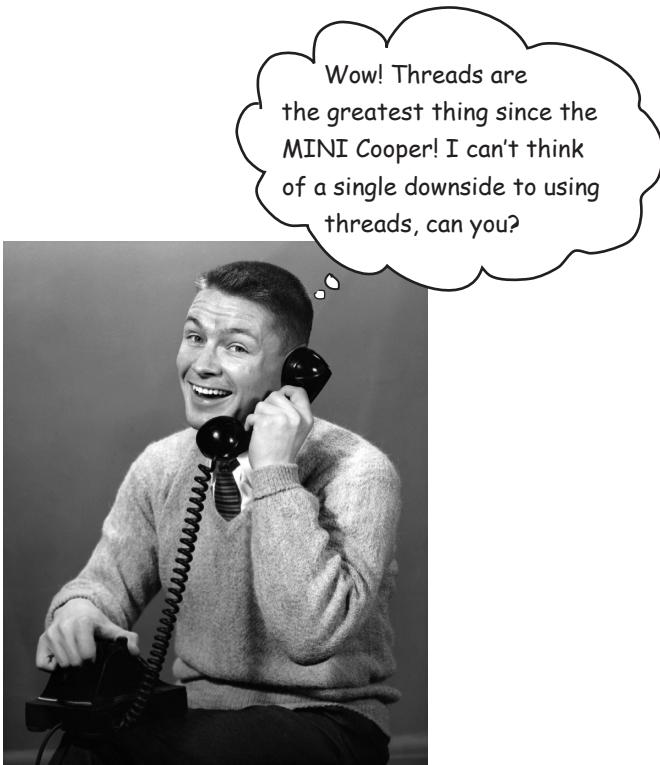
public class ClosingTime {
    public static void main(String[] args) {
        ExecutorService threadPool = Executors.newFixedThreadPool(2);

        threadPool.execute(new LongJob("Long Job"));
        threadPool.execute(new ShortJob("Short Job"));

        threadPool.shutdown();
        try {
            boolean finished = threadPool.awaitTermination(5, TimeUnit.SECONDS);
            System.out.println("Finished? " + finished);
        } catch (InterruptedException e) {
            e.printStackTrace();
        }
        threadPool.shutdownNow();
    }
}

Ask the ExecutorService to shut down. If you call execute() with a job after this, you will get a RejectedExecutionException. The ExecutorService will continue to run all the jobs that are running, and run any waiting jobs too.
Create a thread pool with just two threads.
Start two jobs, a short one that just prints the name and a "long" one that uses a sleep so it can pretend to be a long-running job (LongJob and ShortJob are classes that implement Runnable).
Wait up to 5 seconds for the ExecutorService to finish everything. If this method hits the timeout before everything has finished, it returns "false".
At this point, we tell the ExecutorService to stop everything right now. If everything was already shut down, that's fine; this won't do anything.

```



Um, yes. There IS a dark side. Multithreading can lead to concurrency "issues."

Concurrency issues lead to race conditions. Race conditions lead to data corruption. Data corruption leads to fear...you know the rest.

It all comes down to one potentially deadly scenario: two or more threads have access to a single object's *data*. In other words, methods executing on two different stacks are both calling, say, getters or setters on a single object on the heap.

It's a whole "left-hand-doesn't-know-what-the-right-hand-is-doing" thing. Two threads, without a care in the world, humming along executing their methods, each thread thinking that he is the One True Thread. The only one that matters. After all, when a thread is not running, and in runnable (or blocked) it's essentially knocked unconscious. When it becomes the currently running thread again, it doesn't know that it ever stopped.

BULLET POINTS

- The static `Thread.sleep()` method forces a thread to leave the running state for at least the duration passed to the sleep method. `Thread.sleep(200)` puts a thread to sleep for 200 milliseconds.
- You can also use the sleep method on `java.util.concurrent.TimeUnit`, for example `TimeUnit.SECONDS.sleep(2)`.
- The `sleep()` method throws a checked exception (`InterruptedException`), so all calls to `sleep()` must be wrapped in a try/catch, or declared.
- There are different mechanisms to give threads a chance to wait for each other. You can use `sleep()`, but you can also use `CountDownLatch` to wait for the right number of events to have happened before continuing.
- Managing threads directly can be a lot of work. Use the factory methods in `Executors` to create an `ExecutorService`, and use this service to run `Runnable` jobs.
- Thread pools can manage creation, reuse, and destruction of threads so you don't have to.
- `ExecutorServices` should be shut down correctly so the jobs are finished and threads terminated. Use `shutdown()` for graceful shutdown, and `shutdownNow()` to kill everything.



It's a cliff-hanger!

A dark side? Race conditions?? Data corruption?! But what can we DO about those things? Don't leave us hanging!

Stay tuned for the next chapter, where we address these issues and more...



Exercise

A bunch of Java and network terms, in full costume, are playing a party game, “Who am I?” They’ll give you a clue—you try to guess who they are based on what they say. Assume they always tell the truth about themselves. If they happen to say something that could be true for more than one attendee, then write down all for whom that sentence applies. Fill in the blanks next to the sentence with the names of one or more attendees.

Tonight’s attendees:

InetSocketAddress, SocketChannel, IP address, host name, port, Socket, ServerSocketChannel, Thread, thread pool, Executors, ExecutorService, CountDownLatch, Runnable, InterruptedException, Thread.sleep()



I need to be shut down or I might live forever _____

I let you talk to a remote machine _____

I might be thrown by sleep() and await() _____

If you want to reuse Threads, you should use me _____

You need to know me if you want to connect to another machine _____

I’m like a separate process running on the machine _____

I can give you the ExecutorService you need _____

You need one of me if you want clients to connect to me _____

I can help you make your multithreaded code more predictable _____

I represent a job to run _____

I store the IP address and port of the server _____

—————> Answers on page 636.

New and improved SimpleChatClient

Way back near the beginning of the chapter, we built the SimpleChatClient that could *send* outgoing messages to the server but couldn't receive anything. Remember? That's how we got onto this whole thread topic in the first place, because we needed a way to do two things at once: send messages *to* the server (interacting with the GUI) while simultaneously reading incoming messages *from* the server, displaying them in the scrolling text area.

This is the New Improved chat client that can both send and receive messages, thanks to the power of multithreading! Remember, you need to run the chat server first to run this code.

```

import javax.swing.*;
import java.awt.*;
import java.awt.event.*;
import java.io.*;
import java.net.InetSocketAddress;
import java.nio.channels.*;
import java.util.concurrent.*;

import static java.nio.charset.StandardCharsets.UTF_8;

public class SimpleChatClient {
    private JTextArea incoming;
    private JTextField outgoing;
    private BufferedReader reader;
    private PrintWriter writer;

    public void go() {
        setUpNetworking();

        JScrollPane scroller = createScrollableTextArea();
        outgoing = new JTextField(20);
        JButton sendButton = new JButton("Send");
        sendButton.addActionListener(e -> sendMessage());

        JPanel mainPanel = new JPanel();
        mainPanel.add(scroller);
        mainPanel.add(outgoing);
        mainPanel.add(sendButton);

        ExecutorService executor = Executors.newSingleThreadExecutor();
        executor.execute(new IncomingReader());
    }

    JFrame frame = new JFrame("Ludicrously Simple Chat Client");
    frame.getContentPane().add(BorderLayout.CENTER, mainPanel);
    frame.setSize(400, 350);
    frame.setVisible(true);
    frame.setDefaultCloseOperation(WindowConstants.EXIT_ON_CLOSE);
}

```

Yes, there really IS an end to this chapter.
But not yet.

This is mostly GUI code you've seen before. Nothing special except the highlighted part where we start the new "reader" thread.

We've got a new job, an inner class, which is a Runnable. The job is to read from the server's socket stream, displaying any incoming messages in the scrolling text area. We start this job using a single thread executor since we know we want to run only this one job.

```

private JScrollPane createScrollableTextArea() {
    incoming = new JTextArea(15, 30);
    incoming.setLineWrap(true);
    incoming.setWrapStyleWord(true);
    incoming.setEditable(false);
    JScrollPane scroller = new JScrollPane(incoming);
    scroller.setVerticalScrollBarPolicy(ScrollPaneConstants.VERTICAL_SCROLLBAR_ALWAYS);
    scroller.setHorizontalScrollBarPolicy(ScrollPaneConstants.HORIZONTAL_SCROLLBAR_NEVER);
    return scroller;
}

private void setUpNetworking() {
    try {
        InetSocketAddress serverAddress = new InetSocketAddress("127.0.0.1", 5000);
        SocketChannel socketChannel = SocketChannel.open(serverAddress);

        reader = new BufferedReader(Channels.newReader(socketChannel, UTF_8));
        writer = new PrintWriter(Channels.newWriter(socketChannel, UTF_8));

        System.out.println("Networking established.");
    } catch (IOException ex) {
        ex.printStackTrace();
    }
}

private void sendMessage() {
    writer.println(outgoing.getText());
    writer.flush();
    outgoing.setText("");
    outgoing.requestFocus();
}
}

public class IncomingReader implements Runnable {
    public void run() {
        String message;
        try {
            while ((message = reader.readLine()) != null) {
                System.out.println("read " + message);
                incoming.append(message + "\n");
            }
        } catch (IOException ex) {
            ex.printStackTrace();
        }
    }
}

public static void main(String[] args) {
    new SimpleChatClient().go();
}
}

```

A helper method, like we saw back in Chapter 1b, to create a scrolling text area.

We're using `Channels` to create a new reader and writer for the `SocketChannel` that's connected to the server. The writer sends messages to the server, and now we're using a reader so that the reader job can get messages from the server.

Nothing new here. When the user clicks the send button, this method sends the contents of the text field to the server.

This is what the thread does!!
In the `run()` method, it stays in a loop (as long as what it gets from the server is not null), reading a line at a time and adding each line to the scrolling text area (along with a new line character).

Remember, the Chat Server code was the Ready-Bake Code from page 606.

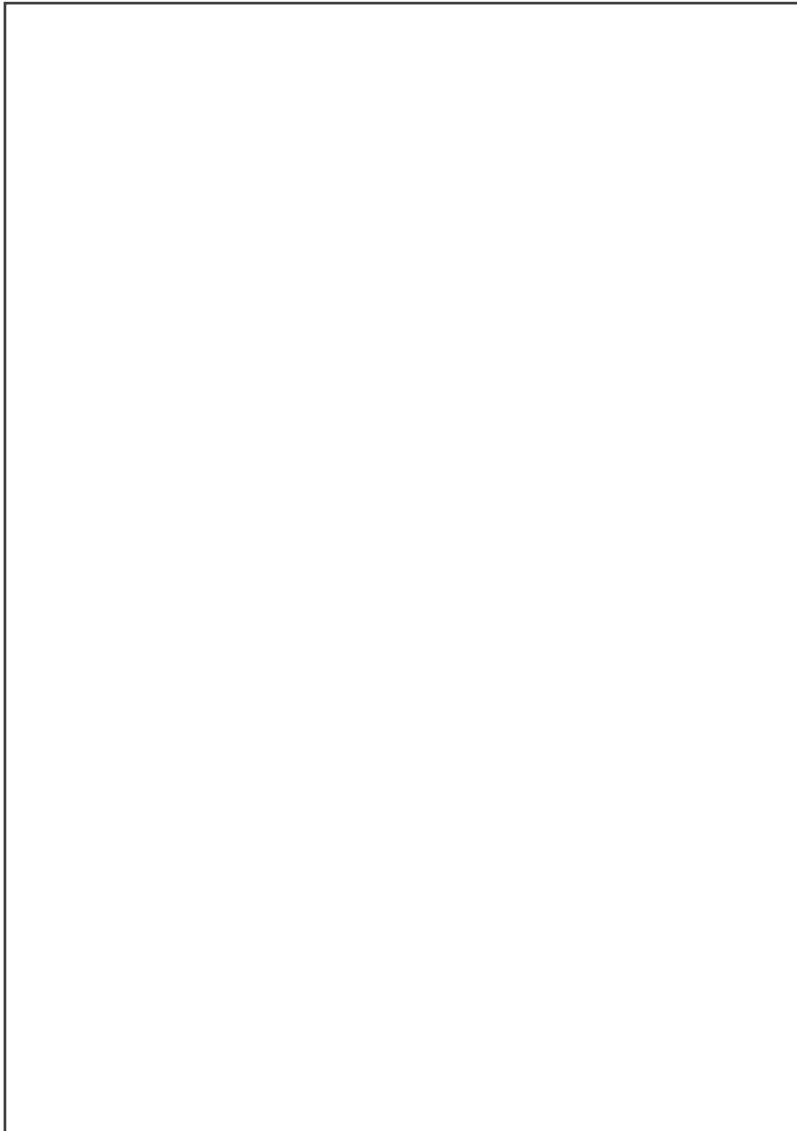
exercise: code magnets



Code Magnets

A working Java program is scrambled up on the fridge (see the next page). Can you reconstruct the code snippets on the next page to make a working Java program that produces the output listed below?

To get it to work, you will need to be running the `SimpleChatServer` from page 606.



→ Answers on page 636.

```
File Edit Window Help StillThere
% java PingingClient
Networking established
09:27:06 Sent ping 0
09:27:07 Sent ping 1
09:27:08 Sent ping 2
09:27:09 Sent ping 3
09:27:10 Sent ping 4
09:27:11 Sent ping 5
09:27:12 Sent ping 6
09:27:13 Sent ping 7
09:27:14 Sent ping 8
09:27:15 Sent ping 9
```

Code Magnets, continued

```

String message = "ping " + i;

try (SocketChannel channel = SocketChannel.open(server)) {
    import static java.nio.charset.StandardCharsets.UTF_8;
    import static java.time.LocalDateTime.now;
    import static java.time.format.DateTimeFormatter.ofLocalizedTime;
    e.printStackTrace();
} catch (IOException | InterruptedException e) {
    public class PingingClient {
        }
        writer.println(message);
    }

PrintWriter writer = new PrintWriter(Channels.newWriter(channel, UTF_8));

import java.io.*;
import java.net.InetSocketAddress;
import java.nio.channels.*;
import java.time.format.FormatStyle;
import java.util.concurrent.TimeUnit;
for (int i = 0; i < 10; i++) {
    System.out.println(currentTime + " Sent " + message);
    writer.flush();
}

InetSocketAddress server = new InetSocketAddress("127.0.0.1", 5000);
}

System.out.println("Networking established");
    TimeUnit.SECONDS.sleep(1);

public static void main(String [] args) {
    String currentTime = now().format(ofLocalizedTime(FormatStyle.MEDIUM));
}

```

exercise solutions

Who Am I? (from page 631)

I need to be shut down or I might live forever
I let you talk to a remote machine
I might be thrown by sleep() and await()
If you want to reuse Threads, you should use me
You need to know me if you want to connect to another machine
I'm like a separate process running on the machine
I can give you the ExecutorService you need
You need one of me if you want clients to connect to me
I can help you make your multithreaded code more predictable
I represent a job to run
I store the IP address and port of the server

ExecutorService
SocketChannel, Socket
InterruptedException
Thread pool, ExecutorService
IP Address, Host name, port
Thread
Executors
ServerSocketChannel
Thread.sleep(), CountDownLatch
Runnable
InetSocketAddress

Exercise Solutions



Code Magnets

(from pages 634–635)

```
import java.io.*;
import java.net.InetSocketAddress;
import java.nio.channels.*;
import java.time.format.FormatStyle;
import java.util.concurrent.TimeUnit;

import static java.nio.charset.StandardCharsets.UTF_8;
import static java.time.LocalDateTime.now;
import static java.time.format.DateTimeFormatter.ofLocalizedTime;

public class PingingClient {

    public static void main(String[] args) {
        InetSocketAddress server = new InetSocketAddress("127.0.0.1", 5000);
        try (SocketChannel channel = SocketChannel.open(server)) {
            PrintWriter writer = new PrintWriter(Channels.newWriter(channel, UTF_8));
            System.out.println("Networking established");

            for (int i = 0; i < 10; i++) {
                String message = "ping " + i;
                writer.println(message);
                writer.flush();
                String currentTime = now().format(ofLocalizedTime(FormatStyle.MEDIUM));
                System.out.println(currentTime + " Sent " + message);
                TimeUnit.SECONDS.sleep(1);
            }
        } catch (IOException | InterruptedException e) {
            e.printStackTrace();
        }
    }
}
```

You should get the same output even if you move the sleep() to somewhere else inside this for loop.

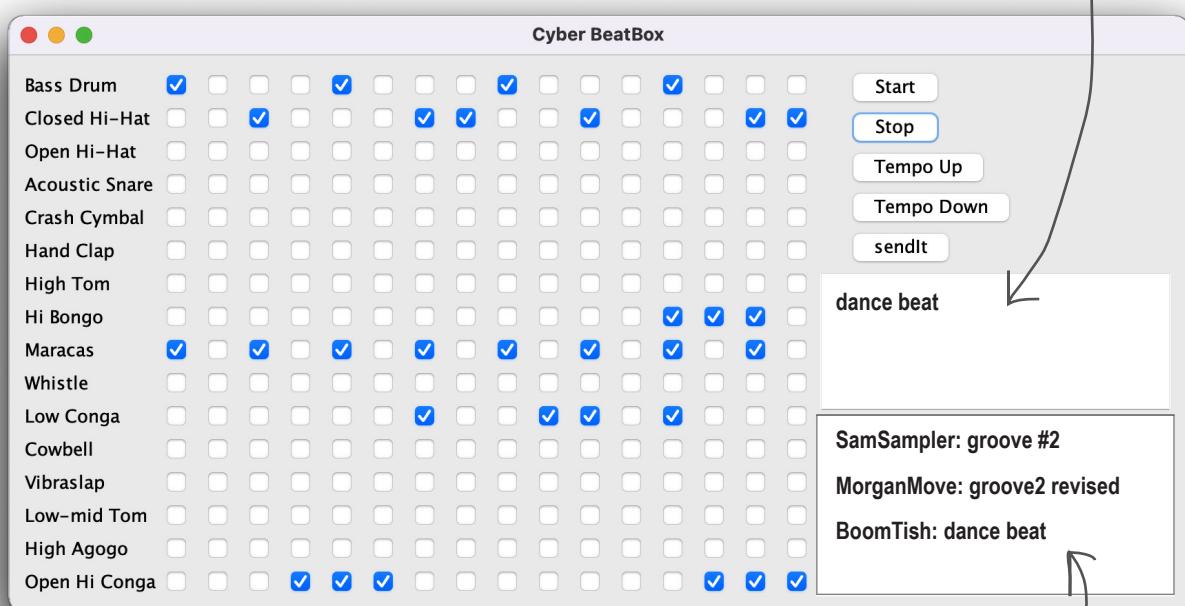
Catching all Exceptions at the end because we do the same thing with them all.

This is one way of getting the current time and turning it into a String in the format of Hours:Minutes:Seconds.

You can catch the InterruptedException thrown by sleep() inside the for loop, or you can catch all the Exceptions at the end of the method.

Code Kitchen

Your message gets sent to the other players, along with your current beat pattern, when you hit "sendIt."



Now you've seen how to build a chat client, we have the last version of the BeatBox!

It connects to a simple MusicServer so that you can send and receive beat patterns with other clients.

The code is really long, so the complete listing is actually in Appendix A.

Dealing with Concurrency Issues

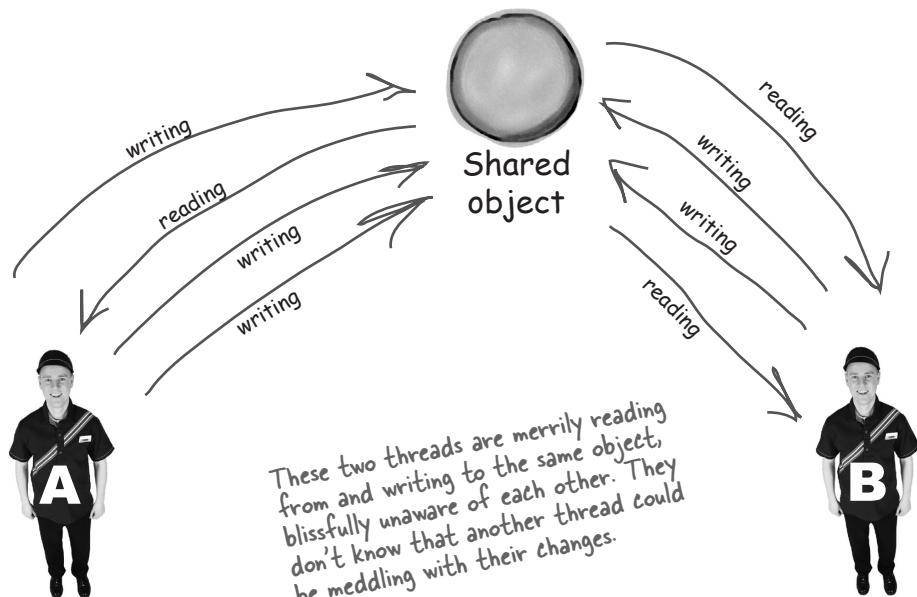


Doing two or more things at once is hard. Writing multithreaded code is easy. Writing multithreaded code that works the way you expect can be much harder. In this final chapter, we’re going to show you some of the things that can go wrong when two or more threads are working at the same time. You’ll learn about some of the tools in `java.util.concurrent` that can help you to write multithreaded code that works correctly. You’ll learn how to create immutable objects (objects that don’t change) that are safe for multiple threads to use. By the end of the chapter, you’ll have a lot of different tools in your toolkit for working with concurrency.

What could possibly go wrong?

At the end of the last chapter we hinted that things may not all be rainbows and sunshine when you're working with multithreaded code. Well, actually, we did more than hint! We outright said:

"It all comes down to one potentially deadly scenario: two or more threads have access to a single object's data."



Brain Barbell

Why is it a problem if two threads are both reading and writing?

If a thread reads the object's data before changing it, why is it a problem that another thread might also be writing at the same time?

Marriage in Trouble.

Can this couple be saved?

Next, on a very special Dr. Steve Show

[Transcript from episode #42]

Welcome to the Dr. Steve show.



We've got a story today that's centered around one of the top reasons why couples split up—finances.

Today's troubled pair, Ryan and Monica, share a bank account. But not for long if we can't find a solution. The problem? The classic "two people—one bank account" thing.

Here's how Monica described it to me:

"Ryan and I agreed that neither of us will overdraw the checking account. So the procedure is, whoever wants to spend money *must* check the balance in the account *before* withdrawing cash or spending on a card. It all seemed so simple. But suddenly we're getting hit with overdraft fees!"

I thought it wasn't possible; I thought our procedure was safe. But then *this* happened:

Ryan had a full online shopping cart totalling \$50. He checked the balance in the account and saw that it was \$100. No problem. So he started the checkout procedure.

And that's where *I* come in; while Ryan was filling in the shipping details, I wanted to spend \$100. I checked the balance, and it's \$100 (because Ryan hasn't clicked the "Pay" button yet), so I think, no problem. So I spend the money, and again no problem. But then Ryan finally pays, and we're suddenly overdrawn! He didn't know that I was spending money at the same time, so he just went ahead and completed his transaction without checking the balance again. You've got to help us, Dr. Steve!"

Is there a solution? Are they doomed? We can't help them with their online shopping addiction, but can we make sure one of them can't start spending while the other one is shopping?

Take a moment and think about that while we go to a commercial break.



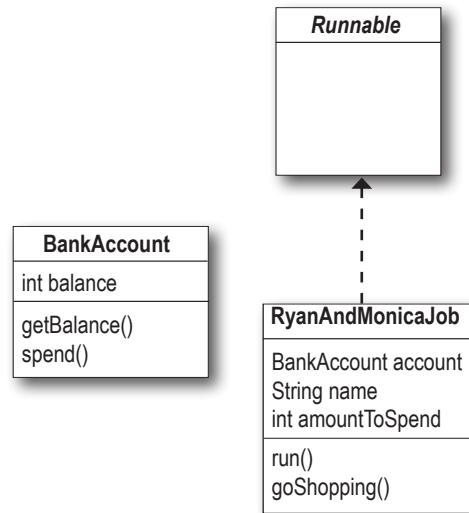
Ryan and Monica: victims of the "two people, one account" problem.

The Ryan and Monica problem, in code

The following example shows what can happen when *two* threads (Ryan and Monica) share a *single* object (the bank account).

The code has two classes, BankAccount and RyanAndMonicaJob. There's also a RyanAndMonicaTest with a main method to run everything. The RyanAndMonicaJob class implements Runnable, and represents the behavior that Ryan and Monica both have—checking the balance and spending money.

The RyanAndMonicaJob class has instance variables for the shared BankAccount, the person's name, and the amount they want to spend. The code works like this:



① Make an instance of the shared bank account

Creating a new one will set up all the defaults correctly.

```
BankAccount account = new BankAccount();
```

② Make one instance of RyanAndMonicaJob for each person

We need one job for each person. We also need to give them access to the BankAccount and tell them how much to spend.

```
RyanAndMonicaJob ryan = new RyanAndMonicaJob("Ryan", account, 50);
RyanAndMonicaJob monica = new RyanAndMonicaJob("Monica", account, 100);
```

③ Create an ExecutorService and give it the two jobs

Since we know we have two jobs, one for Ryan and one for Monica, we can create a fixed-sized thread pool with two threads.

```
ExecutorService executor = Executors.newFixedThreadPool(2);
executor.execute(ryan);
executor.execute(monica);
```

④ Watch both jobs run

One thread represents Ryan, the other represents Monica. Both threads check the balance before spending money. Remember that when more than one thread is running at a time, you can't assume that your thread is the only one making changes to a shared object (e.g., the BankAccount). Even if there's only two lines of code related to the shared object, and they're right next to each other.

```
if (account.getBalance() >= amount) {
    account.spend(amount);
} else {
    System.out.println("Sorry, not enough money");
}
```

In the `goShopping()` method, do exactly what Ryan and Monica would do—check the balance and, if there's enough money, spend.

This should protect against overdrawing the account.

Except...if Ryan and Monica are spending money at the same time, the money in the bank account might be gone before the other one can spend it!

The Ryan and Monica example

```

import java.util.concurrent.*;

public class RyanAndMonicaTest {
    public static void main(String[] args) {
        BankAccount account = new BankAccount();
        RyanAndMonicaJob ryan = new RyanAndMonicaJob("Ryan", account, 50);
        RyanAndMonicaJob monica = new RyanAndMonicaJob("Monica", account, 100);
        ExecutorService executor = Executors.newFixedThreadPool(2);
        executor.execute(ryan); // Start both jobs
        executor.execute(monica); // running.
        executor.shutdown();
    } // Don't forget to shut the pool down.
}

class RyanAndMonicaJob implements Runnable {
    private final String name;
    private final BankAccount account;
    private final int amountToSpend;
    RyanAndMonicaJob(String name, BankAccount account, int amountToSpend) {
        this.name = name;
        this.account = account;
        this.amountToSpend = amountToSpend;
    }
    public void run() {
        goShopping(amountToSpend); // The run() method just calls goShopping()
    }
    private void goShopping(int amount) {
        if (account.getBalance() >= amount) {
            System.out.println(name + " is about to spend");
            account.spend(amount);
            System.out.println(name + " finishes spending");
        } else {
            System.out.println("Sorry, not enough for " + name);
        }
    }
}

class BankAccount {
    private int balance = 100; // The account starts with a
    public int getBalance() { // balance of $100.
        return balance;
    }
    public void spend(int amount) {
        balance = balance - amount;
        if (balance < 0) {
            System.out.println("Overdrawn!");
        }
    }
}

```

There will be only ONE instance of the BankAccount. That means both threads will access this one account.

Make two jobs that will do the withdrawal from the shared bank account, one for Monica and one for Ryan, passing in the amount they're going to spend.

Create a new thread pool with two threads for our two jobs.

The run() method just calls goShopping() with the amount they need to spend.

Check the account balance, and if there's enough money, we go ahead and spend the money, just like Ryan and Monica did.

We put in a bunch of print statements so we can see what's happening as it runs.

Ryan and Monica output

How did this happen? →

```
File Edit Window Help Visa
% java RyanAndMonicaTest
Ryan is about to spend
Monica is about to spend
Overdrawn!
Ryan finishes spending
Monica finishes spending
```

This code is not deterministic; it doesn't always produce the same result every time. You may need to run it a few times before you see the problem.

This is common with multithreaded code, since it depends upon which threads start first and when each thread gets its time on a CPU core.

Sometimes the code works correctly and they don't go overdrawn. →

```
File Edit Window Help WorksOnMyMachine
% java RyanAndMonicaTest
Ryan is about to spend
Ryan finishes spending
Sorry, not enough for Monica
```

The **goShopping()** method always checks the balance before making a withdrawal, but still we went overdrawn.

Here's one scenario:

Ryan checks the balance, sees that there's enough money, and goes to check out.

Meanwhile, Monica checks the balance. She, too, sees that there's enough money. She has no idea that Ryan is about to pay for something.

Ryan completes his purchase.

Monica completes her purchase. Big Problem! In between the time when she checked the balance and spent the money, Ryan had already spent money!

Monica's check of the account was not valid, because Ryan had already checked and was still in the middle of making a purchase.

Monica must be stopped from getting into the account until Ryan has finished, and vice versa.

They need a lock for account access!

The lock works like this:

- ① There's a lock associated with the bank account transaction (checking the balance and withdrawing money). There's only one key, and it stays with the lock until somebody wants to access the account.



The bank account transaction is unlocked when nobody is using the account.

- ② When Ryan wants to access the bank account (to check the balance and withdraw money), he locks the lock and puts the key in his pocket. Now nobody else can access the account, since the key is gone.



When Ryan wants to access the account, he secures the lock and takes the key.

- ③ Ryan keeps the key in his pocket until he finishes the transaction. He has the only key, so Monica can't access the account until Ryan unlocks the account and returns the key.

Now, even if Ryan gets distracted after he checks the balance, he has a guarantee that the balance will be the same when he spends the money, because he kept the key while he was doing something else!



When Ryan is finished, he unlocks the lock and returns the key. Now the key is available for Monica (or Ryan again) to access the account.

We need to check the balance and spend the money as one atomic thing



We need to make sure that once a thread starts a shopping transaction, *it must be allowed to finish* before any other thread changes the bank account.

In other words, we need to make sure that once a thread has checked the account balance, that thread has a guarantee that it can spend the money *before any other thread can check the account balance!*

Use the **synchronized** keyword on a method, or with an object, to lock an object so only one thread can use it at a time.

That's how you protect the bank account! We can put a lock on the bank account inside the method that does the banking transaction. That way, one thread gets to complete the whole transaction, start to finish, even if that thread is taken out of the "running" state by the thread scheduler or another thread is trying to make changes at exactly the same time.

On the next couple of pages we'll look at the different things that we can lock. With the Ryan and Monica example, it's quite simple—we want to wrap our shopping transaction in a block that locks the bank account:

```
private void goShopping(int amount) {  
    synchronized (account) {  
        if (account.getBalance() >= amount) {  
            System.out.println(name + " is about to spend");  
            account.spend(amount);  
            System.out.println(name + " finishes spending");  
        } else {  
            System.out.println("Sorry, not enough for " + name);  
        }  
    }  
}
```

(Note for you physics-savvy readers: yes, the convention of using the word "atomic" here does not reflect the whole subatomic particle thing. Think Newton, not Einstein, when you hear the word "atomic" in the context of threads or transactions. Hey, it's not OUR convention. If WE were in charge, we'd apply Heisenberg's Uncertainty Principle to pretty much everything related to threads.)

The synchronized keyword means that a thread needs a key in order to access the synchronized code.

To protect your data (like the bank account), synchronize the code that acts on that data.

there are no
Dumb Questions

Q: Why don't you just synchronize all the getters and setters from the class with the data you're trying to protect?

A: Synchronizing the getters and setters isn't enough. Remember, the point of synchronization is to make a specific section of code work ATOMICALLY. In other words, it's not just the individual methods we care about; it's methods that require **more than one step to complete!**

Think about it. We added a synchronized block inside the goShopping() method. If getBalance() and spend() were both synchronized instead, it wouldn't help—Ryan (or Monica) would have checked the balance returned the key! The whole point is to keep the key until **both** operations are completed.

Using an object's lock

Every object has a lock. Most of the time, the lock is unlocked, and you can imagine a virtual key sitting with it. Object locks come into play only when there is a **synchronized block** for an object (like in the last page) or a class has **synchronized methods**.

A method is synchronized if it has the synchronized keyword in the method declaration.

When an object has one or more synchronized methods, **a thread can enter a synchronized method only if the thread can get the key to the object's lock!**

The locks are not per *method*, they are per *object*. If an object has two synchronized methods, it doesn't *only* mean two threads can't enter the same method. It means you can't have two threads entering *any* of the synchronized methods. If you have two synchronized methods on the same object, method1() and method2(), if one thread is in method1(), a second thread can't enter method1(), obviously, but it *also can't enter method2()*, or any other synchronized method on that object.

Think about it. If you have multiple methods that can potentially act on an object's instance variables, all those methods need to be protected with synchronized.

The goal of synchronization is to protect critical data. But remember, you don't lock the data itself; you synchronize the methods that access that data.

So what happens when a thread is cranking through its call stack (starting with the run() method) and it suddenly hits a synchronized method? The thread recognizes that it needs a key for that object before it can enter the method. It looks for the key (this is all handled by the JVM; there's no API in Java for accessing object locks), and if the key is available, the thread grabs the key and enters the method.

From that point forward, the thread hangs on to that key like the thread's life depends on it. The thread won't give up the key until it completes the synchronized method or block. So while that thread is holding the key, no other threads can enter *any* of that object's synchronized methods, because the one key for that object won't be available.



Hey, this object's takeMoney() method is synchronized. I need to get this object's key before I can go in...



Every Java object has a lock. A lock has only one key.

Most of the time, the lock is unlocked and nobody cares.

But if an object has synchronized methods, a thread can enter one of the synchronized methods ONLY if the key for the object's lock is available. In other words, only if another thread hasn't already grabbed the one key.

using synchronized

Using synchronized methods

Can we synchronize the goShopping() method to fix Ryan and Monica's problem?

```
private synchronized void goShopping(int amount) {  
    if (account.getBalance() >= amount) {  
        System.out.println(name + " is about to spend");  
        account.spend(amount);  
        System.out.println(name + " finishes spending");  
    } else {  
        System.out.println("Sorry, not enough for " + name);  
    }  
}
```

It does NOT work!

```
File Edit Window Help WaitWhat  
% java RyanAndMonicaTest  
Ryan is about to spend  
Ryan finishes spending  
Monica is about to spend  
Overdrawn!  
Monica finishes spending
```

The synchronized keyword locks an object. The goShopping() method is in RyanAndMonicaJob. Synchronizing an instance method means “lock *this* RyanAndMonicaJob instance.” However, there are *two* instances of RyanAndMonicaJob; one is “ryan,” and the other is “monica.” If “ryan” is locked, “monica” can still make changes to the bank account; she doesn’t care that the “ryan” job is locked.

The object that needs locking, the object these two threads are fighting over, is the BankAccount. Putting synchronized on a method in RyanAndMonicaJob (and locking a RyanAndMonicaJob instance) isn’t going to solve anything.

It's important to lock the correct object

Since it's the BankAccount object that's shared, you could argue it should be the BankAccount that's in charge of making sure it is safe for multiple threads to use. The spend() method on BankAccount could make sure there's enough money *and* debit the account in a single transaction.

```

class RyanAndMonicaJob implements Runnable {
    // ...rest of the RyanAndMonicaJob class
    // the same as before...

    private void goShopping(int amount) {
        System.out.println(name + " is about to spend");
        account.spend(name, amount);
        System.out.println(name + " finishes spending");
    }
}

class BankAccount {
    // other methods in BankAccount...
    public synchronized void spend(String name, int amount) {
        if (balance >= amount) {
            balance = balance - amount;
            if (balance < 0) {
                System.out.println("Overdrawn!");
            } else {
                System.out.println("Sorry, not enough for " + name);
            }
        } else {
            Do the balance check and balance
            decrease in the BankAccount itself. If
            this method is synchronized, it becomes
            an atomic transaction that can be done
            in full by only one thread at a time.
        }
    }
}

```



This would no longer need to check the balance before spending if we know the BankAccount spend() method checks for us.

Locks the BankAccount instance the two threads are using.

Ryan and Monica SHOULDN'T go overdrawn now; this should never be the case.

there are no Dumb Questions

Q: What about protecting static variable state? If you have static methods that change the static variable state, can you still use synchronization?

A: Yes! Remember that static methods run against the class and not against an individual instance of the class. So you might wonder whose object's lock would be used on a static method? After all, there might not even be any instances of that class. Fortunately, just as each *object* has its own lock, each loaded *class* has a lock. That means if you have three Dog objects on your heap, you have a total of four Dog-related locks; three belonging to the three Dog instances, and one belonging to the Dog class itself. When you synchronize a static method, Java uses the lock of the class itself. So if you synchronize two static methods in a single class, a thread will need the class lock to enter either of the methods.

The dreaded “Lost Update” problem

Here's another classic concurrency problem. Sometimes you'll hear them called “race conditions,” where two or more threads are changing the same data at the same time. It's closely related to the Ryan and Monica story, so we'll use this example to illustrate a few more points.

The lost update revolves around one process:

Step 1: Get the balance in the account

```
int i = balance;
```

Step 2: Add 1 to that balance

```
balance = i + 1;
```

Probably not an atomic process

Even if we used the more common syntax of `balance++`, there is no guarantee that the compiled bytecode will be an “atomic process.” In fact, it probably won't—it's actually multiple operations: a read of the current value and then adding one to that value and setting it back into the original variable.

In the “Lost Update” problem, we have many threads trying to increment the balance. Take a look at the code, and then we'll look at the real problem:

```
public class LostUpdate {
    public static void main(String[] args) throws InterruptedException {
        ExecutorService pool = Executors.newFixedThreadPool(6);
        Balance balance = new Balance();
        for (int i = 0; i < 1000; i++) { ← Run 1,000 attempts to update
            pool.execute(() -> balance.increment()); ← the balance, on different threads.
        }
        pool.shutdown();
        if (pool.awaitTermination(1, TimeUnit.MINUTES)) { ← Make sure the pool has finished
            System.out.println("balance = " + balance.balance); ← running all the updates before
        } ← printing the final balance. In
    } ← theory, this should be 1,000. If
} ← it's any less than that, we've lost
an update!
```

Create a thread pool to run all the jobs. If you add more threads here, you may see even more missing updates.

```
class Balance {
    int balance = 0;
    public void increment() { ← Here's the crucial part! We increment the balance
        balance++; ← by adding 1 to whatever the value of balance was AT THE TIME WE READ IT (rather than adding 1 to whatever the CURRENT value is). You might think that “++” is an atomic operation, but it is not.
    }
}
```

Let's run this code...

① Thread A runs for a while

Reads balance: 0
 Set the value of balance to $0 + 1$.
 Now balance is 1



Reads balance: 1
 Set the value of balance to $1 + 1$.
 Now balance is 2

② Thread B runs for a while



Reads balance: 2
 Set the value of balance to $2 + 1$.
 Now balance is 3

Reads balance: 3
[now thread B is sent back to runnable, before it sets the value of balance to 4]

③ Thread A runs again, picking up where it left off

Reads balance: 3
 Set the value of balance to $3 + 1$.
 Now balance is 4



Reads balance: 4
 Set the value of balance to $4 + 1$.
 Now balance is 5

④ Thread B runs again, and picks up exactly where it left off!

Set the value of balance to $3 + 1$.
 Now balance is 4



Yikes!!
 Thread A updated it to 5, but now B came back and stepped on top of the update A made, as if A's update never happened.

**We lost the last updates that Thread A made!
 Thread B had previously done a “read” of the value of balance, and when B woke up, it just kept going as if it never missed a beat.**

Make the increment() method atomic.

Synchronize it!



Synchronizing the increment() method solves the “Lost Update” problem, because it keeps the steps in the method (read of balance and increment of balance) as one unbreakable unit.

```
public synchronized void increment() {
    balance++;
}
```

Classic concurrency gotcha: this looks like a single operation, but it's actually more than one—it's a read of the balance, an increment, and an update to the balance.

Once a thread enters the method, we have to make sure that all the steps in the method complete (as one atomic process) before any other thread can enter the method.

there are no Dumb Questions

Q: Sounds like it's a good idea to synchronize everything, just to be thread-safe.

A: Nope, it's not a good idea. Synchronization doesn't come for free. First, a synchronized method has a certain amount of overhead. In other words, when code hits a synchronized method, there's going to be a performance hit (although typically, you'd never notice it) while the matter of “is the key available?” is resolved.

Second, a synchronized method can slow your program down because synchronization restricts concurrency. In other words, a synchronized method forces other threads to get in line and wait their turn. This might not be a problem in your code, but you have to consider it.

Third, and most frightening, synchronized methods can lead to deadlock! (We'll see this in a couple of pages.)

A good rule of thumb is to synchronize only the bare minimum that should be synchronized. And in fact, you can synchronize at a granularity that's even smaller than a method. Remember, you can use the synchronized keyword to synchronize at the more fine-grained level of one or more statements, rather than at the whole-method level (we used this in our first solution to Ryan and Monica's problem).

doStuff() doesn't need to be synchronized, so we don't synchronize the whole method.

```
public void go() {
    doStuff();
    synchronized(this) {
        criticalStuff();
        moreCriticalStuff();
    }
}
```



```
synchronized(this) {
    criticalStuff();
    moreCriticalStuff();
}
```

Although there are other ways to do it, you will almost always synchronize on the current object (this). That's the same object you'd lock if the whole method were synchronized.

Now, only these two method calls are grouped into one atomic unit. When you use the synchronized keyword WITHIN a method, rather than in a method declaration, you have to provide an argument that is the object whose key the thread needs to get.

① Thread A runs for a while

Attempt to enter the increment() method.

The method is synchronized, so **get the key** for this object

Reads balance: 0

Set the value of balance to $0 + 1$.

Now balance is 1

Return the key (it completed the increment() method).

Re-enter the increment() method and **get the key**.

Reads balance: 1



[now thread A is sent back to runnable, but since it has not completed the synchronized method, Thread A keeps the key]

② Thread B is selected to run



Attempt to enter the increment() method. The method is synchronized, so we need to get the key.

The key is not available.

[now thread B is sent into an “object lock not available” lounge]

③ Thread A runs again, picking up where it left off (remember, it still has the key)

Set the value of balance $1 + 1$.

Now balance is 2

Return the key.

[now thread A is sent back to runnable, but since it has completed the increment() method, the thread does NOT hold on to the key]



④ Thread B is selected to run



Attempt to enter the increment() method. The method is synchronized, so we need to get the key.

This time, the key IS available; get the key.

Reads balance: 2

[continues to run...]

Deadlock, a deadly side of synchronization

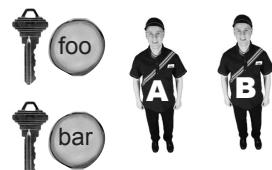
Synchronization saved Ryan and Monica from using their bank account at the same time, and has saved us from losing updates. But we also mentioned that we shouldn't synchronize everything, one reason being that synchronization can slow your program down.

There's another important consideration: we need to be careful using synchronized code, because nothing will bring your program to its knees like thread deadlock. Thread deadlock happens when you have two threads, both of which are holding a key the other thread wants. There's no way out of this scenario, so the two threads will simply sit and wait. And wait. And wait.

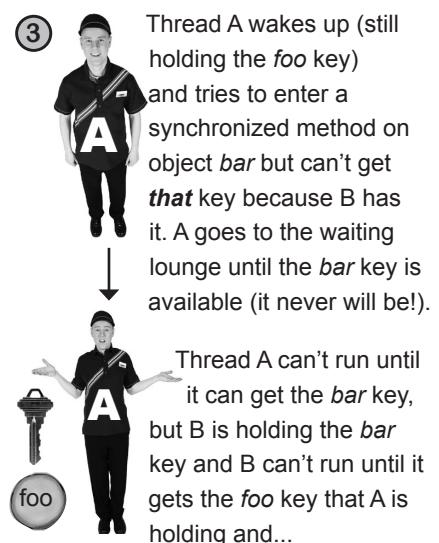
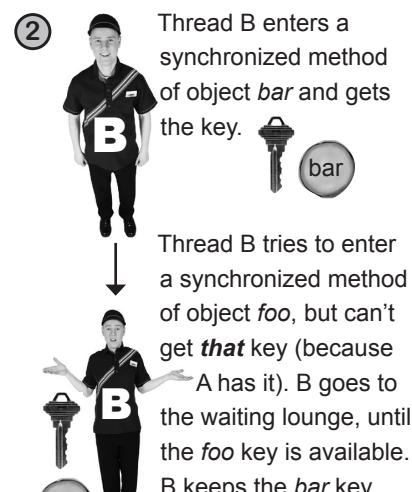
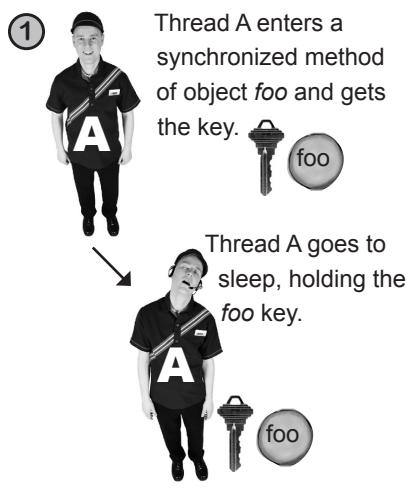
If you're familiar with databases or other application servers, you might recognize the problem; databases often have a locking mechanism somewhat like synchronization. But a real transaction management system can sometimes deal with deadlock. It might assume, for example, that deadlock might have occurred when two transactions are taking too long to complete. But unlike Java, the application server can do a "transaction rollback" that returns the state of the rolled-back transaction to where it was before the transaction (the atomic part) began.

Java has no mechanism to handle deadlock. It won't even *know* deadlock occurred. So it's up to you to design carefully. We're not going to go into more detail about deadlock than you see on this page, so if you find yourself writing multithreaded code, you might want to study *Java Concurrency in Practice* by Brian Goetz, et al. It goes into a lot of detail about the sorts of problems you can face with concurrency (like deadlock), and approaches to address these problems.

All it takes for deadlock are two objects and two threads.



A simple deadlock scenario:



You don't always have to use synchronized

Since synchronization can come with some costs (like performance and potential deadlocks), you should know about other ways to manage data that's shared between threads. The `java.util.concurrent` package has lots of classes and utilities for working with multithreaded code.

Atomic variables

If the shared data is an int, long, or boolean, we might be able to replace it with an *atomic variable*. These classes provide methods that are atomic, i.e., can safely be used by a thread without worrying about another thread changing the object's values at the same time.

There are few types of atomic variable, e.g., **AtomicInteger**, **AtomicLong**, **AtomicBoolean**, and **AtomicReference**.

We can fix our Lost Update problem with an `AtomicInteger`, instead of synchronizing the increment method.

```
class Balance {
    AtomicInteger balance = new AtomicInteger(0);
    public void increment() {
        balance.incrementAndGet();
    }
}
```

No need to add "synchronized" when you're using atomic operations.

Use an `AtomicInteger` initialized to zero, instead of an int value.

 if it's used by multiple threads, it will safely increase the value by one in a single operation. The `incrementAndGet` method returns the new, updated value, but we don't need that for our example; we're not going to use the returned value.

So I can use `AtomicInteger` as long as all I want to do is a simple increment. How does this help me if I want to do normal things like complex calculations?



Atomic variables get more interesting when you use their *compare-and-swap* (**CAS**) operations. CAS is yet another way to make an atomic change to a value. You can use CAS on atomic variables by using the **compareAndSet** method. Yes, it's a slightly different name! Gotta love programming, where naming is always the hardest problem to solve.

The `compareAndSet` method takes a value, which is what you *expect* the atomic variable to be, compares it to the *current* value, and if that matches, *then* the operation will complete.

In fact, we can use this to fix our Ryan and Monica problem, instead of locking the whole bank account with **synchronized**.

Compare-and-swap with atomic variables

How could we make use of atomic variables, and CAS (via `compareAndSet`), to solve Ryan and Monica's problem?

Since Ryan and Monica were both trying to access an int value, the account balance, we could use an `AtomicInteger` to store that balance. We could then use `compareAndSet` to update the balance when someone wants to spend money.

```
private AtomicInteger balance = new AtomicInteger(100);
```

...

```
boolean success = balance.compareAndSet(expectedValue, newValue)
```

True if the balance was updated to the new value. If this is false, the balance wasn't changed and YOU decide what you need to do next.

This is the value you THINK the balance is.

This is the value you want the balance to have.

In plain English:

“Set the balance to this new value only if the current balance is the same as this expected value, and tell me if the balance was actually changed.”

Compare-and-swap uses *optimistic locking*. Optimistic locking means you don't stop all threads from getting to the object; you *try* to make the change, but you embrace the fact that the change **might not happen**. If it doesn't succeed, you decide what to do. You might decide to try again, or to send a message letting the user know it didn't work.

This may be more work than simply locking all other threads out from the object, but it can be faster than locking everything. For example, when the chances of multiple writes happening at the same time are very low or if you have a lot of threads reading and not so many writing, then you may not want to pay the price of a lock on every write.

When you're using CAS operations, you have to deal with the times when the operation does NOT succeed.

Ryan and Monica, going atomic

Let's see the whole thing in action in Ryan and Monica's bank account. We'll put the balance in an AtomicInteger and use compareAndSet to make an *atomic* change to the balance.

```

import java.util.concurrent.atomic.AtomicInteger;
class BankAccount {
    private final AtomicInteger balance = new AtomicInteger(100);

    public int getBalance() {
        return balance.get();
    }
    Not synchronized
    public void spend(String name, int amount) {
        int initialBalance = balance.get(); } Like before, check if there's enough money. This
        if (initialBalance >= amount) { time, keep a record of the balance.

        boolean success = balance.compareAndSet(initialBalance, initialBalance - amount);
        The balance will NOT be changed if
        the initial balance does not match
        the actual balance right now.
        Pass in the balance from
        when we checked if there
        was enough money.
        If success was false, the
        money was NOT spent.
        Tell Ryan or Monica it
        didn't work and they can
        decide what to do.
    }
}

```

```

File Edit Window Help SorryMonica
% java RyanAndMonicaTest
Ryan is about to spend
Monica is about to spend
Ryan finishes spending
Sorry Monica, you can't buy this
Monica finishes spending

```

Monica was able to start her shopping, but by the time she came to pay, the bank said no. At least they didn't go overdrawn!

`java.util.concurrent` has lots of useful classes and utilities for working with multithreaded code. Take a look at what's there!



So if all these problems are caused by writing to a shared object, what if we stopped threads from changing the data in the shared objects? Is there a way to do that?

Make an object immutable if you're going to share it between threads and you don't want the threads to change its data.

The very best way to know *for sure* that another thread isn't changing your data is to make it impossible to change the data in the object. When an object's data cannot be changed, we call it an **immutable object**.

Writing a class for immutable data

All fields should be FINAL. The value will be set once, in the field declaration or constructor, and cannot be changed afterward.

```
public final class ImmutableData {
    private final String name;
    private final int value;

    public ImmutableData(String name, int value) {
        this.name = name;
        this.value = value;
    }

    public String getName() { return name; }

    public int getValue() { return value; }
}
```

We don't want to allow subclasses that might add mutable values, so make this immutable class final.

All fields need to be initialized once, usually in the constructor.

Immutable objects may have getters, but no setters. The values inside the object should not be changed in any method.



Brain Barbell

There are times when adding the final keyword isn't enough to prevent changes. When do you think that might be the case? We'll give you a clue....



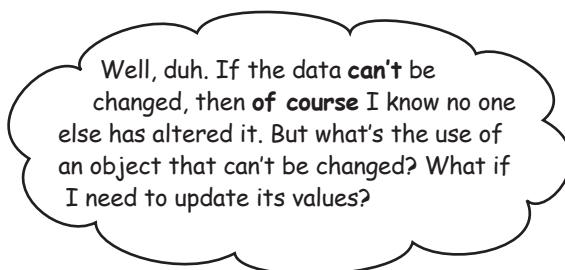
Using immutable objects

It is terribly convenient to be able to change data on a shared object and assume that all the other threads will be able to see these changes.

However, we've also seen that while it's *convenient*, it's not very *safe*.

On the other hand, when a thread is working with an object that cannot be changed, it can make assumptions about the data in that object; e.g., once the thread has read a value from the object, it knows that data can't change.

We don't need to use synchronization or other mechanisms to control who changes the data because it can't change.



Working with immutable objects means thinking in a different way.

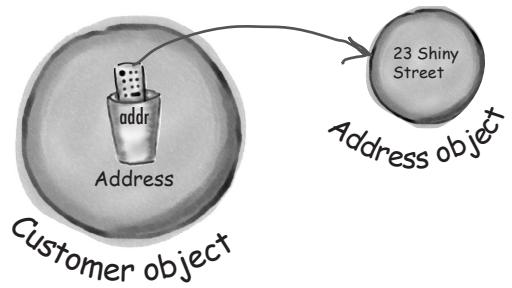
Instead of making changes to the *same* object, we *replace* the old object with a new one. The new object has the updated values, and any threads that need the new values need to use the new object.

What happens to the old object? Well, if it's still being used by something (and it might be—it's perfectly valid sometimes to work with older data), it will hang around on the heap. If it's not being used, it'll be garbage collected, and we don't have to worry about it anymore.

Changing immutable data

Imagine a system that has customers, and that each Customer object has an Address that represents the street address of a customer. If the customer's Address is an immutable object (all its fields are final and the data cannot be changed), how do you change the customer's address when they move?

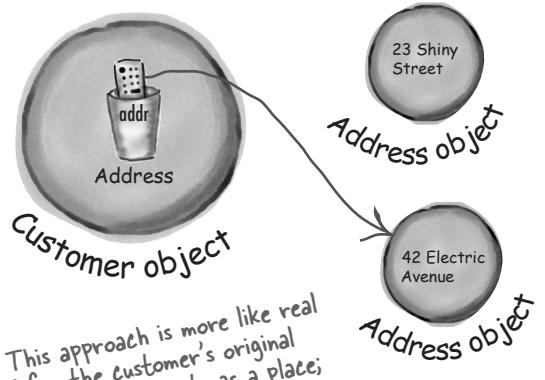
- 1 The Customer has a reference to the original Address object containing the customer's street address data.



- 2 When the customer moves, a brand new Address object is created with the new street address for the customer.



- 3 The Customer object's reference to their address is changed to point to the new Address object.



there are no
Dumb Questions

Q: What happens if other parts of the program have a reference to the old Address?

A: Actually sometimes we want this. Imagine the customer placed an order to be delivered to their original address. We still want the details of that order to have the original address; we don't want it changed to contain the details of the new address.

Once the customer changes their address (and the Customer contains a reference to the new address object), *then* we want new orders to use the new Address object.

This approach is more like real life—the customer's original address still exists as a place; it's just not the place that our customer lives at anymore.



- o o

Wait just a minute! The `Address` object is immutable and doesn't change, but the `Customer` object still has to change.

Absolutely right. If your system has data that changes, those changes do have to happen somewhere. The key idea to take away from this discussion is that not all of the classes in your application have to have data that changes. In fact, we'd argue for minimizing the places where things change. Then, there are far fewer places where you have to think about what happens if multiple threads are making changes at the same time.

There are a number of techniques for working effectively with immutable data classes; we've just scratched the surface here. It is interesting to note that Java 16 introduced *records*, which are immutable data classes provided directly by the language.

Use immutable data classes where you can.
Limit the number of places where data can be changed by multiple threads.

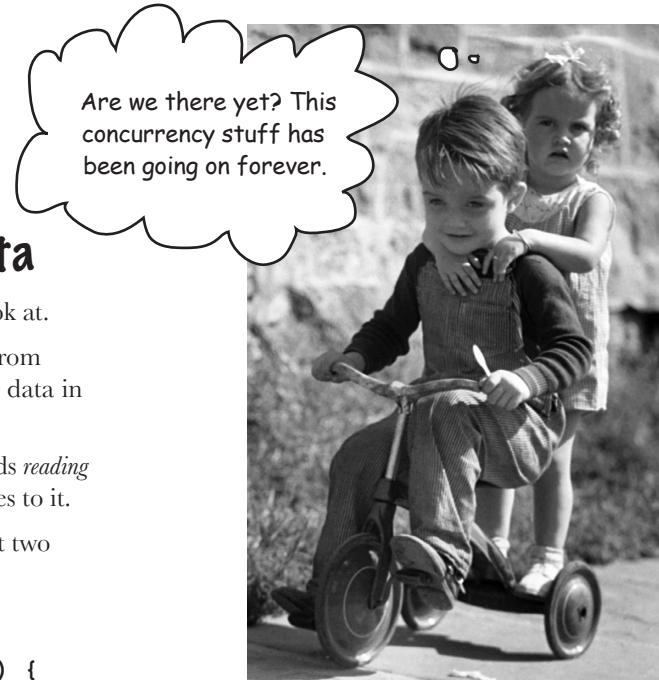
More problems with shared data

We're nearly there, we promise! Just one last thing to look at.

So far we've seen all sorts of problems that can come from many threads writing to the same data. This applies to data in Collections too.

We can even have problems when we have lots of threads *reading* the same data, even if only one thread is making changes to it.

This code has just one thread writing to a collection, but two threads reading it.



```
public class ConcurrentReaders {
    public static void main(String[] args) {
        List<Chat> chatHistory = new ArrayList<>(); ← Stores the Chat
        ExecutorService executor = Executors.newFixedThreadPool(3); objects in an ArrayList,
        for (int i = 0; i < 5; i++) { which is NOT thread
            executor.execute(() -> chatHistory.add(new Chat("Hi there!")));
            executor.execute(() -> System.out.println(chatHistory));
            executor.execute(() -> System.out.println(chatHistory));
        }
        executor.shutdown();
    }

    final class Chat {
        private final String message;
        private final LocalDateTime timestamp; ← Create a writing thread
        public Chat(String message) { that adds to the List, and
            this.message = message;
            timestamp = LocalDateTime.now(); two threads that read
        }
        ← Instances of Chat are immutable.
        public String toString() {
            String time = timestamp.format(ofLocalizedTime(MEDIUM));
            return time + " " + message;
        }
    }
}
```

Making an Object field "final" doesn't guarantee that the data inside that object won't change, just that the reference won't change. String and LocalDateTime are immutable, so this is safe.

Reading from a changing data structure causes an Exception

Running the code on the last page causes an Exception to be thrown, sometimes. By now you know these sorts of issues depend a lot on the whims of the hardware, the operating system, and the JVM.

```
File Edit Window Help PoorBrunonono
% java ConcurrentReaders
[]
[]
[18:43:59 Hi there!, 18:43:59 Hi there!]
[18:43:59 Hi there!, 18:43:59 Hi there!]
[18:43:59 Hi there!, 18:43:59 Hi there!, 18:43:59 Hi there!]
[18:43:59 Hi there!, 18:43:59 Hi there!, 18:43:59 Hi there!]
Exception in thread "pool-1-thread-2" Exception in thread "pool-1-thread-1" java.util.
ConcurrentModificationException
    at java.base/java.util.ArrayList$Itr.checkForComodification(ArrayList.java:1043)
    at java.base/java.util.ArrayList$Itr.next(ArrayList.java:997)
    at java.base/java.util.AbstractCollection.toString(AbstractCollection.java:472)
    at java.base/java.lang.String.valueOf(String.java:2951)
    at java.base/java.io.PrintStream.println(PrintStream.java:897)
    at ConcurrentReaders.lambda$main$2(ConcurrentReaders.java:17)
```

A ConcurrentModificationException is thrown by the reading thread when the List it is reading is changed WHILE this thread is reading it.

If a collection is changed by one thread while another thread is reading that collection, you can get a ConcurrentModificationException.

Use a thread-safe data structure

Clearly good ol' ArrayList just isn't going to cut it if you have threads reading data that's being changed at the same time. Luckily for us, there are other options. We want a thread-safe data structure, one that can be written to, and read from, by multiple threads at the same time.

The java.util.concurrent package has a number of thread-safe data structures, and we're going to look at **CopyOnWriteArrayList** to solve this specific problem.

CopyOnWriteArrayList is a reasonable choice when you have a List that is being **read a lot, but not changed very often**. We'll see why later.

```
public class ConcurrentReaders {
    public static void main(String[] args) {
        List<Chat> chatHistory = new CopyOnWriteArrayList<>();
        ExecutorService executor = Executors.newFixedThreadPool(3);
        for (int i = 0; i < 5; i++) {
            executor.execute(() -> chatHistory.add(new Chat("Hi there!")));
            executor.execute(() -> System.out.println(chatHistory));
            executor.execute(() -> System.out.println(chatHistory));
        }
        executor.shutdown();
    }
}
```

CopyOnWriteArrayList implements the List interface, so we can use it as a drop-in replacement for any List.

The rest of the code is exactly the same as before.

```
File Edit Window Help AyMariposa
% java ConcurrentReaders
[]
[]
[]
[]
[10:26:22 Hi there!, 10:26:22 Hi there!]
[10:26:22 Hi there!, 10:26:22 Hi there!, 10:26:22 Hi there!, 10:26:22 Hi there!]
[10:26:22 Hi there!, 10:26:22 Hi there!, 10:26:22 Hi there!, 10:26:22 Hi there!]
[10:26:22 Hi there!, 10:26:22 Hi there!, 10:26:22 Hi there!, 10:26:22 Hi there!, 10:26:22 Hi there!]
[10:26:22 Hi there!, 10:26:22 Hi there!, 10:26:22 Hi there!, 10:26:22 Hi there!, 10:26:22 Hi there!]

Process finished with exit code 0
```

No Exception!

CopyOnWriteArrayList

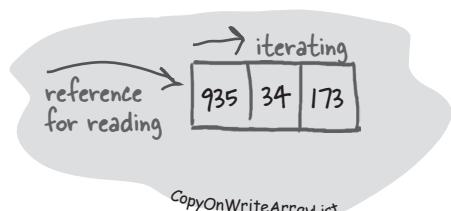
CopyOnWriteArrayList uses immutability to provide safe access for reading threads while other threads are writing.

How does it work? Well, it does what it says on the tin: when a thread is writing to the list, it's actually writing to a *copy* of the list. When the changes have been made, then the new copy replaces the original. In the meantime, any threads that were reading the list before the change are happily (and safely!) reading the original.

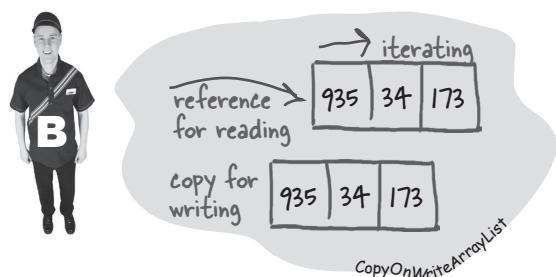
- 1 An instance of `CopyOnWriteArrayList` contains an ordered set of data, like an array.



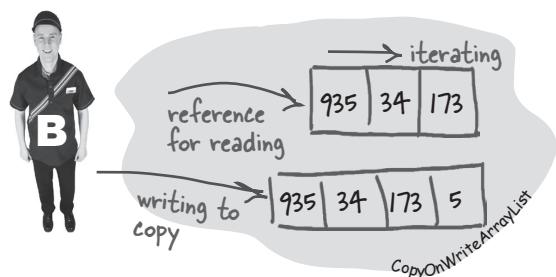
- 2 When Thread A reads the `CopyOnWriteArrayList`, it gets an Iterator that allows it to read a **snapshot** of the list data at that point in time.



- 3 Thread B writes data to the `CopyOnWriteArrayList` by adding a new element, and the `CopyOnWriteArrayList` creates a **copy** of the list data before any changes are made. This is invisible to any of the reading or writing threads.

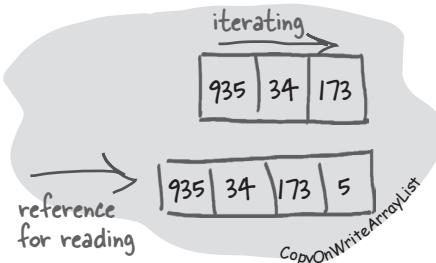


- 4 When Thread B makes changes to "the list," it's actually making changes to this copy. It's happy knowing the changes are being made. The reading threads like Thread A are not affected at all; they're iterating over the snapshot of the original data.



- 5 Once Thread B has finished its updates, then the original data is replaced with the new data.

If Thread A is still reading, it's safely reading the old data. If any other threads start reading after the change, they get the new data.



there are no Dumb Questions

Q: Doesn't using `CopyOnWriteArrayList` mean that some reading threads will be reading old data?

A: Yes, the reading threads will always be working off data that is a snapshot of the data from when they first started reading. This means that potentially the data might be out of date at some point, but at least it's not going to throw a `ConcurrentModificationException`.

Q: Isn't it bad to be using out of date data?

A: Not necessarily. In many systems this is "good enough." Think, for example, about a website that shows live news. Yes, you want it to be pretty up to date, but it doesn't have to be up to date to the latest millisecond; it's probably fine if some news is a few seconds old.

Q: But I don't want my bank statement to be even slightly out of date! How can I make sure that critical shared data is always correct?

A: `CopyOnWriteArrayList` is probably not the right choice if all threads need to be working off exactly the same data. Other data structures, like `Vector`, provide thread-safety by using locks to ensure only one thread at a time has access to the data. This is safe, but can be slow—you're not getting any benefits of multi-threading if your threads need to wait their turn to get to their data.

Q: So `CopyOnWriteArrayList` is a fast thread-safe data structure?

A: Well.. it depends! It's fast (compared to a locking Collection) if you have lots of reading threads and not many writes. But if there's a lot of writing going on, it might not be the best data structure. The cost of creating a new copy of the data every time a single write is made might be too high for some applications.

Q: Why isn't there an easy answer to the best way to do this concurrency stuff?

A: Concurrent programming is all about trade-offs. You need a good understanding of what your application is doing, how it should work, and the hardware and environment that's running it. If you find yourself wondering which approach is better for your application, it's probably a good time to learn about performance testing so you can measure exactly how each approach impacts the performance in your system.

Q: You never told us about the case where adding `final` to a field declaration is not enough to make sure that value is never changed. What's the deal?

A: Good catch! The "deal" is that if your field is a reference to another object, like a `Collection` or one of your own objects, using `final` does not prevent another thread from changing the values inside that object. The only way to make sure that doesn't happen is to make sure all your fields that are references refer only to immutable objects themselves. Otherwise, your immutable object can have data that changes.

See the `LocalDateTime` case on page 662.

Thread-safe collections in early versions of Java were made safe via locking. For example, `java.util.Vector`.

Java 5 introduced concurrent data structures in `java.util.concurrent`. These do NOT use locking.

BULLET POINTS

- You can have serious problems with threads if two or more threads are trying to change the same data.
- Two or more threads accessing the same object can lead to data corruption if one thread, for example, leaves the running state while still in the middle of manipulating an object's critical state.
- To make your objects thread-safe, decide which statements should be treated as one atomic process. In other words, decide which methods must run to completion before another thread enters the same method on the same object.
- Use the keyword **synchronized** to modify a method declaration when you want to prevent two threads from entering that method.
- Every object has a single lock, with a single key for that lock. Most of the time we don't care about that lock; locks come into play only when an object has synchronized methods or use the synchronized keyword with a specified object.
- When a thread attempts to enter a synchronized method, the thread must get the key for the object (the object whose method the thread is trying to run). If the key is not available (because another thread already has it), the thread goes into a kind of waiting lounge until the key becomes available.
- Even if an object has more than one synchronized method, there is still only one key. Once any thread has entered a synchronized method on that object, no thread can enter any other synchronized method on the same object. This restriction lets you protect your data by synchronizing any method that manipulates the data.
- The synchronized keyword isn't the only way to safeguard your data from multi-threaded changes. Atomic variables, with CAS operations, may be suitable if it's just one value that is being changed by multiple threads.
- It's *writing* data from multiple threads that causes the most problems, not *reading*, so consider if your data needs to be changed at all or if it can be immutable.
- Make a class immutable by making the class final, making all the fields final, setting the values just once in the constructor or field declaration, and having no setters or other methods that can change the data.
- Having immutable objects in your application doesn't mean nothing ever changes; it means that you limit the parts of your application where you have to worry about multiple threads changing the data.
- There are thread-safe data structures that let you have multiple threads reading the data while one (or more) threads change the data. Some of these are in `java.util.concurrent`.
- Concurrent programming is difficult! But there are plenty of tools to help you.



BE the JVM



The Java file on this page represents a complete source file. Your job is to play JVM and determine what the output would be when the program runs.

How might you fix it, making sure the output is correct every time?

→ Answers on page 670.

```

import java.util.*;
import java.util.concurrent.*;

public class TwoThreadsWriting {
    public static void main(String[] args) {
        ExecutorService threadPool = Executors.newFixedThreadPool(2);
        Data data = new Data();
        threadPool.execute(() -> addLetterToData('a', data));
        threadPool.execute(() -> addLetterToData('A', data));
        threadPool.shutdown();
    }

    private static void addLetterToData(char letter, Data data) {
        for (int i = 0; i < 26; i++) {
            data.addLetter(letter++);
            try {
                Thread.sleep(50);
            } catch (InterruptedException ignored) {}
        }
        System.out.println(Thread.currentThread().getName() + data.getLetters());
        System.out.println(Thread.currentThread().getName()
                           + " size = " + data.getLetters().size());
    }
}

final class Data {
    private final List<String> letters = new ArrayList<>();

    public List<String> getLetters() {return letters;}

    public void addLetter(char letter) {
        letters.add(String.valueOf(letter));
    }
}

```



Near-miss at the airlock

As Sarah joined the onboard development team's design review meeting, she gazed out the portal at sunrise over the Indian Ocean. Even though the ship's conference room was incredibly claustrophobic, the sight of the growing blue and white crescent overtaking night on the planet below filled Sarah with awe and appreciation.

Five-Minute Mystery



This morning's meeting was focused on the control systems for the orbiter's airlocks. As the final construction phases were nearing their end, the number of spacewalks was scheduled to increase dramatically, and traffic was high both in and out of the ship's airlocks. "Good morning, Sarah," said Tom, "Your timing is perfect; we're just starting the detailed design review."

"As you all know," said Tom, "Each airlock is outfitted with space-hardened GUI terminals, both inside and out. Whenever spacewalkers are entering or exiting the orbiter they will use these terminals to initiate the airlock sequences." Sarah nodded and asked, "Tom, can you tell us what the method sequences are for entry and exit?" Tom rose and floated to the whiteboard, "First, here's the exit sequence method's pseudocode." Tom quickly wrote on the board.

```
orbiterAirlockExitSequence ()  
    verifyPortalStatus ();  
    pressurizeAirlock ();  
    openInnerHatch ();  
    confirmAirlockOccupied ();  
    closeInnerHatch ();  
    decompressAirlock ();  
    openOuterHatch ();  
    confirmAirlockVacated ();  
    closeOuterHatch ();
```

"To ensure that the sequence is not interrupted, we have synchronized all of the methods called by the orbiterAirlockExitSequence() method," Tom explained. "We'd hate to see a returning spacewalker inadvertently catch a buddy with his space pants down!"

Everyone chuckled as Tom erased the whiteboard, but something didn't feel right to Sarah, and it finally clicked as Tom began to write the entry sequence pseudocode on the whiteboard. "Wait a minute, Tom!" cried Sarah, "I think we've got a big flaw in the exit sequence design. Let's go back and revisit it; it could be critical!"

Why did Sarah stop the meeting? What did she suspect?

→ Answers on page 671.



Exercise Solutions

BE the JVM (from page 668)

The answer is the output won't be the same every time. In theory, one might expect the size to always be 52 (2×26 letters in the alphabet), but in fact this is one of those lost-update problems.

```

File Edit Window Help ZeCount
% java TwoThreadsWriting
pool-1-thread-2[a, A, b, B, c, C, d, D, E, F, g, G, h, H, i, j, K, k,
pool-1-thread-1[a, A, b, B, c, C, d, D, E, F, g, G, h, H, i, j, K, k,
pool-1-thread-1 size = 40
pool-1-thread-2 size = 40

```

Example output. Your output probably won't look exactly the same as this, but if you predicted that the size would be less than 52, you win a cookie.

It can be solved in two different ways; both are valid.

Synchronize the write method

```

public synchronized void addLetter(char letter) {
    letters.add(String.valueOf(letter));
}

```

If this method is synchronized, only one thread at a time can write to the data, and therefore no updates will be lost. This will not work if there's a DIFFERENT thread reading at the same time as one of these threads are writing.

Use a thread-safe collection

```
private final List<String> letters = new CopyOnWriteArrayList<>();
```

Using `CopyOnWriteArrayList` will allow the threads to both safely write to the letters List.

Use either solution, you do NOT have to do both!!

With a thread-safe collection, you don't have to synchronize the writing method.



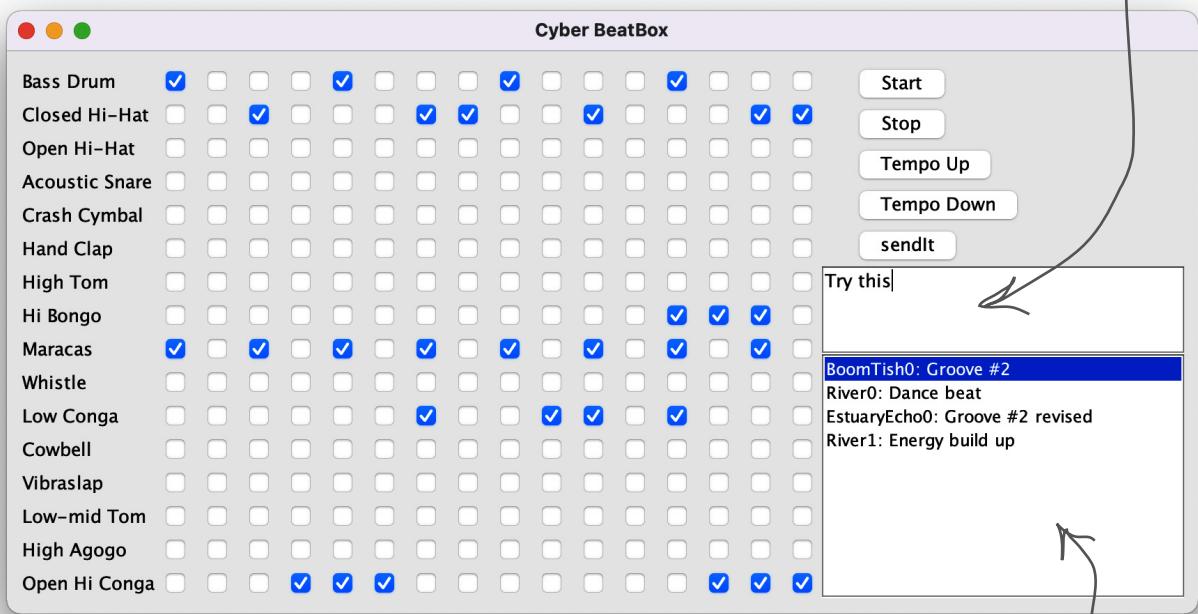
Five-Minute Mystery (from page 669)

What did Sarah know?

Sarah realized that in order to ensure that the entire exit sequence would run without interruption the `orbiterAirlockExitSequence()` method needed to be synchronized. As the design stood, it would be possible for a returning spacewalker to interrupt the Exit Sequence! The Exit Sequence thread couldn't be interrupted in the middle of any of the lower-level method calls, but it could be interrupted in *between* those calls. Sarah knew that the entire sequence should be run as one atomic unit, and if the `orbiterAirlockExitSequence ()` method was synchronized, it could not be interrupted at any point.

Appendix A: Final Code Kitchen

Your message gets sent to the other players, along with your current beat pattern, when you hit "sendit."



Finally, the complete version of the BeatBox!

It connects to a simple MusicServer so that you can send and receive beat patterns with other clients.

Incoming messages from players. Click one to load the pattern that goes with it, and then click 'Start' to play it.

Final BeatBox client program

Most of this code is the same as the code from the Code Kitchens in the previous chapters, so we don't annotate the whole thing again. The new parts include:

GUI: Two new components are added for the text area that displays incoming messages (actually a scrolling list) and the text field.

NETWORKING: Just like the SimpleChatClient in this chapter, the BeatBox now connects to the server and gets an input and output stream.

MULTITHREADED: Again, just like the SimpleChatClient, we start a "reader" job that keeps looking for incoming messages from the server. But instead of just text, the messages coming in include TWO objects: the String message and the serialized array (the thing that holds the state of all the checkboxes).

All the code is available at https://oreil.ly/hfJava_3e_examples.

```
import javax.sound.midi.*;
import javax.swing.*;
import javax.swing.event.*;
import java.awt.*;
import java.io.*;
import java.net.Socket;
import java.util.*;
import java.util.concurrent.*;

import static javax.sound.midi.ShortMessage.*;

public class BeatBoxFinal {
    private JList<String> incomingList;
    private JTextArea userMessage;
    private ArrayList<JCheckBox> checkboxList;

    private Vector<String> listVector = new Vector<>();
    private HashMap<String, boolean[]> otherSeqsMap = new HashMap<>();

    private String userName;
    private int nextNum;

    private ObjectOutputStream out;
    private ObjectInputStream in;

    private Sequencer sequencer;
    private Sequence sequence;
    private Track track;

    String[] instrumentNames = {"Bass Drum", "Closed Hi-Hat",
        "Open Hi-Hat", "Acoustic Snare", "Crash Cymbal", "Hand Clap",
        "High Tom", "Hi Bongo", "Maracas", "Whistle", "Low Conga",
        "Cowbell", "Vibraslap", "Low-mid Tom", "High Agogo",
        "Open Hi Conga"};
    int[] instruments = {35, 42, 46, 38, 49, 39, 50, 60, 70, 72, 64, 56, 58, 47, 67, 63};
```

These are the names of the instruments, as a String array, for building the GUI labels (on each row).

These represent the actual drum "keys." The drum channel is like a piano, except each "key" on the piano is a different drum. So the number '35' is the key for the Bass drum, 42 is Closed Hi-Hat, etc.

```
public static void main(String[] args) {
    new BeatBoxFinal().startUp(args[0]); Add a command-line argument for your screen name.
}
```

Example: % java BeatBoxFinal theFlash

```
public void startUp(String name) {
    userName = name;
    // open connection to the server
    try {
        Socket socket = new Socket("127.0.0.1", 4242);
        out = new ObjectOutputStream(socket.getOutputStream());
        in = new ObjectInputStream(socket.getInputStream());
        ExecutorService executor = Executors.newSingleThreadExecutor();
        executor.submit(new RemoteReader());
    } catch (Exception ex) {
        System.out.println("Couldn't connect - you'll have to play alone.");
    }
    setUpMidi();
    buildGUI();
}
```

Set up the networking, I/O, and make (and start) the reader thread. We're using Sockets here instead of Channels because they work better with Object Input/Output streams.

```
public void buildGUI() {
    JFrame frame = new JFrame("Cyber BeatBox");
    frame.setDefaultCloseOperation(JFrame.EXIT_ON_CLOSE);
    BorderLayout layout = new BorderLayout();
    JPanel background = new JPanel(layout);
    background.setBorder(BorderFactory.createEmptyBorder(10, 10, 10, 10));
    Box buttonBox = new Box(BoxLayout.Y_AXIS);
    JButton start = new JButton("Start");
    start.addActionListener(e -> buildTrackAndStart());
    buttonBox.add(start);
```

You've seen this GUI code before, in Chapter 15.

An "empty border" gives us a margin between the edges of the panel and where the components are placed. (Purely aesthetic.)

```
JButton stop = new JButton("Stop");
stop.addActionListener(e -> sequencer.stop());
buttonBox.add(stop);
```

Lambda expressions call a specific method on this class when the button is pressed.

```
JButton upTempo = new JButton("Tempo Up");
upTempo.addActionListener(e -> changeTempo(1.03f));
buttonBox.add(upTempo);
```

The default tempo is 1.0, so we're adjusting +/- 3% per click.

```
JButton downTempo = new JButton("Tempo Down");
downTempo.addActionListener(e -> changeTempo(0.97f));
buttonBox.add(downTempo);
```

```
JButton sendIt = new JButton("sendIt");
sendIt.addActionListener(e -> sendMessageAndTracks());
buttonBox.add(sendIt);
```

This is new; send the message and the current beat sequence to the music server.

```
userMessage = new JTextArea();
userMessage.setLineWrap(true);
userMessage.setWrapStyleWord(true);
JScrollPane messageScroller = new JScrollPane(userMessage);
buttonBox.add(messageScroller);
```

Create a text area for the user to type their message.

final BeatBox code

We saw `JList` briefly in Chapter 15. This is where the incoming messages are displayed. Only instead of a normal chat where you just LOOK at the messages, in this app you can SELECT a message from the list to load and play the attached beat pattern.

```
incomingList = new JList<>();
incomingList.addListSelectionListener(new myListSelectionListener());
incomingList.setSelectionMode(ListSelectionModel.SINGLE_SELECTION);
JScrollPane theList = new JScrollPane(incomingList);
buttonBox.add(theList);
incomingList.setListData(listVector);
```

```
Box nameBox = new Box(BoxLayout.Y_AXIS);
for (String instrumentName : instrumentNames) {
    JLabel instrumentLabel = new JLabel(instrumentName);
    instrumentLabel.setBorder(BorderFactory.createEmptyBorder(4, 1, 4, 1));
    nameBox.add(instrumentLabel);
}
```

This border on each instrument name helps them line up with the checkboxes.

```
background.add(BorderLayout.EAST, buttonBox);
background.add(BorderLayout.WEST, nameBox);
```

```
frame.getContentPane().add(background);
GridLayout grid = new GridLayout(16, 16);
grid.setVgap(1);
grid.setHgap(2);
```

This layout manager one lets you put the components in a grid with rows and columns.

```
JPanel mainPanel = new JPanel(grid);
background.add(BorderLayout.CENTER, mainPanel);
```

```
checkboxList = new ArrayList<>();
for (int i = 0; i < 256; i++) {
    JCheckBox c = new JCheckBox();
    c.setSelected(false);
    checkboxList.add(c);
    mainPanel.add(c);
}
```

Make the check boxes, set them to "false" (so they aren't checked), and add them to the `ArrayList` AND to the GUI panel.

```
frame.setBounds(50, 50, 300, 300);
frame.pack();
frame.setVisible(true);
}
```

```
private void setUpMidi() {
    try {
        sequencer = MidiSystem.getSequencer();
        sequencer.open();
        sequence = new Sequence(Sequence.PPQ, 4);
        track = sequence.createTrack();
        sequencer.setTempoInBPM(120);
    } catch (Exception e) {
        e.printStackTrace();
    }
}
```

Get the Sequencer, make a Sequence, and make a Track.

```

private void buildTrackAndStart() {
    ArrayList<Integer> trackList; // this will hold the instruments for each
    sequence.deleteTrack(track);
    track = sequence.createTrack();
    for (int i = 0; i < 16; i++) {
        trackList = new ArrayList<>();
        int key = instruments[i];
        for (int j = 0; j < 16; j++) {
            JCheckBox jc = checkboxList.get(j + (16 * i));
            if (jc.isSelected()) {
                trackList.add(key);
            } else {
                trackList.add(null); // because this slot should be empty in the track
            }
        }
        makeTracks(trackList);
        track.add(makeEvent(CONTROL_CHANGE, 1, 127, 0, 16));
    }
    track.add(makeEvent(PROGRAM_CHANGE, 9, 1, 0, 15)); // - so we always go to 16 beats
}

```

```

try {
    sequencer.setSequence(sequence);
    sequencer.setLoopCount(sequencer.LOOP_CONTINUOUSLY);
    sequencer.setTempoInBPM(120);
    sequencer.start();
} catch (Exception e) {
    e.printStackTrace();
}
}

```

```

private void changeTempo(float tempoMultiplier) {
    float tempoFactor = sequencer.getTempoFactor();
    sequencer.setTempoFactor(tempoFactor * tempoMultiplier);
}

```

```

private void sendMessageAndTracks() {
    boolean[] checkboxState = new boolean[256];
    for (int i = 0; i < 256; i++) {
        JCheckBox check = checkboxList.get(i);
        if (check.isSelected()) {
            checkboxState[i] = true;
        }
    }
    try {
        out.writeObject(userName + nextNum++ + ":" + userMessage.getText());
        out.writeObject(checkboxState);
    } catch (IOException e) {
        System.out.println("Terribly sorry. Could not send it to the server.");
        e.printStackTrace();
    }
    userMessage.setText("");
}

```

This is new...it's a lot like the SimpleChatClient, except instead of sending a String message, we serialize two objects (the String message and the beat pattern) and write those two objects to the socket output stream (to the server).

Build a track by walking through the checkboxes to get their state and mapping that to an instrument (and making the MidiEvent for it). This is pretty complex, but it is EXACTLY as it was in the previous chapters, so refer to the Code Kitchen in Chapter 15 to get the full explanation again.

final BeatBox code

```
public class MyListSelectionListener implements ListSelectionListener {
    public void valueChanged(ListSelectionEvent lse) {
        if (!lse.getValueIsAdjusting()) {
            String selected = incomingList.getSelectedValue();
            if (selected != null) {
                // now go to the map, and change the sequence
                boolean[] selectedState = otherSeqsMap.get(selected);
                changeSequence(selectedState);
                sequencer.stop();
                buildTrackAndStart();
            }
        }
    }
}

private void changeSequence(boolean[] checkboxState) {
    for (int i = 0; i < 256; i++) {
        JCheckBox check = checkboxList.get(i);
        check.setSelected(checkboxState[i]);
    }
}

public void makeTracks(ArrayList<Integer> list) {
    for (int i = 0; i < list.size(); i++) {
        Integer instrumentKey = list.get(i);
        if (instrumentKey != null) {
            track.add(makeEvent(NOTE_ON, 9, instrumentKey, 100, i));
            track.add(makeEvent(NOTE_OFF, 9, instrumentKey, 100, i + 1));
        }
    }
}

public static MidiEvent makeEvent(int cmd, int chnl, int one, int two, int tick) {
    MidiEvent event = null;
    try {
        ShortMessage msg = new ShortMessage();
        msg.setMessage(cmd, chnl, one, two);
        event = new MidiEvent(msg, tick);
    } catch (Exception e) {
        e.printStackTrace();
    }
    return event;
}
```

This is also new—a ListSelectionListener that tells us when the user made a selection on the list of messages. When the user selects a message, we IMMEDIATELY load the associated beat pattern (it's in the HashMap called otherSeqsMap) and start playing it. There's some if tests because of little quirky things about getting ListSelectionEvents.

This method is called when the user selects something from the list. We IMMEDIATELY change the pattern to the one they selected.

All the MIDI stuff is exactly the same as it was in the previous versions.

```

public class RemoteReader implements Runnable {
    public void run() {
        try {
            Object obj;
            while ((obj = in.readObject()) != null) {
                System.out.println("got an object from server");
                System.out.println(obj.getClass());

                String nameToShow = (String) obj;
                boolean[] checkboxState = (boolean[]) in.readObject();
                otherSeqsMap.put(nameToShow, checkboxState);

                listVector.add(nameToShow);
                incomingList.setListData(listVector);
            }
        } catch (IOException | ClassNotFoundException e) {
            e.printStackTrace();
        }
    }
}

```

Multi-catch: if you want to catch two different Exceptions but do the same thing with them (like here, we just want to print them out), you can separate the two Exception classes with a pipe.

This is the thread job—it reads in data from the server. In this code, “data” will always be two serialized objects: the String message and the beat pattern (a boolean array of checkbox state values).

When a message comes in, we read (deserialize) the two objects (the message and the array of boolean checkbox state values), which we want to add to the JList component.

Adding to a JList is a two-step thing: you keep a Vector of the lists data (Vector is an old-fashioned ArrayList), and then tell the JList to use that Vector as its source for what to display in the list.



Sharpen your pencil

→ Yours to solve.

What are some of the ways you can improve this program?

Here are a few ideas to get you started:

1. Once you select a pattern, whatever current pattern was playing is blown away. If that was a new pattern you were working on (or a modification of another one), you're out of luck. You might want to pop up a dialog box that asks the user if he'd like to save the current pattern.

2. If you fail to type in a command-line argument, you just get an exception when you run it! Put something in the main method that checks to see if you've passed in a command-line argument. If the user doesn't supply one, either pick a default or print out a message that says they need to run it again, but this time with an argument for their screen name.

3. It might be nice to have a feature where you can click a button and it will generate a random pattern for you. You might hit on one you really like. Better yet, have another feature that lets you load in existing "foundation" patterns, like one for jazz, rock, reggae, etc., that the user can add to.

You don't have to type all the code in! You can clone it from the repository at All the code is available at https://oreil.ly/hfJava_3e_examples.

There's also an alternative BeatBox solution, which uses Maps and Lists instead of the arrays used in this solution. There's more than one way to solve any problem!

Final BeatBox server program

Most of this code is identical to the SimpleChatServer we made in Chapter 17, *Make a Connection*. The only difference, in fact, is that this server receives, and then re-sends, two serialized objects instead of a plain String (although one of the serialized objects happens to *be* a String).

```

import java.io.*;
import java.net.*;
import java.util.*;
import java.util.concurrent.*;

public class MusicServer {
    final List<ObjectOutputStream> clientOutputStreams = new ArrayList<>();

    public static void main(String[] args) {
        new MusicServer().go();
    }

    public void go() {
        try {
            ServerSocket serverSock = new ServerSocket(4242);           Open a server socket at port 4242.
            ExecutorService threadPool = Executors.newCachedThreadPool();

            while (!serverSock.isClosed()) {
                Socket clientSocket = serverSock.accept();               Keep listening for client
                ObjectOutputStream out = new ObjectOutputStream(clientSocket.getOutputStream());   connections; create a
                clientOutputStreams.add(out);                                         new Socket and new
                ClientHandler clientHandler = new ClientHandler(clientSocket);      ClientHandler for each
                threadPool.execute(clientHandler);                                     connected client.
                System.out.println("Got a connection");
            }
        } catch (IOException e) {
            e.printStackTrace();
        }
    }

    public void tellEveryone(Object one, Object two) {
        for (ObjectOutputStream clientOutputStream : clientOutputStreams) {
            try {
                clientOutputStream.writeObject(one);
                clientOutputStream.writeObject(two);
            } catch (IOException e) {
                e.printStackTrace();
            }
        }
    }
}

```

List of all the client output streams to send messages to when a message is received.

Send the message and the beat pattern to all the clients.

final BeatBox code

```
public class ClientHandler implements Runnable {  
    private ObjectInputStream in;  
  
    public ClientHandler(Socket socket) {  
        try {  
            in = new ObjectInputStream(socket.getInputStream());  
        } catch (IOException e) {  
            e.printStackTrace();  
        }  
    }  
  
    public void run() {  
        Object userName;  
        Object beatSequence;  
        try {  
            while ((userName = in.readObject()) != null) {  
                beatSequence = in.readObject();  
  
                System.out.println("read two objects");  
                tellEveryone(userName, beatSequence);  
            }  
        } catch (IOException | ClassNotFoundException e) {  
            e.printStackTrace();  
        }  
    }  
}
```

Create an ObjectInputStream for reading messages from this client.

When the client sends a message, it's made of two parts: a String that contains the username and their message; and a Object that represents the beat sequence (this is actually a boolean array, but the server doesn't care about that).

Once we've got the message and the beat sequence, send these to all the clients (including this one).

Appendix B:

The top ten-ish topics that didn't make it into the rest of the book...



We covered a lot of ground, and you're almost finished with this book. We'll miss you, but before we let you go, we wouldn't feel right about sending you out into JavaLand without a little more preparation. We can't possibly fit everything you'll need to know into this relatively small appendix. Actually, we *did* originally include everything you need to know about Java (not already covered by the other chapters), by reducing the type point size to .00003. It all fit, but nobody could read it. So, we threw most of it away, but kept the best bits for this Top Ten-ish appendix. Yep, there's more than ten Really Useful Things that you still need to know.

This really *is* the end of the book. Except for the index (a must-read!).

#11 JShell (Java REPL)

Why do you care?

A REPL (Read Eval Print Loop) lets you run snippets of code without needing a full application or framework. It's a great way to try out new features, experiment with new ideas, and get immediate feedback. We've put this right at the start of this appendix in case you want to use JShell to try out some of the features we'll be talking about in the following pages.

Starting the REPL

JShell is a command-line tool that comes part of the JDK. If `JAVA_HOME/bin` is on your system's path, you can just type “`jshell`” from the command line (full details on getting started are in Oracle's Introduction to JShell (<https://oreil.ly/Ei3Df>)).

```
File Edit Window Help Ammonite
%jshell
| Welcome to JShell -- Version 17.0.2
| For an introduction type: /help intro

jshell>
```

JShell is available only in **JDK 9 and higher**, but the good news is that even if you're running code and applications on an older version of Java, you can still use JShell from a more recent version, since it's completely independent of your “`JAVA_HOME`” or IDE's version of Java. Just run it directly from the `bin` directory of whichever version of Java you want to use.

Run Java code without a class

Try out some Java from the prompt:

```
File Edit Window Help LookMumNoSemiColons
jshell> System.out.println("Hello")
Hello

jshell>
```

Note:

- No need for a class
- No need for a *public static main* method
- No need for a semicolon on the end of the line

Just start typing Java!

Java 9+

More than just lines of code

You can define variables and methods:

```
File Edit Window Help RealJava
jshell> String message = "Hello there "
message ==> "Hello there "

jshell> void greet(String name) {
...>     System.out.println(message+name);
...> }
| created method greet(String)

jshell> greet("you")
Hello there you
```

Semicolons are needed inside blocks of code like methods.

It supports *forward references*, so you can sketch out the shape of your code without having to define everything immediately.

```
File Edit Window Help ForwardLooking
jshell> void doSomething() {
...>     doSomethingElse();
...> }
| created method doSomething(), however, it
cannot be invoked until method doSomethingElse()
is declared
```

Code suggestions

If you press Tab halfway through typing, you'll get code suggestions. You can also use the up and down arrows to cycle through the lines you've typed so far.

```
File Edit Window Help YouCompleteMe
jshell> System.out.pr
print(   printf(   println(
jshell> System.out.print
```

Commands

There are lots of helpful commands that are part of JShell and not part of Java. For example, type `/vars` to see all the variables you've declared. Type `/exit` to, er, exit. Use `/help` to see a list of commands and to get more information.

Oracle has a very useful JShell User Guide (<https://oreil.ly/Ei3Df>), which also shows how to create and run scripts with JShell.

#10 Packages

Packages prevent class name conflicts

Although packages aren't just for preventing name collisions, that's a key feature. If part of the point of OO is to write reusable components, developers need to be able to piece together components from a variety of sources and build something new out of them. Your components have to be able to "play well with others," including those you didn't write or even know about.

Remember way back in Chapter 6, *Using the Java Library*, when we discussed how a package name is like the full name of a class, technically known as the *fully qualified name*. Class List is really ***java.util.List***, a GUI List is really ***java.awt.List***, and Socket is really ***java.net.Socket***. Hey presto, an example of how package names can help prevent name conflicts—there's a List that's a data structure and a List that's a GUI element, and we can use the package names to tell them apart.

Notice that these classes have *java* as their "first name." In other words, the first part of their fully qualified names is "java"; think of a hierarchy when you think of package structures, and organize your classes accordingly.

Preventing package name conflicts

The standard package naming convention is to prepend every class with your reverse domain name. Remember, domain names are guaranteed to be unique. Two different guys can be named Bartholomew Simpson, but two different domains cannot be named *doh.com*.

The class name is
always capitalized.

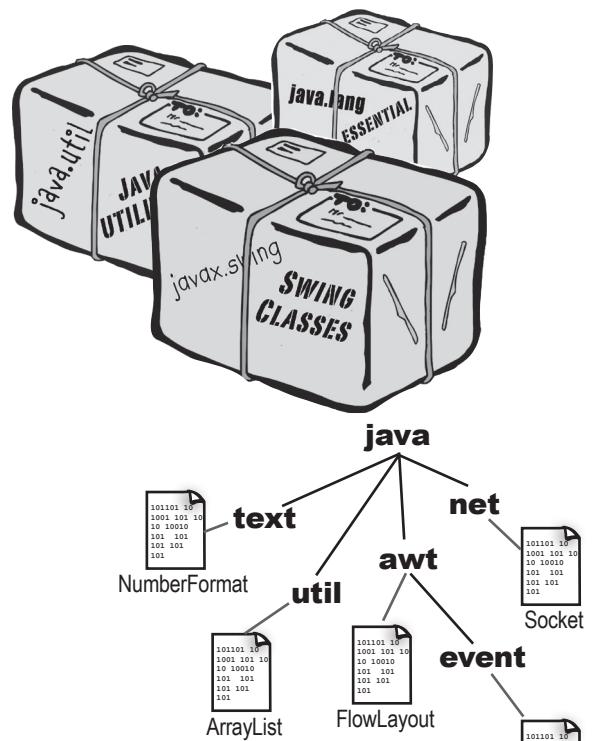
Reverse domain package names

com.headfirstjava.projects.Chart

Start the package with your reverse domain, separated by a dot (.), and then add your own organizational structure after that. Packages are lowercase.

Projects.Chart might be a common name, but adding com.headfirstjava means we have to worry about only our own in-house developers.

When you look at the code samples at https://oreil.ly/hfJava_3e-examples, you'll see we've put the classes into packages named after each chapter to clearly separate the examples.



What does this picture look like to you? Doesn't it look a whole lot like a directory hierarchy?

Packages can prevent name conflicts, but only if you choose a package name that's guaranteed to be unique. The best way to do that is to preface your packages with your reverse domain name.

com.headfirstbooks.Book

package name

class name

#10 Packages, cont.

To put your class in a package:

① Choose a package name

We're using `com.headfirstjava` as our example. The class name is `PackageExercise`, so the fully qualified name of the class is now `com.headfirstjava.PackageExercise`.

② Put a package statement in your class

It must be the first statement in the source code file, above any import statements. There can be only one package statement per source code file, so all classes in a source file must be in the same package. That includes inner classes, of course.

```
package com.headfirstjava;

import javax.swing.*;

public class PackageExercise {
    // life-altering code here
}
```

A note on directories

In the Real World, source files and class files are usually kept in separate directories—you don't want to copy the source code to wherever it's running (a customer's computer or the cloud), only the class files.

The most common structure for Java projects is based off Maven's* convention:

`MyProject/src/main/java` Application sources

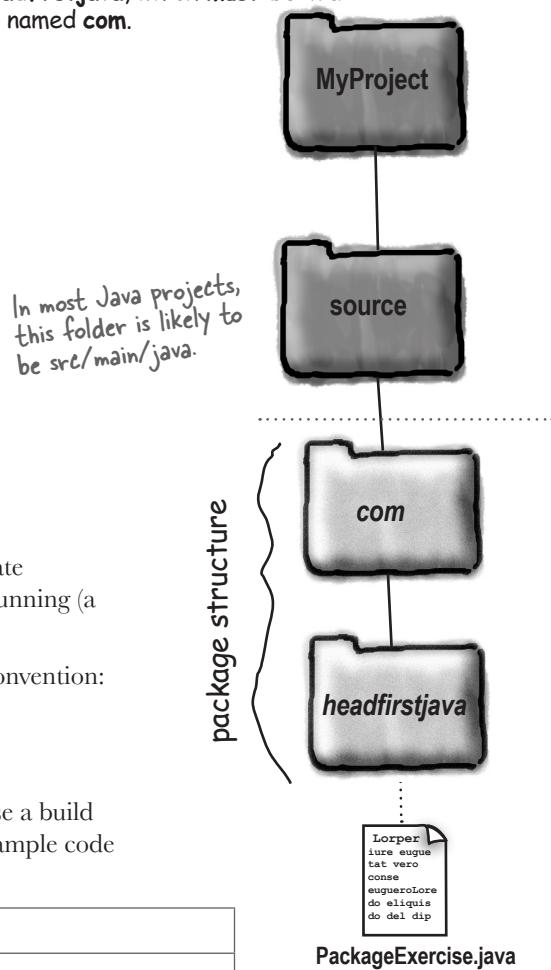
`MyProject/src/test/java` Test sources

The class files are placed elsewhere. Real enterprise systems usually use a build tool like Maven or Gradle to compile and build the application (our sample code uses Gradle). Each build tool puts the classes into different folders:

	Maven	Gradle
Application classes	MyProject/target/classes	MyProject/out/production/classes
Test classes	MyProject/target/test-classes	MyProject/out/test/classes

③ Set up a matching directory structure

It's not enough to *say* your class is in a package by merely putting a package statement in the code. Your class isn't *truly* in a package until you put the class in a matching directory structure. So, if the fully qualified class name is `com.headfirstjava.PackageExercise`, you *must* put the `PackageExercise` source code in a directory named `headfirstjava`, which *must* be in a directory named `com`.



*Maven and Gradle are the most common build tools for Java projects.

#10 Packages, cont.

Compiling and running with packages

We don't need to use a build tool to separate our classes and source files. By using the **-d** flag, you get to decide which **directory** the compiled code lands in, rather than accepting the default of class files landing in the same directory as the source code.

Compiling with the **-d** flag not only lets you send your compiled class files into a directory other than the one where the source file is, but it also knows to put the class into the correct directory structure for the package the class is in. Not only that, compiling with -d tells the compiler to *build* the directories if they don't exist.

Compiling with the **-d (directory)** flag

```
%cd MyProject/source
```

Stay in the source directory! Do NOT cd down into the directory where the .java file is!

```
%javac -d ../classes com/headfirstjava/PackageExercise.java
```

Tells the compiler to put the compiled code (class files) into the classes directory, within the right package structure!! Yes, it knows.

Now you have to specify the PATH to get to the actual source file.

To compile all the .java files in the com.headfirstjava package, use:

```
%javac -d ../classes com/headfirstjava/*.java
```

Compiles every source (.java) file in this directory

Running your code

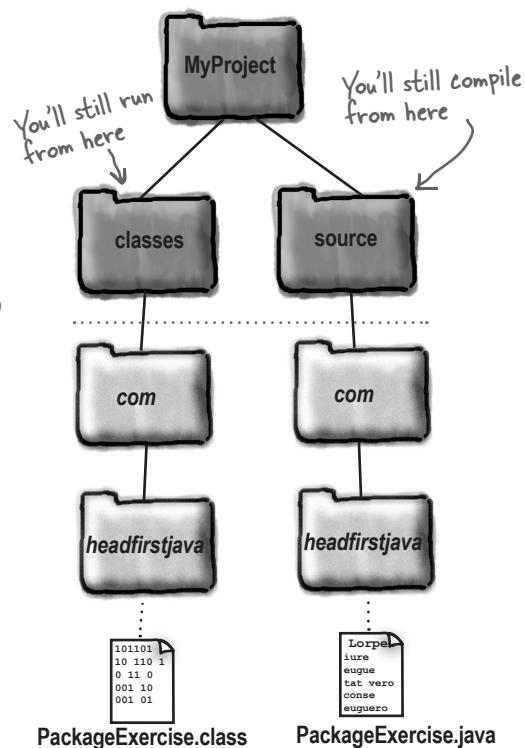
```
%cd MyProject/classes
```

Run your program from the "classes" directory.

```
%java com.headfirstjava.PackageExercise
```

You MUST give the fully qualified class name! The JVM will see that and immediately look inside its current directory (classes) and expect to find a directory named com, where it expects to find a directory named headfirstjava, and in there it expects to find the class. If the class is in the "com" directory, or even in "classes," it won't work!

The **-d flag tells the compiler, “Put the class into its package directory structure, using the class specified after the **-d** as the root directory. But...if the directories aren’t there, create them first and then put the class in the right place!”**



#9 Immutability in Strings and Wrappers

We talked about immutability in Chapter 18, and we'll mention immutability in the last item of this appendix. This section is specifically about immutability in two important Java types: Strings and Wrappers.

Why do you care that Strings are immutable?

For security purposes and for the sake of conserving memory (whether you're running on phones, IoT devices, or the cloud, memory matters), Strings in Java are immutable. What this means is that when you say:

```
String s = "0";
for (int i = 1; i < 10; i++) {
    s = s + i;
}
```

what's actually happening is that you're creating ten String objects (with values "0," "01," "012," through "0123456789"). In the end, *s* is referring to the String with the value "0123456789," but at this point there are *ten different* Strings in existence!

Similarly, if you use methods on String to "change" a String object, it doesn't change that object at all; it creates a new one:

```
String str = "the text";
String upperStr = str.toUpperCase();
```

Reference to the new uppercase String "THE TEXT"

Variable "str" is still "the text"

Creates and returns a NEW String object

How does this save memory?

Whenever you make a new String, the JVM puts it into a special part of memory called the "String Pool" (sounds refreshing, doesn't it?). If there is already a String in the pool with the same value, the JVM *doesn't create a duplicate*; it refers your reference variable to the *existing* entry. So you won't have 500 objects of the word "customer" (for example), but 500 references to the single "customer" String object.

```
String str1 = "customer";
String str2 = "customer";
System.out.println(str1 == str2);
```

These are not only the same value; they're the same object.

Immutability makes reuse possible

The JVM can get away with this because Strings are **immutable**; one reference variable can't change a String's value out from under another reference variable referring to the same String.

What happens to unused Strings?

Our first example created a lot of intermediate Strings that weren't used ("01," "012," etc). These were placed in the String Pool, which is on the heap and therefore eligible for Garbage Collection (see Chapter 9). Strings that aren't used will eventually be garbage-collected.

However, if you have to do a lot of String manipulations (like concatenations, etc.), you can avoid the creation of unnecessary strings by using a *StringBuilder*:

```
StringBuilder s = new StringBuilder("0");
for (int i = 1; i < 10; i++) {
    s.append(i);
}
String finalString = s.toString();
```

This way, the single *mutable* *StringBuilder* is updated every time to represent the intermediate states, instead of ten *immutable* String instances being created and the nine intermediate Strings being thrown away.

Why do you care that Wrappers are immutable?

In Chapter 10, we talked about the two main uses of the wrapper classes:

- Wrapping a primitive so it can act like an object.
- Using the static utility methods (e.g., `Integer.parseInt()`).

It's important to remember that when you create a wrapper object like:

```
Integer iWrap = new Integer(42);
```

that's it for that wrapper object. Its value will *always* be 42. **There is no setter method for a wrapper object.** You can, of course, refer *iWrap* to a *different* wrapper object, but then you'll have *two* objects. Once you create a wrapper object, there's no way to change the *value* of that object!

#8 Access levels and access modifiers (who sees what)

Java has four access levels and three access modifiers. There are only three modifiers because the default (what you get when you don't use any access modifier) is one of the four access levels.

Access levels (in order of how restrictive they are, from least to most restrictive)

public ← public means any code anywhere can access the public thing (by "thing" we mean class, variable, method, constructor, etc.).

protected ← protected works just like default (code in the same package has access), EXCEPT it also allows subclasses outside the package to inherit the protected thing.

default ← default access means that only code within the same package as the class with the default thing can access the default thing.

private ← private means that only code within the same class can access the private thing. Keep in mind it means private to the class, not private to the object. One Dog can see another Dog object's private stuff, but a Cat can't see a Dog's privates.

Access modifiers

public
protected
private

Most of the time you'll use only public and private access levels.

public

Use public for classes, constants (static final variables), and methods that you're exposing to other code (for example getters and setters) and most constructors.

private

Use private for virtually all instance variables, and for methods that you don't want outside code to call (in other words, methods used by the public methods of your class).

Although you might not use the other two (protected and default) much, you still need to know what they do because you'll see them in other code.

#8 Access levels and access modifiers, cont.

default and protected

default

Both protected and default access levels are tied to packages. Default access is simple—it means that only code *within the same package* can access code with default access. So a default class, for example (which means a class that isn't explicitly declared as *public*) can be accessed by only classes within the same package as the default class.

But what does it really mean to *access* a class? Code that does not have access to a class is not allowed to even *think* about the class. And by think, we mean *use* the class in code. For example, if you don't have access to a class, because of access restriction, you aren't allowed to instantiate the class or even declare it as a type for a variable, argument, or return value. You simply can't type it into your code at all! If you do, the compiler will complain.

Think about the implications—a default class with public methods means the public methods aren't really public at all. You can't access a method if you can't *see* the class.

Why would anyone want to restrict access to code within the same package? Typically, packages are designed as a group of classes that work together as a related set. So it might make sense that classes within the same package need to access one another's code, while as a package, only a small number of classes and methods are exposed to the outside world (i.e., code outside that package).

OK, that's default. It's simple—if something has default access (which, remember, means no explicit access modifier!), only code within the same package as the default *thing* (class, variable, method, inner class) can access that *thing*.

Then what's *protected* for?

protected

Protected access is almost identical to default access, with one exception: it allows subclasses to *inherit* the protected thing, *even if those subclasses are outside the package of the superclass they extend*. That's it. That's *all* protected buys you—the ability to let your subclasses be outside your superclass package, yet still *inherit* pieces of the class, including methods and constructors.

Many developers find very little reason to use protected, but it is used in some designs, and some day you might find it to be exactly what you need. One of the interesting things about protected is that—unlike the other access levels—protected access applies only to *inheritance*. If a subclass-outside-the-package has a *reference* to an instance of the superclass (the superclass that has, say, a protected method), the subclass can't access the protected method using that superclass reference! The only way the subclass can access that method is by *inheriting* it. In other words, the subclass-outside-the-package doesn't have *access* to the protected method; it just *has* the method, through inheritance.

Experienced developers writing libraries for other developers to use will find both default and protected access levels very helpful.

These access levels can separate the internals of a library from the API that other developers will call from their code.

#7 Varargs

We saw varargs briefly in Chapter 10, *Numbers Matter*, when we looked at the `String.format()` method. You also saw them in Chapter 11, *Data Structures*, when we looked at *convenience factory methods for Collections*. Varargs let a method take as many arguments as they want, as long as they're of the same type.

Why do you care?

Chances are, you won't write many (or any!) methods with a vararg parameter. But you will likely use them, passing in varargs, since the Java libraries do provide helpful methods, like the ones we just mentioned, that can take as many arguments as they like.

How can I tell if a method takes varargs?

Let's look at the API documentation for `String.format()`:

```
static String format(String format, Object... args)
```

The triple dot (...) says this is method takes an arbitrary number of Objects after the String argument, ***including zero***. For example:

```
String msg = String.format("Message");  
String msgName = String.format("Message for %s", name);  
String msgNumName = String.format("%d, messages for %s", number, name);
```

Pointless for the format()
method, but valid
One varargs argument, "name"
Two varargs arguments,
"number" and "name"

Methods that take varargs generally don't care how many arguments there are; it doesn't matter much. Consider `List.of()`, for example. It doesn't care how many items you want in the List; it will just use add all the arguments into the new list.

Creating a method that takes varargs

You will generally be calling a method that takes varargs, not creating it, but let's take a look anyway. If you wanted to define your own method that, for example, printed out everything passed into it, you could do it like this:

```
void printAllObjects(Object... elements) {  
    for (Object element : elements) {  
        System.out.println(element);  
    }  
}
```

The parameter `elements` is nothing magic; it's actually just an array of Objects. So you can iterate over it the same way as if you'd created the method signature as:

```
void printAllObjects(Object[] elements) {
```

It's the calling code that looks different. Instead of having to create an array of objects to pass in, you get the convenience of passing in an arbitrary number of parameters.

Rules

- A method can have only one varargs parameter.
- The varargs parameter must be the last parameter.

#6 Annotations

Why do you care?

We very briefly mentioned annotations back in Chapter 12, *Lambdas and Streams: What, Not How*, when we said that interfaces that can be implemented as a lambda expression may be marked with a “@FunctionalInterface” annotation.

Adding an annotation to your code can add extra behavior, or an annotation can be a kind of compiler-friendly documentation; i.e., you’re simply tagging the code with some additional information that could optionally be used by the compiler.

You will definitely see annotations used in the Real World, and very likely use them.

Where will you see annotations?

You will see annotations in code that uses libraries and frameworks like Java EE/Jakarta EE, Spring/Spring Boot, Hibernate and Jackson, all of which are very commonly used in the Java world for building large and small applications.

```
@SpringBootApplication ← Class-level annotation
public class HelloSpringApplication {
```

Where you will definitely see annotations is in test code. Back in Chapter 5, *Extra-Strength Methods*, we introduced the idea of testing your code, but what we haven’t shown is the frameworks that make it much easier. The most common one is JUnit. If you look at the code samples at https://oreil.ly/hfJava_3e_examples, you’ll see there are some example test classes in the “test” folder.

```
@Test ← Method-level annotation
void shouldReturnAMessage() {
```

Annotations can be applied to classes and methods, to variables (local and instance) and parameters, and even some other places in the code.

Annotations can have elements

Some annotations include elements, which are like parameters with names.

```
@Table(name="cust")
public class Customer {
```

If the annotation has only one element, you don’t need to give the name.

```
@Disabled("This test isn't finished")
void thisTestIsForIgnoring() {
```

As you saw in the earlier examples, you don’t need to add parentheses to an annotation that doesn’t have elements.

You can add more than one annotation to the class, method, or variable that you’re annotating.

What do they do?

Well, it depends! Some can be used as a sort of compiler-safe documentation. If you add @FunctionalInterface to an interface with more than one abstract method, you’ll get a compiler error.

Other annotations (like @NotNull) can be used by your IDE or by analysis tools to see if your code is correct.

Many libraries provide annotations for you to use to tag parts of your code so the framework knows what to do with your code. For example, the @Test annotation tags methods that need to be run as individual tests by JUnit; @SpringBootApplication tags the class with the *main* method that’s the entry point of a Spring Boot application; @Entity tags a Java class as a data object that needs to be saved to a database by Hibernate.

Some annotations provide behavior on top of your code. For example, Lombok can use annotations to generate common code: add @Data to the top of your class, and Lombok will generate constructors, getters and (if needed) setters, and hashCode, toString, and equals methods.

Sometimes annotations seem to work like magic! The code that does the hard work is hidden away. If you use annotations that are well documented, you'll have a better understanding of what they do and how they work. This will help you fix any issues you may run into.

#5 Lambdas and Maps

Java 8+

Why do you care?

Java 8 famously added lambdas and streams to Java, but what is less well-known is that `java.util.Map` also got a few new methods that take lambda expressions as arguments. These methods make it much easier to do common operations on Maps, which will save you time and brainpower.

Create a new value if there isn't one for the key

Imagine you want to track what a customer does on your website, and you do this using an `Actions` object. You might have a Map of String username to `Actions`. When a customer performs some action that you want to add to their `Actions` object, you want to either:

- Create a new `Actions` object for this customer and add it to the Map
- Get the existing `Actions` object for this customer

It's very common to use an `if` statement and a `null` check to do this (pre-Java 8):

```
Map<String, Actions> custActs = new HashMap<>();
// probably other stuff happens here...
Actions actions = custActs.get(usr);
if (actions == null) { // The value doesn't exist...
    actions = new Actions(usr); ...so create a new Actions and
    custActs.put(usr, actions); add it to the Map with the
} // do something with actions
username as the key.
```

See if there's an Actions object for the username.

It's not a lot of code, but it is a pattern that is used again and again. If you're using Java 8 or higher, you don't need to do this at all. Use `computeIfAbsent`, and give it a lambda expression that says how to "compute" the value that should go into the Map if there isn't an entry for the given key:

```
Actions actions =
    custActs.computeIfAbsent(usr, name -> new Actions(name));
This is EITHER the existing Actions object OR the Actions object created by the lambda if the username wasn't in the Map.
```

This lambda says how to create a new value (an Actions) if the username doesn't have an Actions in the Map yet.

This is the key we're looking for in the Map.

#5 Lambdas and Maps, cont.

Java 8+

Update the value only if it already exists

There may be other scenarios when you want to update a value in the Map only if it exists. For example, you might have a Map of things that you are counting, like metrics, and you want to update only the metrics that you care about. You don't want to add any arbitrary new metric to the Map. Before Java 8, you might use a combination of *contains*, *get*, and *put* to check if the map has a value for this metric and update it if so.

```
Map<String, Integer> metrics = new HashMap<>();
// probably other stuff happens here...
if (metrics.containsKey(metric)) {
    Integer integer = metrics.get(metric); ← If it's in the map, get the value.
    metrics.put(metric, ++integer); ← Increment the value, and put
}
```

See if the metric exists as a key in the map. Alternatively, you could do a "get" and see if the result is null or not.

it back in the map.

Java 8 added **computeIfPresent**, which takes the key you're looking for and a lambda expression, which you can use to describe how to calculate the updated value for the Map. Using this, the code above can be simplified to:

```
metrics.computeIfPresent(metric, (key, value) -> ++value);

```

This also returns the new value if the key was in the map (or null if not), but we didn't need this for our example.

This is the key we're looking for.

The lambda parameters are the key and the value, and we can use these to calculate the new value IF the key exists

Other methods

There are other, more advanced methods on Map that can be useful when you want to “add a new value OR do something with the existing value” (or even remove a value), like *merge* and *compute*. There’s also *replaceAll*, which you can give a lambda expression that calculates a new value for all the values in the map (we could use this, for example, to increment ALL the metrics in our previous example, if we needed to). And, like all the collections, it has a *forEach* that lets us iterate over all the key/value pairs in the Map.

The Java libraries continuously evolve, so even if you think you understand something you’ve used a lot, like List or Map, it’s always worth keeping an eye out for changes that may make your life easier.

Remember, the Java API documentation (<https://oreil.ly/ln5xn>) is a great place to start if you want to see what methods are available on a class, and what they do.

#4 Parallel Streams

Back in Chapter 12, *Lambdas and Streams: What, Not How*, we took a long look at the Streams API. We did not look at one of the really interesting features of streams, which is that you can use them to take advantage of modern multicore, multi-CPU hardware and run your stream operations in parallel. Let's look at that now.

So far, we've used the Streams API to effectively “query” our data structures. Now, imagine those data structures can get big. We mean REALLY big. Like all the data from a database, or like a real-time stream of data from a social media API. We *could* plod over each of these items one by one, in ***serial***, until we get the results we want. Or, we could split the work up into multiple operations and run them at the same time, in ***parallel***, on different CPUs. After Chapters 17 and 18 you might be tempted to run off and write a multithreaded application to do that, but *you don't have to!*

Going parallel

You can simply tell the Streams API you want your stream pipeline to be run on multiple CPU cores. There are two ways to do this.

1. Start a `parallelStream`

```
List<Song> songs = getSongs();
Stream<Song> par = songs.parallelStream();
```

Remember our mock song data from Chapter 12?

2. Add `parallel()` to the stream pipeline

```
List<Song> songs = getSongs();
Stream<Song> par = songs.stream()
    .parallel();
```

They both do the same thing, and you can choose whichever approach you prefer.

OK now what?

Now, you just write a stream pipeline just like we did in Chapter 12, adding the operations you want and finishing off with a terminator. The Java libraries will take care of figuring out:

- How to split the data to run the stream pipeline on multiple CPU cores
- How many parallel operations to run
- How to merge the results of the multiple operations

Multithreading is taken care of

Under the covers, parallel streams use the Fork-Join framework (which we did not cover in this book; see <https://oreil.ly/Xf6eH>), yet another type of thread pool (which we did talk about in Chapter 17, *Make a Connection*). With parallel streams, you'll find the number of threads is equal to the number of cores available wherever your application is running. There are ways to change this setup, but it's recommended to stick with the defaults unless you *really* know what you're doing.

Do not use parallel everywhere!

Before you go running off and making all your stream calls parallel, **wait!** Remember we said back in Chapter 18, *Dealing with Concurrency Issues*, that multithreaded programming was hard, because the solutions you choose depend a lot on your application, your data, and your environment? The same applies to using parallel streams. Going parallel and making use of multiple CPU cores is **not** free and does **not** automatically mean your application will run faster.

There is a cost to running a stream pipeline in parallel. The data needs to be split up, the operations need to be run on each bit of data on separate threads, and then at the end the results of each separate parallel operation need to be combined in some way to give a final result. All of that adds time.

If the data going into your stream pipeline is a simple collection, like the examples we looked at in Chapter 12 (indeed, in *most* places streams are used today), using serial streams is almost definitely going to be faster. Yes, you read that correctly: for most ordinary use cases, you do **not** want to go parallel.

Parallel streams can improve performance when:

- The input collection is BIG (think hundreds of thousands of elements at least)
- The stream pipeline is performing complicated, long-running operations
- The decomposition (splitting) of the data/operations and merging of the results are not too costly.

You should measure the performance with and without parallel before using it. If you want to learn more, Richard Warburton's *Java 8 Lambdas* book has an excellent section on data parallelism.

enumerations

#3 Enumerations (also called enumerated types or enums)

We've talked about constants that are defined in the API, for instance, `JFrame.EXIT_ON_CLOSE`. You can also create your own constants by marking a variable `static final`. But sometimes you'll want to create a set of constant values to represent the **only** valid values for a variable. This set of valid values is commonly referred to as an enumeration. Full-fledged enumerations were introduced way back in Java 5.

Who's in the band?

Let's say that you're creating a website for your favorite band, and you want to make sure that all of the comments are directed to a particular band member.

The old way to fake an “enum”:

```
public static final int KEVIN = 1;
public static final int BOB = 2;
public static final int STUART = 3;

// later in the code
if (selectedBandMember == KEVIN) {
    // do KEVIN related stuff
}
```

We're hoping that by the time we got here "selectedBandMember" has a valid value!

The good news about this technique is that it DOES make the code easier to read. The other good news is that you can't ever change the value of the fake enums you've created; KEVIN will always be 1. The bad news is that there's no easy or good way to make sure that the value of `selectedBandMember` will always be 1, 2, or 3. If some hard-to-find piece of code sets `selectedBandMember` equal to 812, it's pretty likely your code will break.

This IS the OLD way to fake an enum, but you will still see code like this in Real Life (e.g., the older Java libraries like AWT).

However, if you have any control over the code, try to use enums instead of constants like this. See the next page...

#3 Enumerations, cont.

Let's see what the band members would look like with a "real" enum. While this is a very basic enumeration, most enumerations usually are this simple.

An official "enum"

```
public enum Member { KEVIN, BOB, STUART };
```

This kind of looks like a simple class definition, doesn't it? It turns out that enums ARE a special kind of class. Here we've created a new enumerated type called "Member."

```
public class SomeClass {
    public Member selectedBandMember;

    // later in the code...
    void someMethod() {
        if (selectedBandMember == Member.KEVIN) {
            // do KEVIN related stuff
        }
    } No need to worry about this variable's value!
}
```

The "selectedBandMember" variable is of type "Member," and can ONLY have a value of "KEVIN," "BOB," or "STUART."

The syntax to refer to an enum "instance"

Your enum extends java.lang.Enum

When you create an enum, you're creating a new class, and ***you're implicitly extending java.lang.Enum***. You can declare an enum as its own standalone class, in its own source file, or as a member of another class.

Using "if" and "switch" with enums

Using the enum we just created, we can perform branches in our code using either the `if` or `switch` statement. Also notice that we can compare enum instances using either `==` or the `equals()` method. Usually `==` is considered better style.

```
Member member = Member.BOB; ← Assigning an enum value to a variable
if (member.equals(Member.KEVIN)) { Both of these work fine!
    System.out.println("Bellloooo!");
}
if (member == Member.BOB) { "Poochy" is printed.
    System.out.println("Poochy");
}

switch (member) { ← Pop Quiz! What's the output?
    case KEVIN: System.out.print("Uh... la cucaracha?");
    case BOB: System.out.println("King Bob");
    case STUART: System.out.print("Banana!");
}
```

You can add a bunch of things to your enum like a constructor, methods, variables, and something called a constant-specific class body. They're not common, but you might run into them.

Answer: King Bob
Banana!

#2 Local Variable Type Inference (var)

Java 10+

If you're using Java 10 or higher, you can use `var` when you're declaring your *local* variables (i.e., variables inside methods, **not** method parameters or instance variables).

```
var name = "A Name";    name is a String
```

This is another example of *type inference*, where the compiler can use what it already knows about the types to save you from writing more. The compiler knows `name` is a String because it was declared as a String on the right hand side of the equals sign.

```
var names = new ArrayList<>();  
var customers = getCustomers();  
    ↗ If getCustomers() returns List<Customer>,  
    this is a List<Customer>.
```

An ArrayList

Type inference, NOT dynamic types

When you declare your variable using `var`, **it still has a type**. It's not a way of adding dynamic or optional types to Java (it's **not** like Groovy's `def`). It's simply a way of avoiding writing that type twice.

You do have to somehow tell the compiler what the type is when you declare the variable. You can't assign it later. So, you **can't** do this:

```
var name;    Does NOT compile!!
```

because the compiler has no idea what type `name` is.

It also means that you can't change its type later:

```
var someValue = 1;  
someValue = "String"; ← Does NOT compile!!
```

Someone has to read your code

Using `var` does make the code shorter, and an IDE can tell you exactly what type your variable is, so you might be tempted to use `var` everywhere.

However, someone reading your code might not be using an IDE or have the same understanding of the code as you.

We did not use `var` in this book (even though it would have been easier to fit the code on the pages), because we wanted to be explicit to you, the reader, about what the code was doing.

Tip: Better with useful variable names

If you don't have the type information visible in the code, descriptive variable and method names will be extra helpful to a reader.

```
var reader = new BufferedReader(get("/"));
```

We can figure out what this is and what to use it for.

```
var stuff = doTheThing();
```

We have NO IDEA what this is.

Tip: Variable will be the concrete type

In Chapter 11 we started “programming to interfaces”; i.e., we declared our variables as the interface type, not the implementation:

```
List<String> list = new ArrayList<>();
```

If you're using `var`, you can't do this. The type will be the type from the right-hand side:

```
var list = new ArrayList<String>();
```

This is an ArrayList<String>.

Tip: Don't use var with the diamond operator

Look at the last example. We declared `list` first as a `List<String>` and used the diamond operator (`<>`) on the right-hand side. The compiler knows the type of the list element is a String from the lefthand side.

If you use `var`, like we did in the second example, the compiler no longer has this information. If you want the list to still be a list of Strings, you need to declare that on the righthand side; otherwise, it will contain Objects.

```
var list = new ArrayList<>();
```

This is an ArrayList<Object>, probably not what you wanted.

Read all the style guidelines from the OpenJDK developers (<https://oreil.ly/eVfSd>).

#1 Records

Why do you care?

A “simple” Java data object is often not simple at all. Even a data class (sometimes called a Java Bean, for historical reasons) with only a couple of fields requires quite a lot more code than you might expect.

A Java data class, before Java 16

Imagine a basic Customer class, with a name and an ID:

```
public final class Customer {
    private final int id;
    private final String name;

    public Customer(int id, String name) {
        this.id = id;
        this.name = name;
    }

    public int getId() {
        return id;
    }

    public String getName() {
        return name;
    }

    public boolean equals(Object o) { }

    public int hashCode() { }

    public String toString() { }
}
```

We've left out the details of the equals, hashCode, and toString methods, but you would probably want to implement those methods, especially if you're going to use this object in any collections. We've also left off the “setters”; this is an immutable object with final fields, but in some cases you might want setters as well.

That's a lot of code! It's a simple class with two fields, and the full code, including implementation, is 41 lines!

Java 16+

What if there was a special syntax for data classes?

Guess what? If you're using **Java 16** or higher, there is! Instead of creating a *class*, you create a *record*.

These are the record's components. These translate into instance variables and an accessor method for the variable.

Don't use "class," use "record."

```
public record Customer(int id, String name) {}
```

This record header also defines what the constructor looks like (the order of the parameters for the constructor).

That's it. That's all you need to do to replace the 42 lines of code of the “old” Customer data class.

A record like this one has instance variables, a constructor, accessor methods, and equals, hashCode, and toString methods.

Using a record

When you're using a record that's already been defined, it looks exactly the same as it would if the record class was a standard data class:

```
Customer customer = new Customer(7, "me");
System.out.println(customer);
System.out.println(customer.name());
```

The output looks like:

```
File Edit Window Help Vinyl
%java UsingRecords
Customer[id=7, name=me]
me
```

Notice this is NOT
getName().

Records have a pretty
toString by default.

Goodbye “get”

Did you notice something? Records don't use the classic “get” prefix for the methods that let you read the instance variables (hence we carefully called them “accessors” and not “getters”). They just use the name of the record component as the method name.

#1 Records, cont.

You can override constructors

The constructor, accessors, and equals, hashCode, and toString methods are all provided by default, but you can still override their behavior if you need something specific.

Most of the time, you probably won't need to. But if you want, for example, to add validation when you create the record, you can do that by overriding the constructor.

```
public record Customer(int id, String name) {
    public Customer(int id, String name) {
        if (id < 0) {
            throw new ValidationException();
        }
        this.id = id;
        this.name = name;
    }
}
```

Actually, it's even easier than that. The example above is a **canonical constructor**, i.e., the normal kind of constructor we've been using everywhere. But records also have a **compact constructor**. This compact constructor assumes all the normal stuff is taken care of (having the right number of parameters in the right order, and all assigned to the instance variables) and lets you define only the other stuff that matters, like validation:

```
Record header defines what the
constructor looks like when you call it
public record Customer(int id, String name) {
    No need to define the
    constructor parameters
    if (id < 0) { Can still access the parameters
        throw new ValidationException();
    } No need to assign anything to
        the instance variables
}
```

When you call the Customer's constructor, you still need to pass it an ID and a name, and they will still be assigned to the instance variables (that's all defined by the record header). All you need to do to add validation to the constructor is use the compact form and let the compiler take care of all of the rest.

```
Customer customer = new Customer(7, "me");
Even with a compact constructor you must pass
in arguments for all of the record components.
```

Java 16+

You can override or add methods

You can override any of the methods and add your own (public, default, or private) methods. If you are migrating existing data classes to use records, you may want to keep your old equals, hashCode, and toString methods.

```
public record Customer(int id, String name) {

    public boolean equals(Object o) {
        return id == ((Customer) o).id;
    } Overrides the equals method to
        provide custom behavior

    private boolean isValidName(String name) {
        // some implementation
    }
}
```

Records can have
"normal" methods too

You can create a protected method; the compiler won't stop you, but there's no point—records are always *final* classes and can't be subclassed.

Records are immutable

In Chapter 18, we talked about making data objects *immutable*. Immutable objects are safer to use in concurrent applications, because you know that it's impossible for more than one thread to change the data.

It's also easier to reason about what's happening in your application if you know the data classes can't change, so even in applications that aren't multithreaded, you may find immutable data objects being used. And in #9 in this appendix we saw how immutability in Strings can save memory.

Records are immutable. You can't change the values in a record Object after you have created it; there are no "setters" and no way to change the instance variables. You can't access them directly from outside the record, only read them via the accessor method.

If you try to change one of the record's instance variables from inside the record, the compiler will throw an exception. A record's instance variables are *final*.

Find out more about records in Oracle's Record Classes documentation (<https://oreil.ly/D7fh3>). There, you can also read about some of the other new language features available in Java 17 that we didn't get a chance to cover, like Pattern Matching, Sealed Classes, Switch Expressions, and the very useful Text Blocks.

Symbols

!= and != (not equals) 151
% percent sign in format String 297–300
&&, || ('and' and 'or' operators) 151
++ -- (increment/decrement) 106, 115
-> (arrow operator) for lambda expressions 388
. (dot operator) 36, 54, 61, 80
// (comment syntax) 12
:: (method reference) 408
<=, ==, !=, >, >= (comparison operators) 13, 86, 151, 348
<> (diamond operator) 312–313, 698
= (assignment operator) 13
== (equals operator) 13, 86, 348

A

abandoned objects. *See* garbage collection
abstract classes 202–209
 conditions for using 229
 constructors in 253
 interface implementation 226
 and polymorphism 199, 208–209
 and static methods 278–282
abstract class modifier 202
abstract methods 205–206, 222, 226, 396
accept() 601

Index

access control
 levels 689–690
 modifiers 81–82, 689–690
 object locking for 645, 647, 649, 656
polymorphism 192
 and subclasses 191
access levels 192, 193, 689–690
access modifiers 81–82, 251, 689–690
Accessors and Mutators. *See* Getters and Setters
ActionEvent 467–468, 481–482, 523
ActionListener 481, 482, 491
actionPerformed() 467, 481, 482
addActionListener() 467–469
add(anObject), ArrayList 137
Advice Guy 598–599
alphabetical (natural) ordering in Java 315
'and' and 'or' operators (&&, ||) 151
Animal simulation program 172–178, 200–204
animation 492–495
annotations 692
argument list 248–250, 302
arguments 74, 76
 autoboxing 292
 catch 428
 with generic types 358–362
Getters and Setters 79
 and no-arg constructor 247–248
 number formatting 297, 302
 passing to super() 257
 polymorphic 189–190, 192
ArrayList 132–139, 314
 autoboxing 291

- casting 231
 - containing references of type Object 213–215
 - diamond operator 312–313
 - as generic class 322–323
 - HashSet instead of 347
 - relationship to List in sorting project 310
 - StartupBust object 142
 - type parameters with 323
 - and value of generics 320, 321
 - arrays 19, 59–62
 - versus ArrayList 137–140
 - behavior of objects in 83
 - for declaring multiple return values 78
 - dot operator in 61
 - elements of as objects 59
 - polymorphic 188
 - SimpleStartupGame bug fix 128–130
 - assignment
 - arrays 60
 - autoboxing 293
 - lambda expression variables 394
 - object 55, 186–188
 - primitive variables 52
 - reference variables 55
 - assignment operator (=) 13
 - AtomicInteger 656
 - atomic transactions, thread concurrency 646
 - atomic variables 655–657
 - attributes 30
 - autoboxing of primitive types 291–292
 - Autocloseable 577
 - awaitTermination 629
- ## B
- background components (containers), GUI 510, 511
 - BeatBox app 422
 - Chat client and server 588–589
 - client program 674–681
 - GUI for 528–533
 - MIDI Music Player 423
- saving drum pattern 579–582
 - saving objects 540
 - server program 681–682
 - behavior of an object. *See* methods
 - BindException 593
 - Boolean anyMatch 377
 - Boolean expressions 13
 - autoboxing with 292
 - not equals (!= and !) 151
 - querying collection for values 410
 - variables as incompatible with integers 14
 - boolean primitive type 51, 53
 - BorderLayout manager 478–482, 511, 513, 514–517
 - Bottle Song application 16
 - BoxLayout manager 513, 521
 - branching (if/else) 12
 - break statements 105
 - BufferedReader 566, 594
 - BufferedWriter 565, 572
 - buffers 565, 566
 - buttons, GUI 463–468
 - ActionEvent 481–482
 - BorderLayout 514–517
 - FlowLayout 519–520
 - inner class two-button code 487
 - response and timing 464–465
 - bytecode 2
 - byte primitive type 51, 53
- ## C
- canonical constructor. *See* constructors
 - CAS (compare-and-swap) operations 655–656
 - casting 78, 111, 117, 218, 551
 - catch argument 428
 - catch blocks. *See* try/catch blocks
 - chained streams 379, 543, 571
 - Chair Wars scenario 28–35, 168–169

- channels, client and server networking 587
 - reading/receiving with 594–596
 - SocketChannel 591, 594, 595, 596–602
- character (%c) formatting type 301
- char primitive type 51, 53
- Chat client app 588–589, 604–608, 632–633
- Chat server app 588–589, 606–608
- check.addItemListener(this) 526
- check box (JCheckBox) 526
- checked versus runtime exceptions 430
- checkUserGuess() 145
- checkYourself() 102, 104, 130
- classes 8, 28, 41, 72. *See also* inheritance
 - abstract. *See* abstract classes
 - Animal simulation program 173, 174
 - concrete 202–212
 - conditions for making 229
 - and constructors 243
 - deserialization and versioning 556–557
 - designing 34
 - documentation 162
 - efficiency of not saving with object 553
 - elements of 7
 - Executors 615
 - File 564
 - Files 572–573
 - final 285, 286
 - full names in Java library 155–157
 - generic 321–323, 328–330
 - Graphics2D 474–475
 - hierarchy design with inheritance 175
 - for immutable data 658, 661
 - implementing multiple interfaces 228
 - inner. *See* inner classes
 - instance variable declarations inside 238
 - and Java API packages 154–162, 685, 686
 - JComponent 510
 - lambda expressions as 389
 - with main() 9
 - Math 276–280
 - in object creation 36–37
 - versus objects 35
 - Object superclass 210–217
 - Paths 572–573
 - PetShop program class tree modifications 221–228
 - Random 111
 - records 699–700
 - Sequencer 424–427
 - Simple Startup Game 99–124
 - and static variables 283
 - tester class 36
 - Thread 609–610
 - and types 50
 - without type parameters 324
- client application, networking 588–600, 604–608, 632–633
- client-server relationship 589–593
- Code Kitchens
 - BeatBox app 674–682
 - GUI for Beatbox 528–533
 - music with graphics 496–503
 - playing sound 445–453
 - saving BeatBox pattern 579–582
- collect() 377, 378, 410
- Collection API 345–346, 376
- collections
 - ArrayList. *See* ArrayList
 - common operations 374
 - and concurrency 662–666
 - count operation on 410
 - enhanced for loop 116
 - factory methods 356–357
 - generics for type-safe 320–324
 - List. *See* List interface
 - Map 355, 409
 - parameterized types 137
 - streams as queries 385, 387
- Collections class 314, 323
- Collections.sort() 314–315
 - Comparator 331–338
 - compare() 333
 - and List.sort() 332
 - with Objects instead of Strings 317–319

the index

Collectors class 378, 383, 387, 409
Collectors.joining 409
Collectors.toList 387, 409
Collectors.toMap 409
Collectors.toSet 409
Collectors.toUnmodifiableList 356, 387, 409
Collectors.toUnmodifiableMap 409
Collectors.toUnmodifiableSet 409
command-line arguments, MIDIEvent 452
commas, formatting large numbers with 296, 298
comment syntax (//) 12
compact constructor 700
Comparable interface 354
 Collections.sort() 327–330
 versus Comparator 332, 335
Comparator interface
 as SAM type 397
 versus Comparable 332, 335
 lambda expressions with 341–343, 390–394
 and method reference 408
 sorting with 314, 331–338
 and TreeSet 354
compare() 332, 333, 341
compareAndSet method 655–656
compare and swap operations. *See CAS*
compareTo() 329–330, 332, 352
comparison operators 13, 86, 151, 348
compiler 2, 10–11, 687
compiler-safe documentation, annotations as 692
computeIfAbsent 693
computeIfPresent 694
concatenated objects 19
concrete classes 202–212
concurrency 639–669
 atomic variables 655–657
collections 662–666
immutable objects 658–661
locking 645
lost update scenario 650–653
multithreading 630
race conditions 641–643
synchronization 646–654
trade-offs in 666
ConcurrentModificationException 663
conditional branching 15
conditional expressions 13, 15
conditional test, for loop 114
connection, client-server 591–593
connection streams 543
constants 41, 284, 696
constructors 243–259
 chaining 253–259
 in deserialization 552
 function of 244
 initializing state of object 245–248
 overloaded 258–259
 overriding 700
 private 191, 251, 278, 282
 review 251
 superclass 252–259
containers (background components), GUI 510, 511
contains() 403
contract
 modifying class tree 220–226
 public methods as 192–193, 219
ControllerEvent 497, 500
convenience methods 357–358, 387
CopyOnWriteArrayList 665–666
CountDownLatch 625
count operation, on collections 410
C programming language 56
curly braces ({}), for classes and methods 12

D

DailyAdviceServer 602
 data structures. *See* collections
 date formatting 302
 DDD (Deadly Diamond of Death) 225
 -d (directory) flag 687
 deadlock, synchronization 654
 decimal (%d) formatting type 301
 declarations
 exceptions 426, 441–443
 method 78, 144, 205–206, 222, 238
 object 186–188
 variable 50–52, 54, 84, 85, 116, 144, 238
 default access level 689, 690
 default method 396
 default value
 instance variable 84
 static variable 283
 deserialization 551–557
 diamond operator (<>) 312–313, 698
 directories
 File object with 564
 packages 573, 686
 distinct(), Stream 375, 406–407
 dot operator (.) 36, 54, 61, 80
 double primitive type 51, 53
 drawing 2D graphics 471, 472–475, 501–503
 Duck constructor 243–248, 250–251, 281–283
 ducking exceptions 441–443
 duplicate code, avoiding with inheritance 184
 duplicate results, removing 406–407

E

e-flashcards example, saving to text file 560
 encapsulation 80–82
 enhanced for loop 106, 116

enumerations 696–697
 equality 348, 349
 equals() 86, 349–351, 697
 equals operator (==) 13, 86, 348
 event handling 465–471
 getting graphics 477
 JCheckBox 526
 JList 527
 listener interface 466–469
 MIDIEvents 449–452, 497–503
 static methods 498–499
 event object 470
 event source 467–469
 exception handling 421, 426–444
 finally block 433, 444, 574–575
 flow control 432–433
 multiple exceptions 438
 try/catch blocks. *See* try/catch blocks
 try-with-resources statement 576–577
 exceptions
 BindException 593
 catching. *See* try/catch blocks
 checked versus runtime 430
 concurrency and collections 663
 declaring 436, 441–443
 ducking 441–443
 methods 425
 multiple 435
 NumberFormatException 294
 as polymorphic 436
 throwing 429–432
 Executors 626
 Executors class 615
 ExecutorService 615, 626–629
 ExecutorService.shutdown() 615, 629
 ExecutorService.shutdownNow() 629
 exercises
 Be the... 21, 42, 63, 88, 118, 195, 306, 363, 395, 505
 Code Magnets 20, 43, 64, 119, 163, 386, 455, 583,
 634

- Mixed Messages 23, 194, 372
 - Popular Objects 269
 - Sharpen Your Pencil 5–6, 15, 37, 52, 87, 107, 134–141, 145–147, 207, 250–251, 259, 287, 293, 334, 343, 353, 397, 431, 434, 440, 493, 502–503, 580, 600–601
 - True or False 307, 454, 582
 - What’s the Declaration? 233
 - What’s the Picture? 232
 - Which Layout? 534–535
 - Who Am I? 45, 89, 504, 631
 - Who Does What? 374
 - explicit cast 78
 - extending classes. *See* inheritance
 - extends keyword 328–330
 - Extreme Programming 101
- ## F
- factory methods 356–358, 387, 615
 - File class 564
 - FileInputStream 551
 - File object 564
 - FileOutputStream 542, 543
 - FileReader 566
 - Files class 572–573
 - files, source file structure 7
 - FileWriter 559, 565
 - filter() 400–403
 - filter streams. *See* chained streams
 - final keyword
 - adding to field declaration 660, 666
 - classes 191, 285, 286
 - methods 191, 285
 - variables 275, 284–286
 - finally block 433, 444, 574–575
 - findFirst(), Optional 377
 - Fireside Chats
 - for loop versus forEach method 371
- instance versus static variables 304–305
 - JVM and compiler roles 10–11
 - variable discussion on life and death 266–267
 - floating point (%f) formatting type 301
 - float primitive type 51, 53
 - flowchart for Sink a Startup 97
 - flow control, exceptions 432–434
 - FlowLayout manager 513, 518–520
 - forEach() 370, 388, 393, 694
 - Fork-Join framework 695
 - for loops 114–116
 - enhanced 106, 116
 - versus forEach() 370–373
 - SimpleStartup class 105
 - format() 298
 - format specifiers 297–298, 300–302
 - Formatter class 296
 - formatting numbers 296–302
 - frames, GUI 462, 511, 522
 - Friesen, Jeff
 - Java I/O, NIO and NIO.2 597
 - fully qualified name, Java library packages 155–157
 - @FunctionalInterface annotation 396
 - functional interface, lambda expression 389–396
 - Function, in map method 405
- ## G
- GameHelper class 112, 142, 152–153
 - GameHelper object 143
 - Garbage-Collectible Heap 238
 - garbage collection 40, 57–58, 262–265
 - generics 320–324
 - classes 321–323
 - extends or implements 328–330
 - methods 324, 375
 - and polymorphic arguments 358–362
 - type parameters 362

`getPreferredSize()` 522

`getSequencer()` 426

Getters and Setters 79–82, 646

`getUserInput()` 112

gradient blend, graphics object 475

Gradle 686

graphics 471–475. *See also* GUI

Graphics superclass 474

greater than operator (`>`) 13

Guessing Game example 38–40

GUI 461–501

- abstract classes in 204

- BeatBox app 528–533, 674

- BorderLayout 478–482, 511

- BoxLayout 513

- building graphics 471–475

- buttons. *See* buttons, GUI

- components 462, 471–475, 510, 523–527

- event handling 465–471

- FlowLayout 513

- inner classes 484–494

- layout managers 511–522

- listener interface 466–469

- Swing 462, 509–533

H

HAS-A test for inheritance 179–183

`hashCode()` 348–351

heap 40, 57, 238–241

hexadecimal (%x) formatting type 301

Hitchens, Ron

- Java NIO 597

HTML API docs 160

I

if/else statement 12, 15

if statement 15, 697–698

if test 15

images on GUI widget 471, 473

immutability 688, 700

immutable objects 658–661, 665–666

imports, static 303

import statement 155, 157

increment/decrement operators (++) and (--) 106, 115

`increment()`, synchronizing 652

index position, List 345

`InetSocketAddress` 591

inheritance 31, 168–185. *See also* polymorphism

- Animal simulation program 172–178

- benefits of using 184–185

- dos and don'ts 183

- implementing abstract methods 206

- keeping trees shallow 191

- relationship to objects 216

- subclasses. *See* subclasses

- superclass. *See* superclass

inheritance trees 191, 220–231

initializing

- with constructor 246–248

- instance variables 84

- static variables 283

inner classes 191, 337, 484–494

- drawing 2D graphics 501

- lambda expressions with 490–491

- relationship to outer class 484–486

- two-button code 487

`InputStreamReader` 596

instances, inner and outer class 485–486

instance variables 34, 35

- Animal simulation program 173, 174

- declaring 84, 238

- default values 84

- Getters and Setters 79

- on heap 241

- inability to use with static methods 279

- initializing 84

the index

- lack of in Math class 276
- life and scope of 260–267
- versus local variables 85, 238–240
- matching with parameter types 76
- pass-by-value/pass-by-copy 77
- private access modifier for 81
- in serialization process 544–546
- setter methods for 80–82
- versus static variables 304–305
- subclass 169
- transient 549, 550, 553
- int array variable 59
- integers, incompatibility with boolean variables 14
- interactive components, GUI 510
- interface keyword 226–231
- interfaces 199
 - functional interface 389–396, 491
 - naming 154–156
 - polymorphism 226–231
- intermediate operations 376
 - as lazily evaluated 383–384
 - chaining operations 380
 - creating stream pipeline 379
- int primitive type 51, 53
- I/O 540, 571
 - deserialization 551–555
 - exception handling 574–577
 - networking. *See* networking
 - saving data to text file 559–571
 - saving objects 541–558
 - serialization 541–550, 554–557
 - streams 543
- IS-A test 179–180, 183, 188, 253
- iteration expression, for loops 114
- iteration variable declaration, enhanced for loop 116
- J**
- Java
 - basic elements 41
 - code structure 7–8
- history 4
- setting up xxx
- speed and memory usage 4
- version naming conventions 5
- workings of 2–3
- Java API 125–164
 - classes and packages 154–162. *See also* packages, Java API
 - documentation 158–162
- Java Collections Framework 309. *See also* collections
- Java-Enabled House 17
- JavaFX 464
- Java I/O, NIO and NIO.2 (Friesen) 597
- Java Module System 161
- java.nio.channels package 597
- Java NIO (Hitchens) 597
- JavaSound API 421, 423
- java.util API 314
- Java Virtual Machine. *See* JVM
- JCheckBox 526
- JComponent class 510
- JFrame 462, 510, 522
- JPanel 510, 518
- JPanel.paintComponent() 472–473, 495
- JPEG on widget 473
- JScrollPane 524
- JShell 684
- JTextArea 524–525
 - JTextArea.requestFocus() 524
 - JTextArea.selectAll() 524
 - JTextArea.setLineWrap(true) 524
 - JTextArea.setText() 524
- JTextField 523
 - JTextField.requestFocus() 523
 - JTextField.selectAll() 523
- JVM (Java Virtual Machine) 9–11

K

key-value pairs, Map 355

keywords in Java 53, 328

L

lambda expressions 340–346, 388–397

anatomy of 391–392

calling Single Abstract Method 389

forEach method 370

implementing Predicate 402

map operation 405, 408, 693–694

method reference replacement for 408

parameters 393–394

replacing inner class with 490–491

threads 612–614, 621

void return in 393

latch.countDown 625

layout managers 509, 511–522, 526

BorderLayout 478–482, 513–517

BoxLayout 513

changing frames 522

differing policies of 512

FlowLayout 513, 518–520

less than operator (<) 13

limit(), Stream 375–377, 381

list.addListSelectionListener(this) 527

listener 467–469, 481

ActionListener 482, 491

check.addItemListener(this) 526

list.addListSelectionListener(this) 527

non-GUI event 497

listener interface, event handling 466–469

List interface 345

collecting to 383, 400, 409

Comparator with 331

sorting with 310–319

unmodifiable output 356–357, 387

using versus implementation type 313

List.of() 357

list.setSelectionMode(ListSelectionModel.SINGLE_SELECTION) 527

list.setVisibleRowCount(4) 527

List.sort() 332

literals, assigning values 52

local variables

declaring 238

versus instance variables 85, 238–240

life and scope of 260–261, 266–267

parameters as 74

references on heap 240

in Stack 238

type inference 698

locking 645–649, 656

long count() 377

long primitive type 51, 53

loop block 13

loops 12, 13

for 105, 106, 114–116, 370–373

while 13, 115, 566

M

main() 8, 12, 14, 27

classes with 9

ducking exception 442

multithreading 611, 613

SimpleStartupGame class 108, 110–111

StartupBust object 142

in tester class 36

uses of 38

makeEvent() 499–500

Map interface 355

Collection API 346

and collections 409

key-value pairs 345

lambda expressions 693–694

Map.of() 357

Map.ofEntries() 357

the index

- map(), Stream 375
- Math.abs() 288
- Math class 276–280, 288–289
- Math.max() 289
- Math.min() 289
- Math.random() 111, 288
- Math.round() 289
- Math.sqrt() 289
- Maven, Java project structure 686
- memory
 - garbage collection 262–265
 - stack and heap in object lifecycle 238–241
 - String immutability 688
- Message, MIDIEvent 449–451, 498
- metadata 572
- method reference, stream 408
- methods 7, 8, 30, 34. *See also* local variables
 - abstract 205–206, 222, 226, 396
 - Animal simulation program 173
 - arguments. *See* arguments
 - autoboxing 292
 - calling non-static from static 280
 - versus constructors 245
 - declarations and implementations 144–145
 - declaring or ducking 441–443
 - descriptive naming best practice 698
 - exception handling 425–426, 443
 - final 191, 285
 - generic 321, 324, 362
 - inheritability 178
 - listener interface 466
 - local variables declared inside 85, 238
 - making available to all code 41
 - matching variable and parameter types 76
 - Math class 288–289
 - multiple parameters with 76
 - multiple return value declarations 78
 - non-abstract methods with abstract classes 206
 - overloading 193
- overriding. *See* overriding methods
- parameters. *See* parameters
- returning values from 75
- SimpleStartup class test code 102–104
 - in Stack 238, 239
 - static 276–280
 - and static variables 286
 - subclasses 182
 - superclass 230
 - test coding 101
 - and type Object's role in code 212, 215
 - varargs 691
 - working with inheritance 177–178
- MIDIEvents 449–452, 497–503
- MIDI Music Player 423–424, 446–453
- mocking 310
- mock Songs class 311
- MovieTestDrive class 37
- multiple inheritance problem 224–225
- multithreading 609–630, 695. *See also* concurrency
 - coordinating threads 622–625
 - parallel streams 695
 - Runnable interface 612–617
 - scheduling 617–619
 - SimpleChatClient complete 632–633
 - sleep() 622–624
 - stack 610–614
 - thread pools 626–629
 - thread states 616
 - music video 496–503
- Mutators and Accessors. *See* Getters and Setters

N

- naming
 - classes and interfaces 154–156
 - variables 50–52, 53, 61
- natural ordering in Java (alphabetical) 315
- networking 587–635
 - channels. *See* channels

client-server relationship 589–593
 simple Chat client app 604–608
 simple Chat server app 606–608
`nextInt()` 111
`NIO.2` 597
`NIO` (non-blocking I/O) 561
 no-arg constructors 247–248, 250
 non-public class 191
 non short circuit operators (`&`, `|`) 151
`not equals` (`!equals` and `!`) 151
 null reference 58, 265
`NumberFormatException` 294
 numbers, formatting 296–302
 numeric primitive types 51

O

`Object` class 210–217
 object graph 546, 548, 550
`ObjectInputStream` 551
`ObjectOutputStream` 542, 543
`ObjectOutputStream.close()` 542
`ObjectOutputStream.writeObject()` 542
 object references. *See* reference variables
 objects. *See also* arrays; Strings
 array elements as 60, 83
 assignment 55, 186–188
 behavior 539. *See also* classes
 versus classes 35
 and collections 374
 creating 36–37, 55, 242–254
 declaring 186–188
 eligibility for garbage collection 262–265
 equality 86, 348, 349
 in heap 238–240
 immutable 658–661, 665–666
 instance variables as living inside 241
 instantiating 103

lambda expressions as 389
 lifecycle of 253, 260–267
 locking 645–649, 656
 of type `Object` 212
 saving state 539, 541–558
 superclass constructors 252–259
`OO` (object-oriented) development 14, 27–48. *See also* classes; objects
 event handling 481
 inheritance 168–185
 saving object state 539
 operators. *See also* primitive variables
 ‘and’ and ‘or’ operators 151
 and autoboxing 293
 comparison 13, 86, 151, 348
 `equals` (`==`) 13, 86, 348
 increment/decrement (`++` and `--`) 106, 115
 non short circuit 151
 post-increment 105
 short circuit 151
 optimistic locking 656
 Optional value, returning from collection query 410–414
 outer class, relationship to inner class 484–486
`OutputStream` 596
 overloaded constructors 258–259
 overloading methods 193
 overriding constructors 700
 overriding methods 32, 169–194
 `hashCode()` and `equals()` 350–351
 `Object` class 212
 rules to keep contract 192
 superclass 169, 230
 `toString()` 316

P

packages, Java API 154–155
 compiling and running 687
 directory structure 686
 organizing code 686

the index

- preventing class name conflicts 685
- putting classes in 686
- reverse domain package names 685
- panels, GUI 511–522
- parallelStream 695
- parallel streams 695
- parameterized types 137
- parameters. *See also* arguments
 - lambda expressions 402
 - and local variables 85
 - matching with instance variable types 76
 - and methods 74, 76
 - type. *See* type parameters
- parse methods 294
- pass-by-value/pass-by-copy 77
- Path interface 572–573
- Paths class 572–573
- percent sign (%) in format String 297–300
- PetShop program 220–228
- Phrase-O-Matic code 18–19
- pipelines, stream. *See* stream pipeline
- polymorphism
 - abstract classes 202–205, 208–209
 - arguments and return types 189–190
 - and exceptions 436–440
 - with generic types 321, 358–362
 - Graphics superclass 474
 - interface implementation 226–231
 - and List versus ArrayList 313
 - methods 205–206
 - Object class 210–215
 - reference and object types as different 188–189
- post-increment operator 105
- Predicate 375, 402
- prep code 99
 - SimpleStartup class 100–101
 - StartupBust class 144–145
- primitive variables 49, 51
 - in arrays 59
 - bit-size space 241
 - comparing objects 86
 - declaring 50–52
 - ranges for variables 51
 - as reserved words 53
 - saving objects 545
 - wrapping 290–294
- print() 595
- printf() 296
- println() 595
- printStackTrace() 429
- print versus println 15
- PrintWriter 595, 596, 602
- private access modifier 81, 689
- private constructor 191, 251, 278, 282
- protected access level 689, 690
- protected access modifier 689
- pseudocode. *See* prep code
- public access modifier 81–82, 689
- puzzles
 - Five-Minute Mystery 67, 92, 270–271, 415, 669
 - GUI-Cross 536
 - Heap o' Trouble 66
 - JavaCross 22, 120, 164, 456, 536
 - Mixed Messages 90, 121
 - Pool Puzzle 24, 44, 91, 196, 234, 416, 506
- Q**
- queries
 - returning Optional value 410–414
 - stream pipelines as queries on collection 380, 385
 - terminal operation options 410–412
- QuizCardBuilder 560–563
- QuizCardPlayer 567–569